

Generating Images Using a Deep Convolutional GANs



Author: Kai Deng

Supervisor: Serhiy Yanchuk

A thesis submitted in partial fulfilment of the
requirements for the degree of
MSc Mathematical Modelling and Machine Learning

School of Mathematical Sciences,
University College Cork,
Ireland

September 2024

Declaration of Authorship

This report is wholly the work of the author, except where explicitly stated otherwise. The source of any material which was not created by the author has been clearly cited.

Date: 13/09/2024

Signature: Kai Deng

Acknowledgements

I would like to express my deepest gratitude to my supervisor, Dr Serhiy Yanchuk, for his invaluable guidance, continuous support, and insightful feedback throughout the course of my research. His expertise and encouragement have been instrumental in the successful completion of this thesis.

Abstract

Generative Adversarial Networks (GANs) have gained attention for their ability to generate realistic data in various fields. In this thesis, I explore the performance of GANs in generating high-quality cat images, utilizing the Animal Faces-HQ dataset, which includes 16,130 high-resolution images. The focus is on comparing different architectural designs, specifically convolutional and dense layers, and investigating how data augmentation techniques influence model performance.

The findings suggest that convolutional architectures offer advantages in capturing spatial features, contributing to improved image quality. While data augmentation introduces diversity into the dataset, it also presents optimization challenges that can complicate the training process. These results highlight important considerations for future studies aiming to enhance GAN-generated images. The insights gained from this work may guide further exploration of more advanced architectures and larger datasets.

Contents

1	Introduction	5
2	Historical Models of Image Generation	7
2.1	Noise Contrastive Estimation (NCEs)	7
2.1.1	NCE Architecture	8
2.1.2	NCE Objective Function	9
2.1.3	Comparison with GANs	10
2.1.4	Applications of NCE	11
2.1.5	Limitations of NCE	12
2.2	Variational Autoencoders (VAEs)	12
2.2.1	VAE Architecture	13
2.2.2	VAE Objective Function	13
2.2.3	Comparison with GANs	14
2.2.4	Applications of VAEs	15
2.2.5	Limitations of VAEs	15
2.3	Diffusion Models	16
2.3.1	Diffusion Model Architecture	16
2.3.2	Diffusion Model Objective Function	17
2.3.3	Comparison with GANs	17
2.3.4	Applications of Diffusion Models	18
2.3.5	Limitations of Diffusion Models	19
3	Theoretical Background	21
3.1	Generative Adversarial Networks (GANs)	21
3.2	GAN Architecture	22

3.3	Objective Function of GAN	22
3.3.1	Discriminator's Loss Function	23
3.3.2	Generator's Loss Function	23
3.3.3	Modification for Vanishing Gradient Problem	23
3.4	GAN Training Process	24
3.4.1	Distribution Changes During GAN Training	25
3.4.2	Mathematical Formulation during GAN Training	26
3.5	Evaluating GAN Performance	29
3.6	Limitations of Accuracy in Evaluating GANs	30
4	Results and Discussion	32
4.1	Standard GAN Versus Other GAN Realizations	32
4.2	GAN With Convolutional or Dense layers	33
4.3	Exploring Layer Depth	37
4.4	Impact of Data Augmentation on Model Performance	38
4.5	Applying the Model to the Animal Faces-HQ Dataset	39
4.5.1	Model Structure and Training	39
4.5.2	Generated Results	42
4.6	Discussion	44
A	Appendix	46

List of Figures

2.1	NCE Architecture.	9
2.2	VAE Architecture.	13
2.3	Diffusion Model Architecture.	17
3.1	GAN Architecture.	22
3.2	Comparison of $\log(1 - D)$ and $-\log(D)$.	24
3.3	Diagram of GAN's Training Process.	25
3.4	GAN Training Accuracy over Epochs	30
3.5	Generated Images from GAN	31
4.1	Generator Architecture with Dense and Convolutional Layers	34
4.2	Discriminator Architecture with Dense and Convolutional Layers	35
4.3	Comparison of GAN Performance at 3000 Epochs	35
4.4	FID Scores Across 3000 Epochs	36
4.5	Cat Faces Generated by GAN	43

List of Tables

4.1	Average FID Scores for Different GAN Architectures (Lower is Better)	37
4.2	Average FID Scores for Different Data Augmentation Techniques	38

Chapter 1

Introduction

Generative Adversarial Networks (GANs) have emerged as a transformative tool in generative modeling, framing the problem as a competition between two networks: a generator that creates synthetic data from noise, and a discriminator that differentiates between real and generated data [1]. Since their introduction in 2014, GANs have found applications in fields such as materials science, radiology, and computer vision [2], [3], [4]. For instance, CycleGAN has been applied in medical imaging, enhancing tasks like liver lesion classification through synthetic image augmentation, outperforming traditional methods in sensitivity and specificity [5]. Similarly, StyleGAN has shown effectiveness in image deformation and style transfer [6], further expanding the reach of GANs in the computer vision field, where they generate data without explicitly modeling probability density functions [3].

However, despite their success, training GANs presents notable challenges, including issues like mode collapse, training instability, and the high computational demands required for effective performance. Evaluating GANs is also complex, as traditional metrics such as accuracy are insufficient to measure the quality and diversity of generated data. These difficulties have driven the development of various architectures and training methods aimed at improving the stability and effectiveness of GANs.

The objective of this thesis is to contribute to a clearer understanding of the factors that influence GAN performance, particularly in generating real-

istic images. For this purpose, I use the Animal Faces-HQ (AFHQ) dataset, containing 16,130 high-resolution images at 512×512 pixels, to train a GAN model specifically for generating realistic cat images. The high resolution presents both opportunities and challenges, as it requires careful attention to the model’s architecture and training to avoid overfitting or underfitting.

This thesis is structured as follows: Chapter 1 introduces the research background, the motivation behind this work, and outlines the key objectives of the thesis. Chapter 2 offers an overview of previous models for image generation, including methods such as Noise Contrastive Estimation, Variational Autoencoders, and Diffusion Models, emphasizing their architectures, objectives, and applications. Chapter 3 focuses on Generative Adversarial Networks (GANs), discussing their theoretical foundation, including key concepts such as objective functions and training dynamics. Chapter 4 describes the experimental work, covering model selection, architecture exploration, the effects of data augmentation, and the application of the GAN model to the Animal Faces-HQ dataset. Lastly, it presents the results, their implications, and suggests directions for future research.

Chapter 2

Historical Models of Image Generation

This chapter provide an overview of three important generative modeling techniques: Noise Contrastive Estimation (NCE), Variational Autoencoders (VAEs), and Diffusion Models. These models have each played a significant role in the evolution of generative modeling and continue to influence modern approaches in the field.

2.1 Noise Contrastive Estimation (NCEs)

Noise Contrastive Estimation (NCE) was introduced in 2010 by Gutmann and Hyvärinen as a method for estimating parameters in unnormalized probabilistic models. It offers an efficient alternative to Maximum Likelihood Estimation (MLE), particularly in cases where MLE can become computationally expensive, especially with large-scale models. NCE reframes the challenge of normalizing probability distributions into a more manageable binary classification problem [7].

The core idea behind NCE is to treat MLE as a binary classification task. Traditionally, training models for unnormalized probability distributions using MLE involves calculating the partition function, which can be computationally infeasible for large datasets. NCE mitigates this difficulty

by incorporating noise samples drawn from a known distribution. The model is then trained to differentiate between real data and noise samples, with higher probabilities assigned to real data and lower probabilities to noise. This approach allows the model to learn the data distribution without the need for explicit normalization [8].

2.1.1 NCE Architecture

In NCE, the likelihood of a data point x is reformulated as the probability that it comes from the real data distribution rather than from the noise distribution. As depicted in Figure 2.1, the architecture of a typical NCE model includes an input layer, a hidden layer, and an output layer. The input layer receives both real and noise samples, while the hidden layer learns representations of these samples. The output layer functions as a binary classifier, producing probabilities that indicate whether a given sample is real or noise.

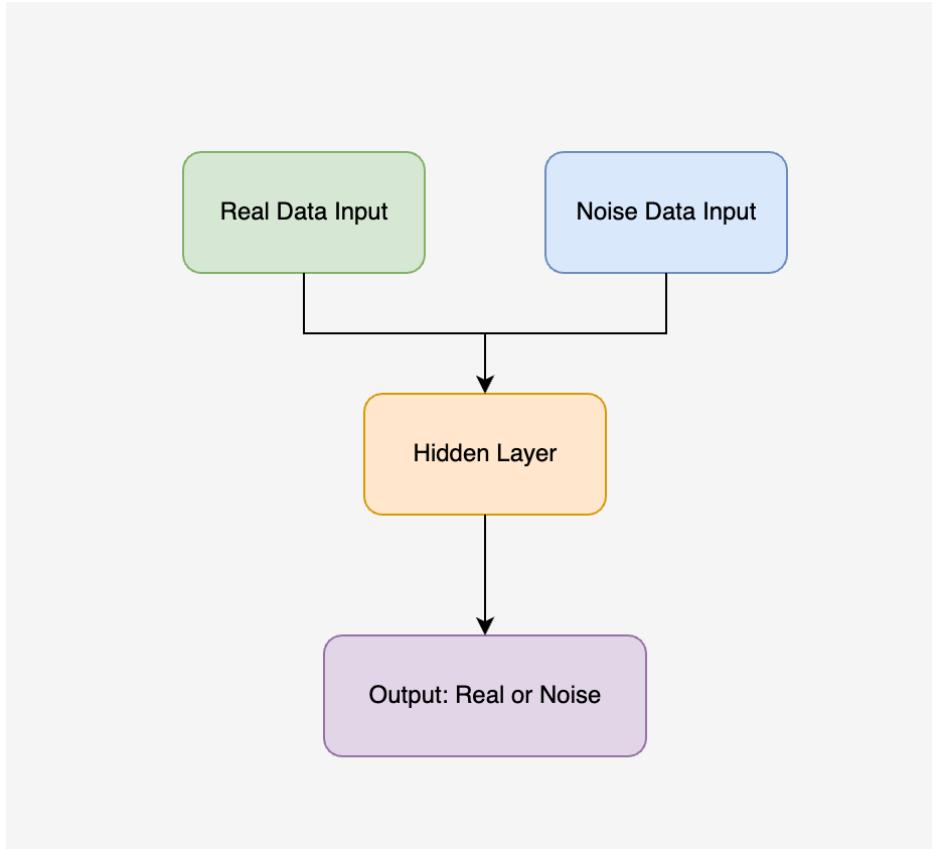


Figure 2.1: NCE Architecture.

- **Data and Noise Samples:** NCE introduces noise samples from a known distribution to compare with actual data. These noise samples act as negative examples in the classification task, while the real data serves as positive examples [7].
- **Binary Classification:** The task of differentiating between real and noise samples is framed as a binary classification problem, which can be optimized using standard logistic regression techniques [9].

2.1.2 NCE Objective Function

In NCE, the probability $P(y = 1|x)$, where $y = 1$ indicates that x is a real sample, is defined as:

$$P(y = 1|x) = \frac{p_\theta(x)}{p_\theta(x) + kp_{\text{noise}}(x)} \quad (2.1)$$

Where $p_\theta(x)$ is the unnormalized probability assigned to the sample x by the model, $p_{\text{noise}}(x)$ is the probability assigned to the noise sample, and k is the ratio of noise samples to real samples.

The corresponding probability that a sample x is from the noise distribution is given by:

$$P(y = 0|x) = \frac{kp_{\text{noise}}(x)}{p_\theta(x) + kp_{\text{noise}}(x)} \quad (2.2)$$

The NCE objective then seeks to maximize the log-probabilities of correctly classifying real and noise samples:

$$\mathcal{L}_{NCE} = \sum_{i=1}^N \left[\log P(y = 1|x_i) + \sum_{j=1}^k \log P(y = 0|x_j) \right] \quad (2.3)$$

This method circumvents the need for computing the partition function, which is typically required in Maximum Likelihood Estimation (MLE), making NCE particularly useful for large-scale models [8].

2.1.3 Comparison with GANs

Noise Contrastive Estimation (NCE) and Generative Adversarial Networks (GANs) are both prominent methods in generative modeling, but they take different approaches to training and model estimation.

- **Training Stability:** One of NCE's key advantages over GANs is the stability it provides during training. GANs frequently suffer from challenges such as mode collapse and instability due to the adversarial nature of the training process, where the generator and discriminator are pitted against each other. NCE, by contrast, frames the training process as a straightforward binary classification task, which tends to be more stable and can be optimized with logistic regression [9].
- **Computational Efficiency:** Training GANs requires simultaneously

optimizing two models—the generator and the discriminator—which can lead to higher computational costs, particularly when tuning their interactions for better performance. NCE, on the other hand, simplifies the problem by reducing it to a comparison between real data and noise samples, avoiding the adversarial framework and offering a more computationally efficient solution, especially for large-scale models [10].

- **Handling Unnormalized Models:** NCE is particularly suited for training unnormalized probabilistic models, where computing the partition function is impractical or impossible. GANs, however, are focused on generating realistic samples and, while they do not require explicit normalization, they do not address the problem of unnormalized models as directly as NCE does [11].
- **Model Interpretability:** In NCE, the model explicitly learns to estimate the probability of real data versus noise, providing insights into the data distribution. In comparison, GANs focus on generating realistic samples without explicitly modeling probabilities, which can make their internal latent space less interpretable than that of NCE-based models.
- **Applications in Language Models and Word Embeddings:** NCE is particularly beneficial in tasks such as word embeddings and large-scale language models, where normalizing the likelihood function over vast vocabularies is computationally prohibitive. GANs have been applied in text generation, but NCE’s efficiency in handling large-scale vocabulary models makes it more suitable for these problems [11].

2.1.4 Applications of NCE

NCE is highly effective in a variety of machine learning tasks, particularly those that involve large datasets and unnormalized models:

- **Word Embeddings:** NCE is widely used in training word embeddings, where the size of the vocabulary makes normalization computationally infeasible.

- **Language Models:** Beyond word embeddings, NCE has been applied to training large-scale language models, where traditional likelihood-based methods may become computationally expensive.
- **Energy-Based Models:** NCE is also effective in training energy-based models, which typically require an intractable partition function for normalization.

NCE allows models to scale efficiently, making it an invaluable tool in areas ranging from natural language processing to computer vision.

2.1.5 Limitations of NCE

While NCE has proven to be a powerful estimation technique, it is not without limitations. A key challenge lies in selecting an appropriate noise distribution. Poor choices in this regard can lead to suboptimal parameter estimates and slower convergence during training [8].

- **Noise Distribution Sensitivity:** The success of NCE relies heavily on choosing a noise distribution that is sufficiently different from the real data distribution. If the noise distribution is poorly chosen, the model may struggle to differentiate between real and noise samples [8].
- **Handling Complex Data Distributions:** NCE may encounter difficulties with highly complex data distributions, particularly in cases where defining an appropriate noise distribution is challenging [8].

2.2 Variational Autoencoders (VAEs)

Variational Autoencoders (VAEs), introduced in 2013, represent one of the foundational approaches to generative modeling. The primary goal of VAEs is to model the underlying distribution of data by learning a compressed latent representation, denoted as z , from the input data x .

2.2.1 VAE Architecture

VAEs consist of two main components: an encoder and a decoder. The encoder maps input data into a latent space, where the latent variables are generally assumed to follow a Gaussian distribution. This assumption simplifies the training process and enables techniques such as the reparameterization trick, which facilitates backpropagation through stochastic layers [12]. The decoder then reconstructs the input data from the latent variables, ensuring that the essential characteristics of the data distribution are captured.

As illustrated in Figure 2.2, the VAE architecture comprises:

- **Encoder:** The encoder takes the input data x and compresses it into a latent representation z . The latent variables are sampled from a Gaussian distribution, which aids in simplifying the optimization process.
- **Decoder:** The decoder reconstructs the original data x' from the latent representation z , aiming to produce data that closely resembles the input.

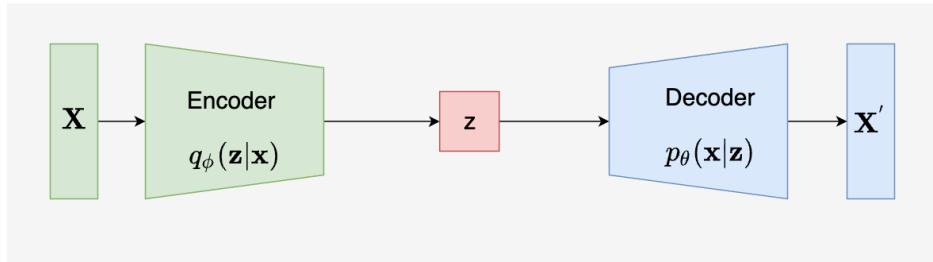


Figure 2.2: VAE Architecture.

2.2.2 VAE Objective Function

The VAE loss function is designed to balance two objectives: accurate data reconstruction and regularization of the latent space. The total loss combines a reconstruction loss with a Kullback-Leibler (KL) divergence term, which helps the model learn a smooth, continuous representation of the latent space [13].

The VAE aims to maximize the variational lower bound, encouraging both accurate reconstruction and a well-structured latent space. This, in turn, enables the model to generate new data by sampling from the latent space and reconstructing it using the decoder.

The loss function of a VAE is expressed as follows:

$$\mathcal{L} = \mathbb{E}_{q_\phi(z|x)}[\log p_\theta(x|z)] - D_{KL}(q_\phi(z|x)\|p(z)) \quad (2.4)$$

Where:

- \mathcal{L} : The total loss that the model seeks to minimize.
- $\mathbb{E}_{q_\phi(z|x)}[\log p_\theta(x|z)]$: This term represents the expected log-likelihood of the reconstructed data x' given the latent variable z . The distribution $q_\phi(z|x)$ corresponds to the encoder's approximation of the posterior, while $p_\theta(x|z)$ represents the decoder's likelihood of reconstructing the input data.
- $D_{KL}(q_\phi(z|x)\|p(z))$: The KL divergence measures the difference between the encoder's learned latent distribution $q_\phi(z|x)$ and the prior distribution $p(z)$, which is typically Gaussian.

2.2.3 Comparison with GANs

When compared to Generative Adversarial Networks (GANs), VAEs exhibit several advantages. One of the most notable strengths of VAEs is their training stability. In contrast to GANs, which involve training both a generator and a discriminator—a process that can lead to issues such as mode collapse—VAEs have a single objective function. This objective combines reconstruction loss and KL divergence, making the training process more straightforward and less prone to instability [12].

Another key strength of VAEs is their ability to generate smooth transitions between data points. This capability is particularly valuable for applications that require meaningful interpolation in the latent space, such as image generation, anomaly detection, and data imputation [14][15]. GANs,

on the other hand, do not explicitly model the latent space, which can limit their interpretability and flexibility in certain tasks.

However, GANs are often favored for generating sharper and more realistic images, particularly in high-resolution tasks. The adversarial loss in GANs encourages the generator to produce outputs that closely resemble real data, while VAEs, due to their reliance on Gaussian priors, tend to generate slightly blurrier images [16]. Nonetheless, VAEs offer greater flexibility and scalability, as they can be trained using standard gradient descent methods and do not require the complex adversarial framework inherent to GANs.

2.2.4 Applications of VAEs

The versatility of VAEs extends to a wide range of applications, many of which benefit from the structured latent space that VAEs provide:

- **Image Generation:** VAEs can generate new images by sampling from the latent space and decoding the latent vectors into realistic image representations.
- **Anomaly Detection:** VAEs are often used to identify outliers in data by examining reconstruction errors. Data points with high reconstruction loss may be flagged as anomalies.
- **Data Imputation:** VAEs are capable of filling in missing data by reconstructing the incomplete data from latent representations, making them useful for handling datasets with gaps or missing values.

2.2.5 Limitations of VAEs

While VAEs have proven effective in many applications, they also have certain limitations:

- **Blurred Outputs:** Due to the Gaussian prior assumption in the latent space, VAEs often generate blurrier outputs compared to GANs, which may affect performance in tasks requiring high-resolution and sharp images [16].

- **Difficulty with Discrete Data:** VAEs struggle to effectively model discrete data, as backpropagation through continuous latent variables is not well-suited to handle discrete structures [17].
- **Limited Diversity:** VAEs may not capture the full complexity and diversity of the data distribution as effectively as GANs, particularly when generating high-resolution images [18].

2.3 Diffusion Models

Diffusion models, emerging in the early 2020s, represent a significant advancement in generative modeling. These models progressively add noise to data and then learn to reverse this process, effectively "denoising" it back to its original form. This iterative process distinguishes diffusion models from traditional approaches like GANs and VAEs [19].

2.3.1 Diffusion Model Architecture

The core idea behind diffusion models relies on a Markov chain where noise is added in the forward process, starting from a simple distribution (e.g., Gaussian), and reversed through a learned denoising mechanism [20], [21]. This approach has proven effective in generating high-quality samples across various domains, such as image synthesis, audio generation, and medical imaging [22], [23]. In several benchmarks, diffusion models have outperformed GANs [19], [24].

The process involves two key steps, as shown in Figure 2.3:

- **Forward Process:** Gradually adds Gaussian noise to data x_0 , creating noisy versions of the data x_1, x_2, \dots, x_T . Each step increases the noise level, leading to a completely noisy version z .
- **Reverse Process:** The model learns to reverse the noise addition process, starting from the fully noisy version z , and progressively denoising it to recover data that resembles the original input x_0 .

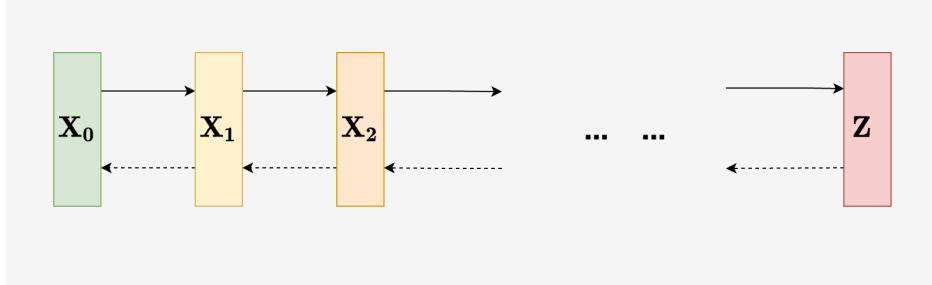


Figure 2.3: Diffusion Model Architecture.

2.3.2 Diffusion Model Objective Function

The training objective for diffusion models is to minimize the difference between the real data distribution and the distribution of generated data across all time steps:

$$L = \sum_{t=1}^T \mathbb{E}_{x_0, \epsilon} [\|\epsilon - \epsilon_\theta(x_t, t)\|^2] \quad (2.5)$$

where:

- L : The loss function to be minimized.
- T : The total number of time steps in the diffusion process.
- x_0 : The original data sample.
- x_t : The data at time step t , after adding noise.
- ϵ : The noise added to the data at each step.
- $\epsilon_\theta(x_t, t)$: The model's estimate of the noise at time step t .

2.3.3 Comparison with GANs

Diffusion models and Generative Adversarial Networks (GANs) represent two distinct approaches in generative modeling, each with its strengths and weaknesses. A key advantage of diffusion models is their training stability.

Unlike GANs, which often suffer from unstable dynamics due to the adversarial competition between the generator and discriminator, diffusion models follow a simpler, noise-reversal process that avoids these issues [25]. This stability helps diffusion models avoid problems like mode collapse, where GANs fail to capture the diversity of the data distribution [26], [27].

Diffusion models also excel in generating diverse, high-quality samples through a gradual denoising process, allowing for controlled generation [28]. In contrast, GANs tend to produce sharper but less diverse images, often overfitting to specific modes of the data [29], [30]. This trade-off between sharpness and diversity is a known limitation of GANs [31].

However, GANs remain preferred for tasks requiring extremely high-resolution and photorealistic images, such as human face generation, where they outperform diffusion models in terms of detail and realism [32], [33].

2.3.4 Applications of Diffusion Models

Diffusion models have found a wide range of applications, especially in fields where stability and quality of generation are important. Some notable applications include:

- **Image Generation:** Diffusion models have proven effective in generating photorealistic images, similar to GANs, but with more stable training dynamics. Recent innovations, such as classifier-free guidance, have further enhanced the quality of generated samples by allowing the model to generate data without relying on an explicit classifier [34].
- **Text-to-Image Generation:** Advancements in diffusion models have also been applied to text-to-image synthesis, where a text prompt is converted into a corresponding image. These models can produce diverse outputs based on input descriptions. Additionally, techniques like Denoising Diffusion Implicit Models (DDIMs) have accelerated the sampling process, making diffusion models more practical for real-time applications [35], [36].

- **Speech Synthesis:** Diffusion models have been applied to generating high-quality audio data. For example, they are used in text-to-speech systems to generate realistic human speech. These models have demonstrated remarkable performance in generating high-fidelity audio outputs [22], [20].
- **Anomaly Detection and Medical Imaging:** Similar to VAEs, diffusion models can be used to detect anomalies by evaluating how well a noisy sample can be denoised. Poor reconstructions may indicate that the input is anomalous or different from the training data. Diffusion models have also been applied in medical image segmentation tasks, such as the MedSegDiff model, which enhances segmentation by leveraging diffusion processes [23]. These models have shown their ability to handle complex data structures, proving their versatility across different domains.

Recent innovations in diffusion models have further enhanced their applicability and efficiency. By reducing the computational burden associated with the iterative sampling process, models like DDIMs maintain the generative capabilities of traditional diffusion models while improving their practicality for real-time applications [37].

2.3.5 Limitations of Diffusion Models

Despite the significant advancements that diffusion models have brought to generative modeling, they are not without limitations, which can be categorized into three primary areas: computational intensity, generation speed, and sample sharpness.

- **Computationally Intensive:** Diffusion models are computationally expensive due to their iterative nature, which involves progressively adding and removing noise from data. This process requires significantly more computational resources compared to GANs, which generate data in a single forward pass through the generator [38], [39].

This computational burden can limit their practical use, especially in scenarios requiring rapid generation [40].

- **Generation Speed:** Diffusion models are slower than GANs in terms of sample generation. Producing a single sample often requires hundreds or even thousands of time steps, significantly increasing the generation time compared to the near-instantaneous output of GANs [41], [35]. This can be a critical drawback in real-time applications.
- **Sample Sharpness:** While diffusion models generate diverse outputs, they may not always match the photorealism and fine detail achieved by GANs, especially in tasks requiring intricate details [19], [42]. This can affect their suitability for applications such as high-resolution image generation or medical imaging [40].

Chapter 3

Theoretical Background

This chapter provides an overview of Generative Adversarial Networks (GANs), focusing on their core architecture, objective function, and training dynamics. It introduces the interplay between the generator and discriminator in the adversarial training process, as well as the Fréchet Inception Distance (FID) metric for evaluating the quality and diversity of generated images.

3.1 Generative Adversarial Networks (GANs)

Generative Adversarial Networks (GANs) were introduced by Goodfellow et al. in 2014 and have become a widely used approach in generative modeling. GANs consist of two neural networks: a generator (G) and a discriminator (D), trained simultaneously through an adversarial process. The generator is tasked with producing synthetic data samples, while the discriminator tries to distinguish between real and generated samples. This architecture leads to a game-like dynamic between the two networks, where the generator improves its ability to create realistic data while the discriminator becomes better at distinguishing real from fake data [43], [44] and provide feedback to the generator, helping it to improve and generate more realistic samples over time [45].

3.2 GAN Architecture

The structure of GANs consists of two key components:

- **Generator (G):** The generator maps random noise z from a prior distribution, such as Gaussian or uniform, to synthetic data samples $G(z)$ that aim to replicate real data distributions.
- **Discriminator (D):** The discriminator receives input data, which can be either real data x or generated data $G(z)$, and outputs the probability that the data is real. The discriminator is trained to classify real samples as real and generated samples as fake [45].

The following figure illustrates the GAN architecture, where the generator produces synthetic data, and the discriminator learns to differentiate real from fake data.

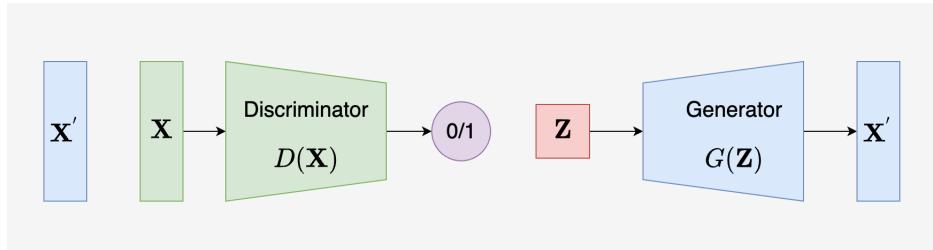


Figure 3.1: GAN Architecture.

3.3 Objective Function of GAN

The training of GANs involves optimizing two neural networks simultaneously through a minimax game. The objective function is defined as follows:

$$\min_G \max_D V(D, G) = \mathbb{E}_{x \sim p_{data}(x)}[\log D(x)] + \mathbb{E}_{z \sim p_z(z)}[\log(1 - D(G(z)))] \quad (3.1)$$

where :

- $z \sim p_z(z)$: The distribution from which the noise z is sampled. The generator transforms this noise into complex, high-dimensional data.

- $x \sim p_{data}(x)$: A sample x drawn from the true data distribution $p_{data}(x)$.
- $D(G(z))$: The discriminator's prediction for generated data $G(z)$.
- $D(x)$: The discriminator's prediction for real data x .
- $\mathbb{E}_{x \sim p_{data}(x)}[\log D(x)]$: The expectation of $\log D(x)$ over the real data distribution.
- $\mathbb{E}_{z \sim p_z(z)}[\log(1 - D(G(z)))]$: The expectation of $\log(1 - D(G(z)))$ over the noise distribution.

3.3.1 Discriminator's Loss Function

The discriminator's loss function measures its ability to distinguish between real and generated data. It is expressed as:

$$\mathcal{L}_D = -\mathbb{E}_{x \sim p_{data}(x)}[\log D(x)] - \mathbb{E}_{z \sim p_z(z)}[\log(1 - D(G(z)))] \quad (3.2)$$

The discriminator minimizes this loss function to improve its ability to classify real and generated data correctly.

3.3.2 Generator's Loss Function

The generator's loss function is designed to encourage it to generate data that can fool the discriminator. It is given by:

$$\mathcal{L}_G = -\mathbb{E}_{z \sim p_z(z)}[\log D(G(z))] \quad (3.3)$$

The generator minimizes this loss function to produce data that the discriminator classifies as real.

3.3.3 Modification for Vanishing Gradient Problem

To address the vanishing gradient problem during GAN training, the generator's loss is often modified. The original form $\log(1 - D(G(z)))$ can result

in small gradients when the discriminator accurately identifies fake data. Therefore, the alternative loss function for the generator is:

$$\mathcal{L}_G = \mathbb{E}_{z \sim p_z(z)}[-\log(D(G(z)))] \quad (3.4)$$

This modification ensures larger gradients and more stable training.

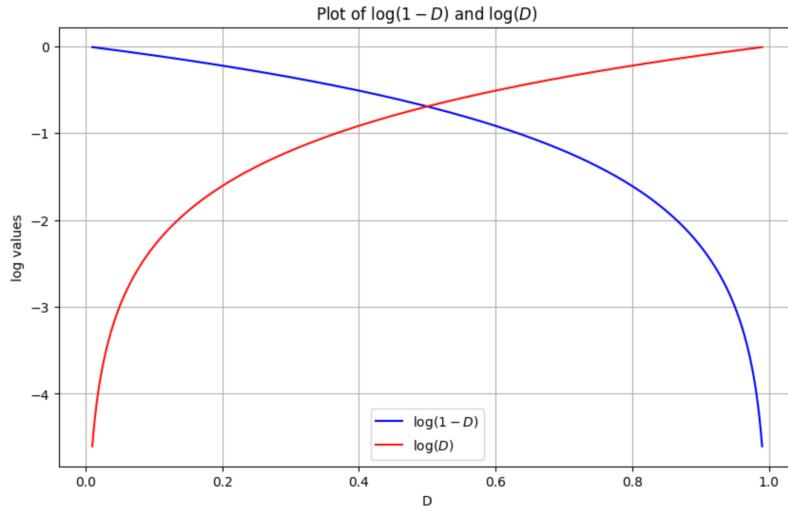


Figure 3.2: Comparison of $\log(1 - D)$ and $-\log(D)$.

3.4 GAN Training Process

During the training process of the GAN model, the generator and the discriminator are in constant competition. The generator aims to produce increasingly realistic samples to deceive the discriminator, while the discriminator works to accurately distinguish between real and generated samples. Initially, the generator produces samples that are of low quality and deviate significantly from the real data, as illustrated by the blue sine wave in Figure (a). As training progresses, the generator improves, and the generated samples become more refined, as shown by the smoother blue lines in Figures (b) and (c). Simultaneously, the distribution of the real samples and generated samples gradually align, and finally, in the optimal state represented by Figure (d), $p_{data}(x) = p_g(x)$, meaning that the generated samples closely

match the real samples. Throughout the training process, the distribution of the generated samples transitions from an initially noisy and scattered state to one that closely mirrors the real data distribution, indicating a significant enhancement in the quality of the generated samples.

3.4.1 Distribution Changes During GAN Training

A simple diagram shows how distribution change in GAN training.

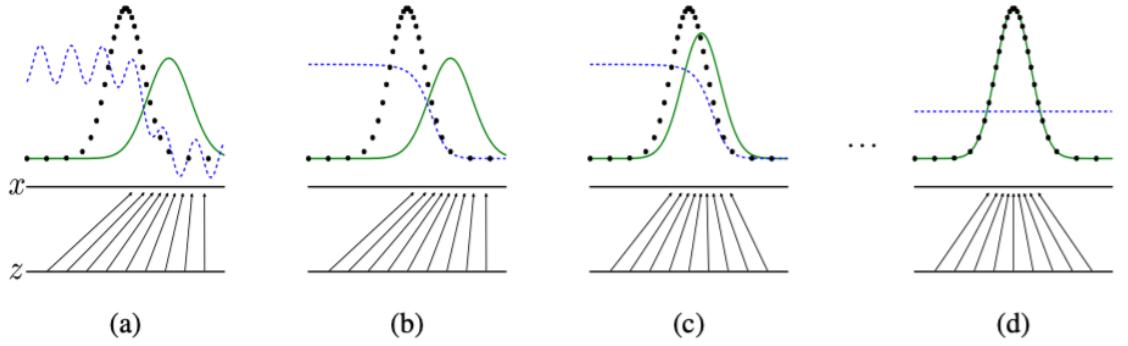


Figure 3.3: The green curve represents the distribution of generated samples. Initially, the generated samples may differ significantly from the real samples. As training progresses, the generated samples' distribution gradually approaches the real samples. The black dots represent the distribution of real samples, which remains unchanged throughout the training process and represents the target distribution. The blue dashed line represents the discriminator's output probability distribution. At the beginning of training, the discriminator can easily distinguish between real and generated data, resulting in a strong classification boundary. As training progresses and the generated data becomes more realistic, the discriminator's ability to differentiate between the two distributions weakens. Eventually, the discriminator's output approaches 0.5, indicating it can no longer effectively distinguish between real and generated data. The lines labeled x and z below represent the distribution of samples in the latent space. During GAN training, samples from the latent space z are mapped to the data space x through the generator.

Source: [46]

As training progresses, the samples generated by the generator gradually

become indistinguishable from the real data, leading to a minimized error between the two distributions. Initially, the generated samples exhibit significant differences from the real data, starting from a noisy and dispersed state. However, over time, the generator learns to map the latent space to the real data distribution more accurately. This results in the generated samples converging toward the real data distribution, ultimately aligning closely with the target distribution.

3.4.2 Mathematical Formulation during GAN Training

1. Problem Setup

In a Generative Adversarial Network (GAN), the objective is to train a generator G and a discriminator D to generate data that resembles the real data distribution. The objective function to be maximized is given by:

$$V(D, G) = \mathbb{E}_{x \sim p_{\text{data}}(x)}[\log D(x)] + \mathbb{E}_{z \sim p_z(z)}[\log(1 - D(G(z)))] \quad (3.5)$$

2. Rewriting the Objective Function

First, the objective function is rewritten in integral form:

$$V(D, G) = \int p_{\text{data}}(x) \log D(x) dx + \int p_z(z) \log(1 - D(G(z))) dz \quad (3.6)$$

By changing variables $x' = G(z)$, the generated data is represented as x' , which corresponds to samples produced by the generator G . This allows the second term to be rewritten as an integral over the generated data distribution $p_g(x)$. The integral now reflects the contribution of the generated samples to the objective function.

$$\int p_g(x) \log(1 - D(x)) dx \quad (3.7)$$

Thus, the objective function becomes:

$$V(D, G) = \int [p_{\text{data}}(x) \log D(x) + p_g(x) \log(1 - D(x))] dx. \quad (3.8)$$

3. Deriving the Optimal Discriminator

To find the optimal discriminator D^* , it needs to take the derivative of the objective function with respect to $D(x)$ and set it to zero.

Let:

$$f(D(x)) = p_{\text{data}}(x) \log D(x) + p_g(x) \log(1 - D(x)). \quad (3.9)$$

Taking the derivative with respect to $D(x)$:

$$\frac{d}{dD(x)} f(D(x)) = \frac{p_{\text{data}}(x)}{D(x)} - \frac{p_g(x)}{1 - D(x)}. \quad (3.10)$$

Setting the derivative to zero:

$$\frac{p_{\text{data}}(x)}{D(x)} = \frac{p_g(x)}{1 - D(x)} \quad (3.11)$$

Solving equation (3.11):

$$D(x) = \frac{p_{\text{data}}(x)}{p_{\text{data}}(x) + p_g(x)}. \quad (3.12)$$

4. Optimal Discriminator Formula

Therefore, the optimal discriminator D^* is given by:

$$D^*(x) = \frac{p_{\text{data}}(x)}{p_{\text{data}}(x) + p_g(x)}. \quad (3.13)$$

- When $p_{\text{data}}(x)$ is much larger than $p_g(x)$, $D^*(x) \approx 1$, indicating that the data point is almost certainly from the real data.
- When $p_{\text{data}}(x)$ is much smaller than $p_g(x)$, $D^*(x) \approx 0$, indicating that the data point is almost certainly from the generated data.
- When $p_{\text{data}}(x)$ is close to $p_g(x)$, $D^*(x) \approx 0.5$, indicating that the discriminator cannot confidently determine whether the data point is real

or generated, giving each a 50% probability.

5. Verifying the Optimal Discriminator

To verify that this D^* maximizes the objective function, substitute D^* back into the objective function:

$$V(D^*, G) = \int \left[p_{\text{data}}(x) \log \left(\frac{p_{\text{data}}(x)}{p_{\text{data}}(x) + p_g(x)} \right) + p_g(x) \log \left(1 - \frac{p_{\text{data}}(x)}{p_{\text{data}}(x) + p_g(x)} \right) \right] dx. \quad (3.14)$$

Since:

$$1 - D^*(x) = 1 - \frac{p_{\text{data}}(x)}{p_{\text{data}}(x) + p_g(x)} = \frac{p_g(x)}{p_{\text{data}}(x) + p_g(x)}, \quad (3.15)$$

substituting this in:

$$V(D^*, G) = \int \left[p_{\text{data}}(x) \log \left(\frac{p_{\text{data}}(x)}{p_{\text{data}}(x) + p_g(x)} \right) + p_g(x) \log \left(\frac{p_g(x)}{p_{\text{data}}(x) + p_g(x)} \right) \right] dx. \quad (3.16)$$

This objective function represents the negative of the cross-entropy, which is maximized when $D(x) = D^*(x)$. When $D(x) = 0.5$, the discriminator cannot distinguish between real and generated data, indicating that the generator has produced samples that closely resemble the real data. Maximizing the negative cross-entropy aligns the generated data distribution with the real data distribution.

6. Conclusion

Through the above derivation, it has shown that given the generator G , the optimal form of the discriminator D is:

$$D^*(x) = \frac{p_{\text{data}}(x)}{p_{\text{data}}(x) + p_g(x)}. \quad (3.17)$$

This demonstrates that the optimal discriminator D^* outputs the probability that the input data comes from the real data distribution. This formula provides a theoretical foundation for training GANs, guiding the updates to the generator G so that its generated data gradually approaches the real data

distribution.

3.5 Evaluating GAN Performance

Fréchet Inception Distance (FID) is a metric commonly used in Generative Adversarial Network (GAN) models to quantify the dissimilarity between two image distributions [47]. It measures the distance between the distributions of real images and generated images, providing a numerical assessment of the quality of generated images. FID has gained prominence in evaluating the performance of GANs due to its ability to capture both the quality and diversity of generated images [48]. The following is the objective function for FID:

$$\text{FID} = \|\mu_r - \mu_g\|^2 + \text{Tr}(\Sigma_r + \Sigma_g - 2(\Sigma_r \Sigma_g)^{1/2}) \quad (3.18)$$

$$\mu_r = \frac{1}{N} \sum_{i=1}^N f(x_i), \quad \Sigma_r = \frac{1}{N} \sum_{i=1}^N (f(x_i) - \mu_r)(f(x_i) - \mu_r)^T \quad (3.19)$$

$$\mu_g = \frac{1}{M} \sum_{i=1}^M f(G(z_i)), \quad \Sigma_g = \frac{1}{M} \sum_{i=1}^M (f(G(z_i)) - \mu_g)(f(G(z_i)) - \mu_g)^T \quad (3.20)$$

where:

- μ_r and μ_g : The feature means of the real and generated images, respectively.
- Σ_r and Σ_g : The feature covariance matrices of the real and generated images, respectively.
- Tr : The trace (the sum of the diagonal elements of the matrix).
- f : The feature extraction function, which extracts feature vectors from images using the Inception network. These feature vectors are used to

compute the mean and covariance for both real and generated images.

A lower FID value indicates that the distribution of the generated images is closer to that of real images, reflecting higher quality and diversity in the generated images [49]. Specifically, a FID value below 10 is considered to represent very high-quality generated images, while values between 10 and 50 indicate good quality, and values above 50 suggest average or poor quality [49].

3.6 Limitations of Accuracy in Evaluating GANs

Accuracy is not suitable for evaluating GANs because accuracy is an indicator of classification tasks, which is used to measure the prediction accuracy of the model in classification tasks, and cannot measure the quality and diversity of generated data. In generation tasks, there is no clear "correct answer" and the generated data has no "real label", so it is impossible to directly compare the correspondence between the generated data and a real sample.

The following is a result for a standard GANs model with high accuracy but generate low quality images.

```
1900 [D loss: 9.02800047697383e-06 | D accuracy: 100.0] [G loss: 0.0005307616665959358] [FID: -1.49196970922936e+87]
2/2 [=====] - 0s 28ms/step
1901 [D loss: 7.511995590903098e-06 | D accuracy: 100.0] [G loss: 0.0006589822005480528] [Epoch time: 0.51 seconds]
2/2 [=====] - 0s 29ms/step
1902 [D loss: 9.099765065911924e-06 | D accuracy: 100.0] [G loss: 0.0008528590551577508] [Epoch time: 0.48 seconds]
2/2 [=====] - 0s 29ms/step
1903 [D loss: 7.715634183114162e-06 | D accuracy: 100.0] [G loss: 0.0008156942203640938] [Epoch time: 0.48 seconds]
2/2 [=====] - 0s 29ms/step
1904 [D loss: 7.386817742371932e-06 | D accuracy: 100.0] [G loss: 0.0005978870904073119] [Epoch time: 0.49 seconds]
2/2 [=====] - 0s 29ms/step
1905 [D loss: 1.5066923879203387e-05 | D accuracy: 100.0] [G loss: 0.0014342099893838167] [Epoch time: 0.48 seconds]
2/2 [=====] - 0s 29ms/step
1906 [D loss: 1.9851730939990375e-05 | D accuracy: 100.0] [G loss: 0.0007973920437507331] [Epoch time: 0.49 seconds]
2/2 [=====] - 0s 29ms/step
1907 [D loss: 1.692915657258709e-05 | D accuracy: 100.0] [G loss: 0.0009673223830759525] [Epoch time: 0.49 seconds]
2/2 [=====] - 0s 29ms/step
1908 [D loss: 1.5433080079674255e-05 | D accuracy: 100.0] [G loss: 0.0010331417433917522] [Epoch time: 0.49 seconds]
2/2 [=====] - 0s 29ms/step
1909 [D loss: 3.74487547105673e-06 | D accuracy: 100.0] [G loss: 0.0008501751581206918] [Epoch time: 0.52 seconds]
```

Figure 3.4: GAN Training Accuracy over Epochs


```
▶ generate_images(generator, 10, 100)
⇒ 1/1 [=====] - 1s 844ms/step

```

Figure 3.5: Generated Images from GAN

FID (Fréchet Inception Distance) quantifies the difference between generated data and real data by comparing their distribution in the Inception network feature space. FID takes into account the overall distribution of generated data and real data, and can reflect the quality and diversity of generated data. A low FID value indicates that the distribution of generated data is very close to the distribution of real data, that is, the generated data is both realistic and diverse, so FID is a more suitable indicator for evaluating the performance of generative models.

Chapter 4

Results and Discussion

This chapter details the numerical processes undertaken to evaluate the performance of various GAN architectures. It includes model selection, structural adjustments, exploration of layer depth, and the impact of data augmentation on model performance. Additionally, the chapter discusses the application of the chosen model to a new dataset, providing insights into the practical effectiveness of the models in generating high-quality images.

4.1 Standard GAN Versus Other GAN Realizations

Approximately a decade has passed since Goodfellow introduced Generative Adversarial Networks (GANs), during which numerous variants of GAN models have been developed. For the purpose of this study, I have selected seven distinct GAN models for examination and minimal implementation: Standard GANs, Conditional GANs, Auxiliary Classifier GANs, Cycle GANs, Domain Transfer Network GANs, Coupled GANs, and Style GANs. Based on their foundational nature, extensive research and documentation, training and implementation efficiency, and flexibility and versatility, Standard GANs were selected for this study. The following outlines the reasons for this choice.

1. Foundational Nature: Standard GANs, introduced by Ian Goodfel-

low et al. in 2014, serve as the foundational model for all subsequent GAN variants. Understanding the principles and mechanics of Standard GANs is crucial for comprehending more complex versions like Conditional GANs, Cycle GANs, and Style GANs. By focusing on the Standard GAN, this study lays a solid foundation for exploring more advanced models.

2. Widely Studied and Well-Documented: Standard GANs have been extensively researched, with a vast amount of literature available. This wealth of resources provides a robust theoretical background and a variety of implementation strategies, facilitating a more thorough and well-supported analysis. This also means that there is ample precedent for common challenges and solutions, making it easier to troubleshoot and refine the model during the study.
3. Training and Implementation Efficiency: Compared to more complex GAN variants, Standard GANs typically require less computational power and shorter training times, making them more accessible for experimentation and analysis. This efficiency allows for multiple experiments and parameter tuning within the constraints of the study, leading to more reliable and reproducible results.
4. Flexibility and Versatility: Standard GANs are highly versatile and can be adapted to a wide range of tasks and datasets. This flexibility makes them an excellent choice for a detailed study that may involve exploring various applications or extending the model to new domains.

4.2 GAN With Convolutional or Dense layers

On the MNIST dataset for 3000 epochs, I compared the generated images from two GAN architectures: one using dense layers and the other using convolutional layers (CNN). It was observed that the GAN with the CNN

architecture outperformed the dense layer architecture in terms of image quality, producing clearer and more realistic images.

As shown in Figures 4.1 and 4.2, the generator and discriminator architectures for both dense and convolutional layers are compared. Figure 4.1a illustrates the generator with dense layers, which consists of a series of fully connected layers, while Figure 4.1b shows the generator with convolutional layers, where convolutional and transpose convolutional layers are used to capture spatial features more effectively. The same comparison applies to the discriminator architectures shown in Figure 4.2. The dense-layer discriminator (Figure 4.2a) is built with fully connected layers, whereas the convolutional discriminator (Figure 4.2b) leverages convolutional layers to better identify patterns in the data.

Layer (type)	Output Shape	Param #
dense_30 (Dense)	(None, 256)	25856
leaky_re_lu_22 (LeakyReLU)	(None, 256)	0
batch_normalization_18 (Batch Normalization)	(None, 256)	1024
dense_31 (Dense)	(None, 512)	131584
leaky_re_lu_23 (LeakyReLU)	(None, 512)	0
batch_normalization_19 (Batch Normalization)	(None, 512)	2048
dense_32 (Dense)	(None, 1024)	525312
leaky_re_lu_24 (LeakyReLU)	(None, 1024)	0
batch_normalization_20 (Batch Normalization)	(None, 1024)	4096
dense_33 (Dense)	(None, 784)	803600
reshape_6 (Reshape)	(None, 28, 28, 1)	0
<hr/>		
Total params:	1493520 (5.78 MB)	
Trainable params:	1489936 (5.68 MB)	
Non-trainable params:	3584 (14.00 KB)	

(a) Generator with Dense Layer

Layer (type)	Output Shape	Param #
dense_38 (Dense)	(None, 6272)	633472
leaky_re_lu_28 (LeakyReLU)	(None, 6272)	0
reshape_8 (Reshape)	(None, 7, 7, 128)	0
batch_normalization_24 (Batch Normalization)	(None, 7, 7, 128)	512
conv2d_transpose (Conv2DTranspose)	(None, 14, 14, 128)	262272
leaky_re_lu_29 (LeakyReLU)	(None, 14, 14, 128)	0
batch_normalization_25 (Batch Normalization)	(None, 14, 14, 128)	512
conv2d_transpose_1 (Conv2DTranspose)	(None, 28, 28, 64)	131136
leaky_re_lu_30 (LeakyReLU)	(None, 28, 28, 64)	0
batch_normalization_26 (Batch Normalization)	(None, 28, 28, 64)	256
conv2d (Conv2D)	(None, 28, 28, 1)	3137
<hr/>		
Total params:	1031297 (3.93 MB)	
Trainable params:	1030657 (3.93 MB)	
Non-trainable params:	640 (2.50 KB)	

(b) Generator with Convolution Layer

Figure 4.1: Generator Architecture with Dense and Convolutional Layers

Layer (type)	Output Shape	Param #
flatten_1 (Flatten)	(None, 784)	0
dense_23 (Dense)	(None, 512)	401920
leaky_re_lu_17 (LeakyReLU)	(None, 512)	0
dropout_2 (Dropout)	(None, 512)	0
dense_24 (Dense)	(None, 256)	131328
leaky_re_lu_18 (LeakyReLU)	(None, 256)	0
dropout_3 (Dropout)	(None, 256)	0
dense_25 (Dense)	(None, 1)	257
Total params:	533505 (2.04 MB)	
Trainable params:	533505 (2.04 MB)	
Non-trainable params:	0 (0.00 Byte)	

Layer (type)	Output Shape	Param #
conv2d_2 (Conv2D)	(None, 14, 14, 64)	640
leaky_re_lu_39 (LeakyReLU)	(None, 14, 14, 64)	0
dropout_6 (Dropout)	(None, 14, 14, 64)	0
conv2d_3 (Conv2D)	(None, 7, 7, 128)	73856
leaky_re_lu_40 (LeakyReLU)	(None, 7, 7, 128)	0
dropout_7 (Dropout)	(None, 7, 7, 128)	0
flatten_3 (Flatten)	(None, 6272)	0
dense_47 (Dense)	(None, 1)	6273

(a) Discriminator with Dense Layer (b) Discriminator with Convolution Layer

Figure 4.2: Discriminator Architecture with Dense and Convolutional Layers

Notably, the dense-layer models contain significantly more parameters compared to the convolutional models. The generator with dense layers has approximately 1.5 million parameters, whereas the convolutional generator has only about 1 million parameters. Similarly, the dense-layer discriminator has over 500,000 parameters, while the convolutional discriminator has only 80,000. Despite the larger parameter size, the dense-layer architectures were less effective in producing high-quality images, highlighting the advantage of convolutional layers in capturing spatial dependencies in the data.

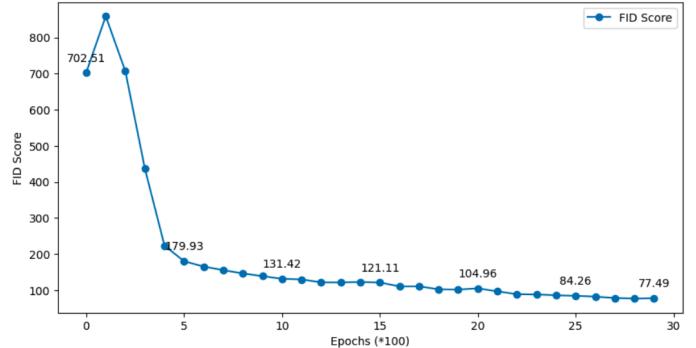


(a) Images Generated by GAN with Dense Layers

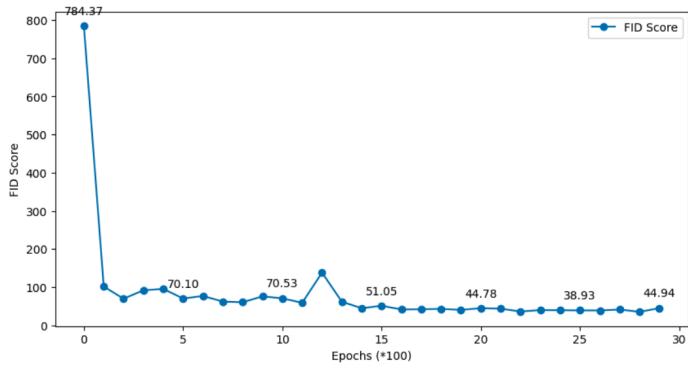


(b) Images Generated by GAN with Convolution Layers

Figure 4.3: Comparison of GAN Performance at 3000 Epochs



(a) FID Score for Dense Layer



(b) FID Score for Convolutional Layer

Figure 4.4: FID Scores Across 3000 Epochs

Based on the results, the GAN model with convolutional layers demonstrated superior performance in generating high-quality images compared to the dense layer model. As shown in the generated images (Figure 4.3), the convolutional GAN produced more coherent and realistic outputs, whereas the dense layer GAN struggled to generate clear and recognizable digits. Additionally, the FID scores, which measure the dissimilarity between generated and real images, further validate this observation. The FID scores for the convolutional GAN (Figure 4.4b) are consistently lower across training epochs compared to the dense layer model (Figure 4.4a), indicating better quality in the generated images.

Given these results, I selected the standard GAN model with convolutional layers for further experiments, as it not only generated higher-quality

images but also used fewer parameters, making it a more efficient and effective choice.

4.3 Exploring Layer Depth

In this section, I explored the impact of changing convolutional layers in the generator and discriminator of a GAN.

Firstly, I trained the basic convolutional GAN model three times and recorded the FID scores. Secondly, based on the basic structure, I gradually added convolutional layers to the generator, from one to three layers, while keeping the discriminator structure fixed. I trained the model three times for each configuration and recorded the FID scores, as shown in Table 4.1.

Finally, using the base GAN with three additional convolutional layers in the generator, I gradually added convolutional layers to the discriminator, from one to three layers, while keeping the generator structure fixed. I trained the model three times for each configuration and recorded the FID scores.

Table 4.1: Average FID Scores for Different GAN Architectures (Lower is Better)

Layers in G \ Layers in D	3	4	5	6
3	68.03	72.45	78.56	85.67
4	78.95	70.89	74.23	79.45
5	99.17	85.36	37.09	65.54
6	194.66	71.54	55.21	35.78

Note: A FID value below 10 is considered to represent very high-quality generated images, while values between 10 and 50 indicate good quality, and values above 50 suggest average or poor quality.

Upon comparing, I found that increasing the number of layers in either the generator or the discriminator alone worsened the model's performance. This imbalance disrupts the dynamic equilibrium between the generator and discriminator in the GAN framework. However, simultaneously increasing the layers in both the generator and the discriminator improved the model's

performance, as demonstrated by the decreasing FID scores with higher layer counts. For example, Table 4.1 shows that increasing both the generator and discriminator layers to five results in a significantly lower FID score of 37.09 compared to the unbalanced configurations.

This finding aligns with the concept of maintaining a balance between the generator and discriminator during training, as highlighted in the literature [50]. The importance of this balance is crucial for the effective operation of GANs, ensuring stable learning and improved image quality [51].

4.4 Impact of Data Augmentation on Model Performance

The impact of various data augmentation techniques on model performance was evaluated by comparing the average FID scores across different augmentation configurations (Table 4.2). The results indicate that data augmentation generally led to poorer performance in terms of image quality compared to training without any augmentation.

Table 4.2: Average FID Scores for Different Data Augmentation Techniques

Data Augmentation Technique	Average FID Score
Rotation 10° + Shifting 0.1 + Flipping	103.49
Rotation 10° + Shifting 0.1	76.04
Shifting 0.1	70.74
Shifting 0.05	66.06
Without Data Augmentation	58.94

Note: A FID value below 10 is considered to represent very high-quality generated images, while values between 10 and 50 indicate good quality, and values above 50 suggest average or poor quality.

Among the techniques tested, applying a combination of rotation, shifting, and flipping resulted in the highest FID score (103.49), suggesting a significant negative impact on the generated image quality. Rotation and

shifting alone slightly improved the FID score to 76.04, while reducing the shift range to 0.1 and 0.05 further improved performance, yielding FID scores of 70.74 and 66.06, respectively. The best performance was observed with no data augmentation, achieving an FID score of 58.94.

These findings suggest that, while data augmentation is commonly employed to improve model generalization, it can potentially disrupt the data distribution and hinder the optimization process when not carefully selected or overused.

4.5 Applying the Model to the Animal Faces-HQ Dataset

In this section, the Animal Faces-HQ (AFHQ) dataset was used to train a standard GAN focused on generating realistic cat images. The dataset consists of 16,130 high-quality images with a resolution of 512x512 pixels. Due to the high resolution of the original images, I downsampled them to 128x128 pixels to avoid GPU memory issues during training.

The GAN was then trained for 4000 epochs. The following results showcase the images generated by the model’s generator. Figure 4.5 displays the cat faces created by the GAN.

4.5.1 Model Structure and Training

The GAN model consists of two main components: a generator and a discriminator. The generator takes a random noise vector and progressively transforms it into a full-resolution image, while the discriminator is trained to distinguish between real and generated images.

The generator architecture starts by expanding the noise vector (100 dimensions) into a 16x16x256 feature map. Several transpose convolutional layers are then used to upsample this feature map to 128x128 pixels. The discriminator processes these images using convolutional layers, which progressively downsample the input. The final output of the discriminator is a classification (real or fake).

Below is a brief summary of the generator and discriminator models:

```
def build_generator():
    model = Sequential()

    # Expand noise vector and reshape
    model.add(Dense(16*16*256, input_dim=100))
    model.add(LeakyReLU(alpha=0.2))
    model.add(Reshape((16, 16, 256)))
    model.add(BatchNormalization(momentum=0.8))

    # Upsample to 128x128
    model.add(Conv2DTranspose(256, kernel_size=4, strides=2, padding='same'))
    model.add(LeakyReLU(alpha=0.2))
    model.add(BatchNormalization(momentum=0.8))

    model.add(Conv2DTranspose(128, kernel_size=4, strides=2, padding='same'))
    model.add(LeakyReLU(alpha=0.2))
    model.add(BatchNormalization(momentum=0.8))

    model.add(Conv2DTranspose(64, kernel_size=4, strides=2, padding='same'))
    model.add(LeakyReLU(alpha=0.2))
    model.add(BatchNormalization(momentum=0.8))

    model.add(Conv2D(3, kernel_size=7, activation='tanh', padding='same'))

    return model
```

The discriminator processes the generated images through a series of convolutional layers and finally classifies them as real or fake.

```
def build_discriminator():
    model = Sequential()
```

```

model.add(Conv2D(64, kernel_size=3, strides=2,
                input_shape=(128, 128, 3), padding='same'))
model.add(LeakyReLU(alpha=0.2))
model.add(Dropout(0.25))

model.add(Conv2D(128, kernel_size=3, strides=2, padding='same'))
model.add(LeakyReLU(alpha=0.2))
model.add(Dropout(0.25))

model.add(Conv2D(256, kernel_size=3, strides=2, padding='same'))
model.add(LeakyReLU(alpha=0.2))
model.add(Dropout(0.25))

model.add(Conv2D(512, kernel_size=3, strides=2, padding='same'))
model.add(LeakyReLU(alpha=0.2))
model.add(Dropout(0.25))

model.add(Flatten())
model.add(Dense(1, activation='sigmoid'))

return model

```

The training was conducted over 4000 epochs with a batch size of 128. Both real and generated images were used to train the discriminator, while the generator was trained to produce images that could fool the discriminator into classifying them as real.

```

def train(generator, discriminator, gan, x_train, epochs, batch_size=128):
    valid = np.ones((batch_size, 1))
    fake = np.zeros((batch_size, 1))

    for epoch in range(epochs):
        start_time = time.time()

```

```

# Train the discriminator
idx = np.random.randint(0, x_train.shape[0], batch_size)
imgs = x_train[idx]

noise = np.random.normal(0, 1, (batch_size, 100))
gen_imgs = generator.predict(noise)

d_loss_real = discriminator.train_on_batch(imgs, valid)
d_loss_fake = discriminator.train_on_batch(gen_imgs, fake)
d_loss = 0.5 * np.add(d_loss_real, d_loss_fake)

# Train the generator
noise = np.random.normal(0, 1, (batch_size, 100))
g_loss = gan.train_on_batch(noise, valid)

end_time = time.time()
epoch_time = end_time - start_time

print(f"{{epoch}} [D loss: {d_loss[0]} | D accuracy: "
      f"{{100 * d_loss[1]}}] [G loss: {g_loss}] "
      f"[Epoch time: {{epoch_time:.2f}} seconds]")

```

4.5.2 Generated Results

Figure 4.5 shows examples of the cat images generated by the GAN after training for 4000 epochs. While the model was able to capture many of the essential features of cat faces, the generated outputs still show some variation compared to the original dataset.



Figure 4.5: Cat Faces Generated by GAN

The GAN's architecture and training details are included in the appendix. Additionally, the full Python code for this implementation is available to provide a detailed understanding of how the model was trained and evaluated.

4.6 Discussion

In this study, I conducted a series of experiments to evaluate the performance of different GAN architectures and explored various factors influencing the quality of generated images. Specifically, the experiments compared the effects of using dense layers versus convolutional layers in a standard GAN, examined the impact of varying the number of convolutional layers in both the generator and discriminator, analyzed the effects of data augmentation on GAN training, and applied a GAN model to a dataset to generate high-quality cat face images.

The experiments revealed that GAN models utilizing convolutional layers outperformed those using dense layers in terms of image quality, confirming the strength of convolutional layers in capturing spatial features. Additionally, I found that increasing the number of layers in the generator or discriminator independently led to a decline in performance, suggesting that the balance between these two components is critical for stable GAN training. Simultaneously increasing the number of layers in both components improved the model's output.

Regarding data augmentation, the results indicated that certain augmentation techniques negatively impacted GAN training, likely by introducing noise or disrupting the data distribution. This suggests that when applying data augmentation to GANs, careful selection and parameter tuning are necessary to avoid adverse effects on training stability and performance.

When applying the model to the Animal Faces-HQ dataset, I successfully trained the GAN to generate cat images. However, due to hardware constraints, I had to downscale the images from 512x512 to 128x128 pixels, which may have limited the generated images' resolution and quality.

The limitations of this study include computational restrictions, which prevented experiments on higher-resolution images, and a focus on standard GANs rather than more advanced architectures such as StyleGAN or BigGAN. Future work could explore these more complex GAN architectures and improve data augmentation strategies to enhance training performance and image quality. Additionally, training GANs on higher-resolution datasets

and experimenting with advanced GAN models, such as Conditional GANs or Adaptive GANs, could provide further insights into enhancing image diversity and fidelity.

In conclusion, this study demonstrates the effectiveness of GAN models in generating high-quality images, while highlighting areas for future improvements in architecture and training methods.

Appendix A

Appendix

Code

Listing A.1: GAN model with dense layers

```
1 import tensorflow as tf
2 from tensorflow.keras.layers import Dense, Reshape, Flatten,
3     Dropout, LeakyReLU, BatchNormalization
4 from tensorflow.keras.models import Sequential
5 from tensorflow.keras.optimizers import Adam
6 from scipy.linalg import sqrtm
7 import numpy as np
8 import time
9 import matplotlib.pyplot as plt
10
11 np.random.seed(1000)
12 tf.random.set_seed(1000)
13
14 # input 100
15 # output 28*28*1
16
17 def build_generator():
18     model = Sequential()
19
20     # increase the dimension
21     model.add(Dense(256, input_dim=100))
22     model.add(LeakyReLU(alpha=0.2))
```

```

22     model.add(BatchNormalization(momentum=0.8))

23

24     model.add(Dense(512))
25     model.add(LeakyReLU(alpha=0.2))
26     model.add(BatchNormalization(momentum=0.8))

27

28     model.add(Dense(1024))
29     model.add(LeakyReLU(alpha=0.2))
30     model.add(BatchNormalization(momentum=0.8))

31

32     model.add(Dense(28*28, activation='tanh'))
33     model.add(Reshape((28, 28, 1)))

34

35     return model

36

37 # input 28*28*1
38 # output 1

39

40 def build_discriminator():
41     model = Sequential()

42

43     model.add(Flatten(input_shape=(28, 28, 1)))

44

45     model.add(Dense(512))
46     model.add(LeakyReLU(alpha=0.2))
47     model.add(Dropout(0.25))

48

49     model.add(Dense(256))
50     model.add(LeakyReLU(alpha=0.2))
51     model.add(Dropout(0.25))

52

53     model.add(Dense(1, activation='sigmoid'))

54

55     return model

56

57 # Load and preprocess the MNIST dataset
58 (x_train, _), (_, _) = tf.keras.datasets.mnist.load_data()
59 x_train = (x_train - 127.5) / 127.5
60 x_train = np.expand_dims(x_train, axis=3)

```

```

61
62 # Calculate FID function
63 def calculate_fid(real_images, fake_images):
64     act1 = real_images.reshape((real_images.shape[0], -1))
65     mu1, sigma1 = act1.mean(axis=0), np.cov(act1, rowvar=
66                               False)
67
68     act2 = fake_images.reshape((fake_images.shape[0], -1))
69     mu2, sigma2 = act2.mean(axis=0), np.cov(act2, rowvar=
70                               False)
71
72     ssdiff = np.sum((mu1 - mu2)**2.0)
73     covmean = sqrtm(sigma1.dot(sigma2))
74
75     if np.iscomplexobj(covmean):
76         covmean = covmean.real
77
78     fid = ssdiff + np.trace(sigma1 + sigma2 - 2.0 * covmean)
79
80     return fid
81
82
83 def train_gan(epochs=1000, batch_size=64, p_epoch=100):
84     generator = build_generator()
85     discriminator = build_discriminator()
86
87     discriminator.compile(loss='binary_crossentropy',
88                           optimizer=Adam(0.0002, 0.5), metrics=['accuracy'])
89     discriminator.trainable = False
90
91     gan_input = tf.keras.Input(shape=(100,))
92     gan_output = discriminator(generator(gan_input))
93     gan = tf.keras.Model(gan_input, gan_output)
94     gan.compile(loss='binary_crossentropy', optimizer=Adam
95                  (0.0002, 0.5))
96
97     half_batch = int(batch_size / 2)
98
99     d_losses = []
100    g_losses = []

```

```

96     d_acc = []
97     fid_scores = [] # List to store FID scores
98
99     start_time = time.time() # Record the start time
100
101    for epoch in range(epochs):
102        # Select a random half batch of real images
103        idx = np.random.randint(0, x_train.shape[0],
104                                half_batch)
105        real_images = x_train[idx]
106
107        # Generate a half batch of new fake images
108        noise = np.random.normal(0, 1, (half_batch, 100))
109        fake_images = generator.predict(noise)
110
111        # Train the discriminator
112        real_labels = np.ones((half_batch, 1))
113        fake_labels = np.zeros((half_batch, 1))
114
115        d_loss_real = discriminator.train_on_batch(
116            real_images, real_labels)
117        d_loss_fake = discriminator.train_on_batch(
118            fake_images, fake_labels)
119        d_loss = 0.5 * np.add(d_loss_real, d_loss_fake)
120
121
122        # Train the generator
123        noise = np.random.normal(0, 1, (batch_size, 100))
124        valid_y = np.ones((batch_size, 1))
125
126        g_loss = gan.train_on_batch(noise, valid_y)
127
128        # Record the losses
129        d_losses.append(d_loss[0])
130        g_losses.append(g_loss)
131        d_acc.append(d_loss[1] * 100)

# Calculate and print FID every p_epoch epochs
if epoch % p_epoch == 0:
    noise = np.random.normal(0, 1, (1000, 100))

```

```

132         fake_images = generator.predict(noise)
133         fid = calculate_fid(x_train[:1000], fake_images)
134         fid_scores.append(fid)
135         print(f"Epoch {epoch} [D loss: {d_loss[0]}, acc.: {100 * d_loss[1]}%] [G loss: {g_loss}] [FID: {fid}]")
136
137     end_time = time.time() # Record the end time
138     total_time = end_time - start_time
139     print(f"Total training time: {total_time:.2f} seconds")
140     return generator, d_losses, g_losses, d_acc, fid_scores

```

Listing A.2: GAN model with Convolutional layers

```

1 import tensorflow as tf
2 from tensorflow.keras.layers import Dense, Reshape,
3     Flatten, Dropout, LeakyReLU, Conv2D, Conv2DTranspose,
4     BatchNormalization
5 from tensorflow.keras.models import Sequential
6 from tensorflow.keras.optimizers import Adam
7 from scipy.linalg import sqrtm
8 import numpy as np
9 import time
10 import matplotlib.pyplot as plt
11
12 np.random.seed(1000)
13 tf.random.set_seed(1000)
14
15 # input 100
16 # output 28*28*1
17
18 def build_generator():
19     model = Sequential()
20
21     # increase the dimension
22     model.add(Dense(7*7*128, input_dim=100))
23     model.add(LeakyReLU(alpha=0.2))
24     model.add(Reshape((7, 7, 128)))
25     model.add(BatchNormalization(momentum=0.8))
26
27
28
29
30
31
32
33
34
35
36
37
38
39
40
41
42
43
44
45
46
47
48
49
50
51
52
53
54
55
56
57
58
59
60
61
62
63
64
65
66
67
68
69
70
71
72
73
74
75
76
77
78
79
80
81
82
83
84
85
86
87
88
89
90
91
92
93
94
95
96
97
98
99
100
101
102
103
104
105
106
107
108
109
110
111
112
113
114
115
116
117
118
119
120
121
122
123
124
125
126
127
128
129
130
131
132
133
134
135
136
137
138
139
140
141
142
143
144
145
146
147
148
149
150
151
152
153
154
155
156
157
158
159
160
161
162
163
164
165
166
167
168
169
170
171
172
173
174
175
176
177
178
179
180
181
182
183
184
185
186
187
188
189
190
191
192
193
194
195
196
197
198
199
200
201
202
203
204
205
206
207
208
209
210
211
212
213
214
215
216
217
218
219
220
221
222
223
224
225
226
227
228
229
230
231
232
233
234
235
236
237
238
239
240
241
242
243
244
245
246
247
248
249
250
251
252
253
254
255
256
257
258
259
260
261
262
263
264
265
266
267
268
269
270
271
272
273
274
275
276
277
278
279
280
281
282
283
284
285
286
287
288
289
290
291
292
293
294
295
296
297
298
299
300
301
302
303
304
305
306
307
308
309
310
311
312
313
314
315
316
317
318
319
320
321
322
323
324
325
326
327
328
329
330
331
332
333
334
335
336
337
338
339
340
341
342
343
344
345
346
347
348
349
350
351
352
353
354
355
356
357
358
359
360
361
362
363
364
365
366
367
368
369
370
371
372
373
374
375
376
377
378
379
380
381
382
383
384
385
386
387
388
389
390
391
392
393
394
395
396
397
398
399
400
401
402
403
404
405
406
407
408
409
410
411
412
413
414
415
416
417
418
419
420
421
422
423
424
425
426
427
428
429
430
431
432
433
434
435
436
437
438
439
440
441
442
443
444
445
446
447
448
449
450
451
452
453
454
455
456
457
458
459
460
461
462
463
464
465
466
467
468
469
470
471
472
473
474
475
476
477
478
479
480
481
482
483
484
485
486
487
488
489
490
491
492
493
494
495
496
497
498
499
500
501
502
503
504
505
506
507
508
509
510
511
512
513
514
515
516
517
518
519
520
521
522
523
524
525
526
527
528
529
530
531
532
533
534
535
536
537
538
539
540
541
542
543
544
545
546
547
548
549
550
551
552
553
554
555
556
557
558
559
560
561
562
563
564
565
566
567
568
569
570
571
572
573
574
575
576
577
578
579
580
581
582
583
584
585
586
587
588
589
589
590
591
592
593
594
595
596
597
598
599
599
600
601
602
603
604
605
606
607
608
609
609
610
611
612
613
614
615
616
617
618
619
619
620
621
622
623
624
625
626
627
628
629
629
630
631
632
633
634
635
636
637
637
638
639
639
640
641
642
643
644
645
645
646
647
647
648
648
649
649
650
651
651
652
652
653
653
654
654
655
655
656
656
657
657
658
658
659
659
660
660
661
661
662
662
663
663
664
664
665
665
666
666
667
667
668
668
669
669
670
670
671
671
672
672
673
673
674
674
675
675
676
676
677
677
678
678
679
679
680
680
681
681
682
682
683
683
684
684
685
685
686
686
687
687
688
688
689
689
690
690
691
691
692
692
693
693
694
694
695
695
696
696
697
697
698
698
699
699
700
700
701
701
702
702
703
703
704
704
705
705
706
706
707
707
708
708
709
709
710
710
711
711
712
712
713
713
714
714
715
715
716
716
717
717
718
718
719
719
720
720
721
721
722
722
723
723
724
724
725
725
726
726
727
727
728
728
729
729
730
730
731
731
732
732
733
733
734
734
735
735
736
736
737
737
738
738
739
739
740
740
741
741
742
742
743
743
744
744
745
745
746
746
747
747
748
748
749
749
750
750
751
751
752
752
753
753
754
754
755
755
756
756
757
757
758
758
759
759
760
760
761
761
762
762
763
763
764
764
765
765
766
766
767
767
768
768
769
769
770
770
771
771
772
772
773
773
774
774
775
775
776
776
777
777
778
778
779
779
780
780
781
781
782
782
783
783
784
784
785
785
786
786
787
787
788
788
789
789
790
790
791
791
792
792
793
793
794
794
795
795
796
796
797
797
798
798
799
799
800
800
801
801
802
802
803
803
804
804
805
805
806
806
807
807
808
808
809
809
810
810
811
811
812
812
813
813
814
814
815
815
816
816
817
817
818
818
819
819
820
820
821
821
822
822
823
823
824
824
825
825
826
826
827
827
828
828
829
829
830
830
831
831
832
832
833
833
834
834
835
835
836
836
837
837
838
838
839
839
840
840
841
841
842
842
843
843
844
844
845
845
846
846
847
847
848
848
849
849
850
850
851
851
852
852
853
853
854
854
855
855
856
856
857
857
858
858
859
859
860
860
861
861
862
862
863
863
864
864
865
865
866
866
867
867
868
868
869
869
870
870
871
871
872
872
873
873
874
874
875
875
876
876
877
877
878
878
879
879
880
880
881
881
882
882
883
883
884
884
885
885
886
886
887
887
888
888
889
889
890
890
891
891
892
892
893
893
894
894
895
895
896
896
897
897
898
898
899
899
900
900
901
901
902
902
903
903
904
904
905
905
906
906
907
907
908
908
909
909
910
910
911
911
912
912
913
913
914
914
915
915
916
916
917
917
918
918
919
919
920
920
921
921
922
922
923
923
924
924
925
925
926
926
927
927
928
928
929
929
930
930
931
931
932
932
933
933
934
934
935
935
936
936
937
937
938
938
939
939
940
940
941
941
942
942
943
943
944
944
945
945
946
946
947
947
948
948
949
949
950
950
951
951
952
952
953
953
954
954
955
955
956
956
957
957
958
958
959
959
960
960
961
961
962
962
963
963
964
964
965
965
966
966
967
967
968
968
969
969
970
970
971
971
972
972
973
973
974
974
975
975
976
976
977
977
978
978
979
979
980
980
981
981
982
982
983
983
984
984
985
985
986
986
987
987
988
988
989
989
990
990
991
991
992
992
993
993
994
994
995
995
996
996
997
997
998
998
999
999
1000
1000
1001
1001
1002
1002
1003
1003
1004
1004
1005
1005
1006
1006
1007
1007
1008
1008
1009
1009
1010
1010
1011
1011
1012
1012
1013
1013
1014
1014
1015
1015
1016
1016
1017
1017
1018
1018
1019
1019
1020
1020
1021
1021
1022
1022
1023
1023
1024
1024
1025
1025
1026
1026
1027
1027
1028
1028
1029
1029
1030
1030
1031
1031
1032
1032
1033
1033
1034
1034
1035
1035
1036
1036
1037
1037
1038
1038
1039
1039
1040
1040
1041
1041
1042
1042
1043
1043
1044
1044
1045
1045
1046
1046
1047
1047
1048
1048
1049
1049
1050
1050
1051
1051
1052
1052
1053
1053
1054
1054
1055
1055
1056
1056
1057
1057
1058
1058
1059
1059
1060
1060
1061
1061
1062
1062
1063
1063
1064
1064
1065
1065
1066
1066
1067
1067
1068
1068
1069
1069
1070
1070
1071
1071
1072
1072
1073
1073
1074
1074
1075
1075
1076
1076
1077
1077
1078
1078
1079
1079
1080
1080
1081
1081
1082
1082
1083
1083
1084
1084
1085
1085
1086
1086
1087
1087
1088
1088
1089
1089
1090
1090
1091
1091
1092
1092
1093
1093
1094
1094
1095
1095
1096
1096
1097
1097
1098
1098
1099
1099
1100
1100
1101
1101
1102
1102
1103
1103
1104
1104
1105
1105
1106
1106
1107
1107
1108
1108
1109
1109
1110
1110
1111
1111
1112
1112
1113
1113
1114
1114
1115
1115
1116
1116
1117
1117
1118
1118
1119
1119
1120
1120
1121
1121
1122
1122
1123
1123
1124
1124
1125
1125
1126
1126
1127
1127
1128
1128
1129
1129
1130
1130
1131
1131
1132
1132
1133
1133
1134
1134
1135
1135
1136
1136
1137
1137
1138
1138
1139
1139
1140
1140
1141
1141
1142
1142
1143
1143
1144
1144
1145
1145
1146
1146
1147
1147
1148
1148
1149
1149
1150
1150
1151
1151
1152
1152
1153
1153
1154
1154
1155
1155
1156
1156
1157
1157
1158
1158
1159
1159
1160
1160
1161
1161
1162
1162
1163
1163
1164
1164
1165
1165
1166
1166
1167
1167
1168
1168
1169
1169
1170
1170
1171
1171
1172
1172
1173
1173
1174
1174
1175
1175
1176
1176
1177
1177
1178
1178
1179
1179
1180
1180
1181
1181
1182
1182
1183
1183
1184
1184
1185
1185
1186
1186
1187
1187
1188
1188
1189
1189
1190
1190
1191
1191
1192
1192
1193
1193
1194
1194
1195
1195
1196
1196
1197
1197
1198
1198
1199
1199
1200
1200
1201
1201
1202
1202
1203
1203
1204
1204
1205
1205
1206
1206
1207
1207
1208
1208
1209
1209
1210
1210
1211
1211
1212
1212
1213
1213
1214
1214
1215
1215
1216
1216
1217
1217
1218
1218
1219
1219
1220
1220
1221
1221
1222
1222
1223
1223
1224
1224
1225
1225
1226
1226
1227
1227
1228
1228
1229
1229
1230
1230
1231
1231
1232
1232
1233
1233
1234
1234
1235
1235
1236
1236
1237
1237
1238
1238
1239
1239
1240
1240
1241
1241
1242
1242
1243
1243
1244
1244
1245
1245
1246
1246
1247
1247
1248
1248
1249
1249
1250
1250
1251
1251
1252
1252
1253
1253
1254
1254
1255
1255
1256
1256
1257
1257
1258
1258
1259
1259
1260
1260
1261
1261
1262
1262
1263
1263
1264
1264
1265
1265
1266
1266
1267
1267
1268
1268
1269
1269
1270
1270
1271
1271
1272
1272
1273
1273
1274
1274
1275
1275
1276
1276
1277
1277
1278
1278
1279
1279
1280
1280
1281
1281
1282
1282
1283
1283
1284
1284
1285
1285
1286
1286
1287
1287
1288
1288
1289
1289
1290
1290
1291
1291
1292
1292
1293
1293
1294
1294
1295
1295
1296
1296
1297
1297
1298
1298
1299
1299
1300
1300
1301
1301
1302
1302
1303
1303
1304
1304
1305
1305
1306
1306
1307
1307
1308
1308
1309
1309
1310
1310
1311
1311
1312
1312
1313
1313
1314
1314
1315
1315
1316
1316
1317
1317
1318
1318
1319
1319
1320
1320
1321
1321
1322
1322
1323
1323
1324
1324
1325
1325
1326
1326
1327
1327
1328
1328
1329
1329
1330
1330
1331
1331
1332
1332
1333
1333
1334
1334
1335
1335
1336
1336
1337
1337
1338
1338
1339
1339
1340
1340
1341
1341
1342
1342
1343
1343
1344
1344
1345
1345
1346
1346
1347
1347
1348
1348
1349
1349
1350
1350
1351
1351
1352
1352
1353
1353
1354
1354
1355
1355
1356
1356
1357
1357
1358
1358
1359
1359
1360
1360
1361
1361
1362
1362
1363
1363
1364
1364
1365
1365
1366
1366
1367
1367
1368
1368
1369
1369
1370
1370
1371
1371
1372
1372
1373
1373
1374
1374
1375
1375
1376
1376
1377
1377
1378
1378
1379
1379
1380
1380
1381
1381
1382
1382
1383
1383
1384
1384
1385
1385
1386
1386
1387
1387
1388
1388
1389
1389
1390
1390
1391
1391
1392
1392
1393
1393
1394
1394
1395
1395
1396
1396
1397
1397
1398
1398
1399
1399
1400
1400
1401
1401
1402
1402
1403
1403
1404
1404
1405
1405
1406
1406
1407
1407
1408
1408
1409
1409
1410
1410
1411
1411
1412
1412
1413
1413
1414
1414
1415
1415
1416
1416
1417
1417
1418
1418
1419
1419
1420
1420
1421
1421
1422
1422
1423
1423
1424
1424
1425
1425
1426
1426
1427
1427
1428
1428
1429
1429
1430
1430
1431
1431
1432
1432
1433
1433
1434
1434
1435
1435
1436
1436
1437
1437
1438
1438
1439
1439
1440
1440
1441
1441
1442
1442
1443
1443
1444
1444
1445
1445
1446
1446
1447
1447
1448
1448
1449
1449
1450
1450
1451
1451
1452
1452
1453
1453
1454
1454
1455
1455
1456
1456
1457
1457
1458
1458
1459
1459
1460
1460
1461
1461
1462
1462
1463
1463
1464
1464
1465
1465
1466
1466
1467
1467
1468
1468
1469
1469
1470
1470
1471
1471
1472
1472
1473
1473
1474
1474
1475
1475
1476
1476
1477
1477
1478
1478
1479
1479
1480
1480
1481
1481
1482
1482
1483
1483
1484
1484
1485
1485
1486
1486
1487
1487
1488
1488
1489
1489
1490
1490
1491
1491
1492
1492
1493
1493
1494
1494
1495
1495
1496
1496
1497
1497
1498
1498
1499
1499
1500
1500
1501
1501
1502
1502
1503
1503
1504
1504
1505
1505
1506
1506
1507
1507
1508
1508
1509
1509
1510
1510
1511
1511
1512
1512
1513
1513
1514
1514
1515
1515
1516
1516
1517
1517
1518
1518
1519
1519
1520
1520
1521
1521
1522
1522
1523
1523
1524
1524
1525
1525
1526
1526
1527
1527
1528
1528
1529
1529
1530
1530
1531
1531
1532
1532
1533
1533
1534
1534
1535
1535
1536
1536
1537
1537
1538
1538
1539
1539
1540
1540
1541

```

```

25     model.add(Conv2DTranspose(128, kernel_size=4, strides
26         =2, padding='same'))
27     model.add(LeakyReLU(alpha=0.2))
28     model.add(BatchNormalization(momentum=0.8))

29     model.add(Conv2DTranspose(64, kernel_size=4, strides
30         =2, padding='same'))
31     model.add(LeakyReLU(alpha=0.2))
32     model.add(BatchNormalization(momentum=0.8))

33     model.add(Conv2D(1, kernel_size=7, activation='tanh',
34         padding='same'))

35     return model

36

37 # input 28*28*1
38 # output 1

39

40 def build_discriminator():
41     model = Sequential()

42

43     model.add(Conv2D(64, kernel_size=3, strides=2,
44         input_shape=(28, 28, 1), padding='same'))
45     model.add(LeakyReLU(alpha=0.2))
46     model.add(Dropout(0.25))

47     model.add(Conv2D(128, kernel_size=3, strides=2,
48         padding='same'))
49     model.add(LeakyReLU(alpha=0.2))
50     model.add(Dropout(0.25))

51     model.add(Flatten())
52     model.add(Dense(1, activation='sigmoid'))

53

54     return model

55

56 # Load and preprocess the MNIST dataset
57 (x_train, _), (_, _) = tf.keras.datasets.mnist.load_data()

```

```

58     x_train = (x_train - 127.5) / 127.5
59     x_train = np.expand_dims(x_train, axis=3)
60
61     # Calculate FID function
62     def calculate_fid(real_images, fake_images):
63         act1 = real_images.reshape((real_images.shape[0], -1))
64         mu1, sigma1 = act1.mean(axis=0), np.cov(act1, rowvar=False)
65
66         act2 = fake_images.reshape((fake_images.shape[0], -1))
67         mu2, sigma2 = act2.mean(axis=0), np.cov(act2, rowvar=False)
68
69         ssdiff = np.sum((mu1 - mu2)**2.0)
70         covmean = sqrtm(sigma1.dot(sigma2))
71
72         if np.iscomplexobj(covmean):
73             covmean = covmean.real
74
75         fid = ssdiff + np.trace(sigma1 + sigma2 - 2.0 * covmean)
76
77     return fid
78
79     def train_gan(epochs=1000, batch_size=64, p_epoch=100):
80         generator = build_generator()
81         discriminator = build_discriminator()
82
83         discriminator.compile(loss='binary_crossentropy',
84                                optimizer=Adam(0.0002, 0.5), metrics=['accuracy'])
85         discriminator.trainable = False
86
87         gan_input = tf.keras.Input(shape=(100,))
88         gan_output = discriminator(generator(gan_input))
89         gan = tf.keras.Model(gan_input, gan_output)
90         gan.compile(loss='binary_crossentropy', optimizer=
91                     Adam(0.0002, 0.5))

```

```

90
91     # using half batches for the discriminator ensures
92     # balanced and efficient training,
93     # better memory management, and more stable training
94     # dynamics in GANs.
95     half_batch = int(batch_size / 2)
96
97
98     d_losses = []
99     g_losses = []
100    d_acc = []
101    fid_scores = [] # List to store FID scores
102
103    start_time = time.time() # Record the start time
104
105    for epoch in range(epochs):
106        # Select a random half batch of real images
107        idx = np.random.randint(0, x_train.shape[0],
108                               half_batch)
109        real_images = x_train[idx]
110
111        # Generate a half batch of new fake images
112        noise = np.random.normal(0, 1, (half_batch, 100))
113        fake_images = generator.predict(noise)
114
115        # Train the discriminator
116        real_labels = np.ones((half_batch, 1))
117        fake_labels = np.zeros((half_batch, 1))
118
119        d_loss_real = discriminator.train_on_batch(
120            real_images, real_labels)
121        d_loss_fake = discriminator.train_on_batch(
122            fake_images, fake_labels)
123        d_loss = 0.5 * np.add(d_loss_real, d_loss_fake)
124
125        # Train the generator
126        noise = np.random.normal(0, 1, (batch_size, 100))
127        valid_y = np.ones((batch_size, 1))
128
129        g_loss = gan.train_on_batch(noise, valid_y)

```

```

124
125     # Record the losses
126     d_losses.append(d_loss[0])
127     g_losses.append(g_loss)
128     d_acc.append(d_loss[1] * 100)

129
130     # Calculate FID every p_epoch epochs
131     if epoch % p_epoch == 0:
132         noise = np.random.normal(0, 1, (1000, 100))
133         fake_images = generator.predict(noise)
134         fid = calculate_fid(x_train[:1000],
135                               fake_images)
136         fid_scores.append(fid)
137         print(f"epoch:[D_loss:{d_loss[0]},acc.:"
138               f"{100*d_loss[1]}%][G_loss:{g_loss}]"
139               f" FID:{fid}]")
140
141     end_time = time.time() # Record the end time
142     total_time = end_time - start_time
143     print(f"Total training time:{total_time:.2f} seconds"
144          ")
145
146     return generator, d_losses, g_losses, d_acc,
147           fid_scores

```

Listing A.3: Explore data augmentation 1

```

1 import tensorflow as tf
2 from tensorflow.keras.layers import Dense, Reshape,
3     Flatten, Dropout, LeakyReLU, Conv2D, Conv2DTranspose,
4     BatchNormalization
5 from tensorflow.keras.models import Sequential
6 from tensorflow.keras.optimizers import Adam
7 from scipy.linalg import sqrtm
8 import numpy as np
9 import time
10 import matplotlib.pyplot as plt
11
12 np.random.seed(1000)
13 tf.random.set_seed(1000)

```

```

13 # input 100
14 # output 28*28*1
15
16 def build_generator():
17     model = Sequential()
18
19     # increase the dimension
20     model.add(Dense(7*7*128, input_dim=100))
21     model.add(LeakyReLU(alpha=0.2))
22     model.add(Reshape((7, 7, 128)))
23     model.add(BatchNormalization(momentum=0.8))
24
25     model.add(Conv2DTranspose(128, kernel_size=4, strides
26                             =2, padding='same'))
27     model.add(LeakyReLU(alpha=0.2))
28     model.add(BatchNormalization(momentum=0.8))
29
30     model.add(Conv2DTranspose(64, kernel_size=4, strides
31                             =2, padding='same'))
32     model.add(LeakyReLU(alpha=0.2))
33     model.add(BatchNormalization(momentum=0.8))
34
35     model.add(Conv2D(1, kernel_size=7, activation='tanh',
36                      padding='same'))
37
38     return model
39
40     # input 28*28*1
41     # output 1
42
43 def build_discriminator():
44     model = Sequential()
45
46     model.add(Conv2D(64, kernel_size=3, strides=2,
47                      input_shape=(28, 28, 1), padding='same'))
48     model.add(LeakyReLU(alpha=0.2))
49     model.add(Dropout(0.25))
50
51     model.add(Conv2D(128, kernel_size=3, strides=2,

```

```

        padding='same'))
48 model.add(LeakyReLU(alpha=0.2))
49 model.add(Dropout(0.25))

50
51 model.add(Flatten())
52 model.add(Dense(1, activation='sigmoid'))

53
54 return model

55
56 (x_train, _), (_, _) = tf.keras.datasets.mnist.load_data()
57 ()
58 x_train = (x_train - 127.5) / 127.5
59 x_train = np.expand_dims(x_train, axis=3)

60 datagen = tf.keras.preprocessing.image.ImageDataGenerator(
61 (
62     rotation_range=10,
63     width_shift_range=0.1,
64     height_shift_range=0.1,
65     horizontal_flip=True
66 )
67
68 def calculate_fid(real_images, fake_images):
69     act1 = real_images.reshape((real_images.shape[0], -1)
70         )
71     mu1, sigma1 = act1.mean(axis=0), np.cov(act1, rowvar=
72         False)
73
74     act2 = fake_images.reshape((fake_images.shape[0], -1)
75         )
76     mu2, sigma2 = act2.mean(axis=0), np.cov(act2, rowvar=
77         False)
78
79     ssdiff = np.sum((mu1 - mu2)**2.0)
    covmean = sqrtm(sigma1.dot(sigma2))

if np.iscomplexobj(covmean):
    covmean = covmean.real

```

```

80         fid = ssdiff + np.trace(sigma1 + sigma2 - 2.0 *
81                         covmean)
82
83     return fid
84
85
86     def train_gan(epochs=1000, batch_size=64, p_epoch=100):
87         generator = build_generator()
88         discriminator = build_discriminator()
89
90         discriminator.compile(loss='binary_crossentropy',
91                               optimizer=Adam(0.0002, 0.5), metrics=['accuracy'])
92         discriminator.trainable = False
93
94         gan_input = tf.keras.Input(shape=(100,))
95         gan_output = discriminator(generator(gan_input))
96         gan = tf.keras.Model(gan_input, gan_output)
97         gan.compile(loss='binary_crossentropy', optimizer=
98                     Adam(0.0002, 0.5))
99
100
101        half_batch = int(batch_size / 2)
102
103        d_losses = []
104        g_losses = []
105        d_acc = []
106        fid_scores = []
107
108        start_time = time.time() # Record the start time
109
110        for epoch in range(epochs):
111            # Select a random half batch of real images
112            idx = np.random.randint(0, x_train.shape[0],
113                                   half_batch)
114            real_images = x_train[idx]
115
116            real_images_augmented = next(datagen.flow(
117                real_images, batch_size=half_batch))
118
119            # Generate a half batch of new fake images
120            noise = np.random.normal(0, 1, (half_batch, 100))

```

```

114     fake_images = generator.predict(noise)
115
116     # Train the discriminator
117     real_labels = np.ones((half_batch, 1))
118     fake_labels = np.zeros((half_batch, 1))
119
120     d_loss_real = discriminator.train_on_batch(
121         real_images_augmented, real_labels)
122     d_loss_fake = discriminator.train_on_batch(
123         fake_images, fake_labels)
124     d_loss = 0.5 * np.add(d_loss_real, d_loss_fake)
125
126     # Train the generator
127     noise = np.random.normal(0, 1, (batch_size, 100))
128     valid_y = np.ones((batch_size, 1))
129
130     g_loss = gan.train_on_batch(noise, valid_y)
131
132     # Record the losses
133     d_losses.append(d_loss[0])
134     g_losses.append(g_loss)
135     d_acc.append(d_loss[1] * 100)
136
137     # Calculate FID every p_epoch epochs
138     if epoch % p_epoch == 0:
139         noise = np.random.normal(0, 1, (1000, 100))
140         fake_images = generator.predict(noise)
141         fid = calculate_fid(x_train[:1000],
142                             fake_images)
143         fid_scores.append(fid)
144         print(f"epoch:[D_loss:{d_loss[0]},acc.:"
145               f"{100*d_loss[1]}%][G_loss:{g_loss}]"
146               f" FID:{fid}]")
147
148     end_time = time.time() # Record the end time
149     total_time = end_time - start_time
150     print(f"Total training time:{total_time:.2f} seconds"
151          ")
152
153     return generator, d_losses, g_losses, d_acc,

```

```

    fid_scores

147
148 # Training the GAN with data augmentation and FID
149   calculation
generator, d_losses, g_losses, d_acc, fid_scores =
    train_gan(epochs=1000, batch_size=64, p_epoch=100)

```

Listing A.4: Explore data augmentation 2

```

1 import tensorflow as tf
2 from tensorflow.keras.layers import Dense, Reshape,
3   Flatten, Dropout, LeakyReLU, Conv2D, Conv2DTranspose,
4   BatchNormalization
5 from tensorflow.keras.models import Sequential
6 from tensorflow.keras.optimizers import Adam
7 from scipy.linalg import sqrtm
8 import numpy as np
9 import time
10 import matplotlib.pyplot as plt
11
12 np.random.seed(1000)
13 tf.random.set_seed(1000)
14
15 # input 100
16 # output 28*28*1
17
18 def build_generator():
19   model = Sequential()
20
21   # increase the dimension
22   model.add(Dense(7*7*128, input_dim=100))
23   model.add(LeakyReLU(alpha=0.2))
24   model.add(Reshape((7, 7, 128)))
25   model.add(BatchNormalization(momentum=0.8))
26
27   model.add(Conv2DTranspose(128, kernel_size=4, strides
28     =2, padding='same'))
29   model.add(LeakyReLU(alpha=0.2))
30   model.add(BatchNormalization(momentum=0.8))

```

```

29         model.add(Conv2DTranspose(64, kernel_size=4, strides
30                         =2, padding='same'))
31         model.add(LeakyReLU(alpha=0.2))
32         model.add(BatchNormalization(momentum=0.8))

33         model.add(Conv2D(1, kernel_size=7, activation='tanh',
34                         padding='same'))

35     return model

36

37 # input 28*28*1
38 # output 1

39

40 def build_discriminator():
41     model = Sequential()

42

43     model.add(Conv2D(64, kernel_size=3, strides=2,
44                     input_shape=(28, 28, 1), padding='same'))
45     model.add(LeakyReLU(alpha=0.2))
46     model.add(Dropout(0.25))

47     model.add(Conv2D(128, kernel_size=3, strides=2,
48                     padding='same'))
49     model.add(LeakyReLU(alpha=0.2))
50     model.add(Dropout(0.25))

51     model.add(Flatten())
52     model.add(Dense(1, activation='sigmoid'))

53

54     return model

55

56 (x_train, _), (_, _) = tf.keras.datasets.mnist.load_data()
57             ()
58 x_train = (x_train - 127.5) / 127.5
59 x_train = np.expand_dims(x_train, axis=3)

60 datagen = tf.keras.preprocessing.image.ImageDataGenerator(
61             (
62                 rotation_range=10,

```

```

62         width_shift_range=0.1,
63         height_shift_range=0.1,
64     )
65
66     def calculate_fid(real_images, fake_images):
67         act1 = real_images.reshape((real_images.shape[0], -1)
68             )
69         mu1, sigma1 = act1.mean(axis=0), np.cov(act1, rowvar=
70             False)
71
72         act2 = fake_images.reshape((fake_images.shape[0], -1)
73             )
74         mu2, sigma2 = act2.mean(axis=0), np.cov(act2, rowvar=
75             False)
76
77         ssdiff = np.sum((mu1 - mu2)**2.0)
78         covmean = sqrtm(sigma1.dot(sigma2))
79
80         if np.iscomplexobj(covmean):
81             covmean = covmean.real
82
83         fid = ssdiff + np.trace(sigma1 + sigma2 - 2.0 *
84             covmean)
85
86         return fid
87
88     def train_gan(epochs=1000, batch_size=64, p_epoch=100):
89         generator = build_generator()
90         discriminator = build_discriminator()
91
92         discriminator.compile(loss='binary_crossentropy',
93             optimizer=Adam(0.0002, 0.5), metrics=['accuracy'])
94         discriminator.trainable = False
95
96         gan_input = tf.keras.Input(shape=(100,))
97         gan_output = discriminator(generator(gan_input))
98         gan = tf.keras.Model(gan_input, gan_output)
99         gan.compile(loss='binary_crossentropy', optimizer=
100             Adam(0.0002, 0.5))

```

```

94
95     half_batch = int(batch_size / 2)
96
97     d_losses = []
98     g_losses = []
99     d_acc = []
100    fid_scores = []
101
102    start_time = time.time() # Record the start time
103
104    for epoch in range(epochs):
105        # Select a random half batch of real images
106        idx = np.random.randint(0, x_train.shape[0],
107                               half_batch)
108        real_images = x_train[idx]
109
110        real_images_augmented = next(datagen.flow(
111            real_images, batch_size=half_batch))
112
113        # Generate a half batch of new fake images
114        noise = np.random.normal(0, 1, (half_batch, 100))
115        fake_images = generator.predict(noise)
116
117        # Train the discriminator
118        real_labels = np.ones((half_batch, 1))
119        fake_labels = np.zeros((half_batch, 1))
120
121        d_loss_real = discriminator.train_on_batch(
122            real_images_augmented, real_labels)
123        d_loss_fake = discriminator.train_on_batch(
124            fake_images, fake_labels)
125        d_loss = 0.5 * np.add(d_loss_real, d_loss_fake)
126
127        # Train the generator
128        noise = np.random.normal(0, 1, (batch_size, 100))

```

```

129     # Record the losses
130     d_losses.append(d_loss[0])
131     g_losses.append(g_loss)
132     d_acc.append(d_loss[1] * 100)
133
134     # Calculate FID every p_epoch epochs
135     if epoch % p_epoch == 0:
136         noise = np.random.normal(0, 1, (1000, 100))
137         fake_images = generator.predict(noise)
138         fid = calculate_fid(x_train[:1000],
139                               fake_images)
140         fid_scores.append(fid)
141         print(f'{epoch}[D_loss:{d_loss[0]}, acc.:'
142               f'{100*d_loss[1]}%][G_loss:{g_loss}]'
143               f' FID:{fid}]')
144
145     end_time = time.time() # Record the end time
146     total_time = end_time - start_time
147     print(f'Total training time:{total_time:.2f} seconds')
148
149     return generator, d_losses, g_losses, d_acc,
150           fid_scores
151
152
153     # Training the GAN with data augmentation and FID
154     # calculation
155     generator, d_losses, g_losses, d_acc, fid_scores =
156         train_gan(epochs=1000, batch_size=64, p_epoch=100)

```

Listing A.5: Explore data augmentation 3

```

1 import tensorflow as tf
2 from tensorflow.keras.layers import Dense, Reshape,
3   Flatten, Dropout, LeakyReLU, Conv2D, Conv2DTranspose,
4   BatchNormalization
5 from tensorflow.keras.models import Sequential
6 from tensorflow.keras.optimizers import Adam
7 from scipy.linalg import sqrtm
8 import numpy as np
9 import time
10 import matplotlib.pyplot as plt

```

```

9
10    np.random.seed(1000)
11    tf.random.set_seed(1000)
12
13    # input 100
14    # output 28*28*1
15
16    def build_generator():
17        model = Sequential()
18
19        # increase the dimension
20        model.add(Dense(7*7*128, input_dim=100))
21        model.add(LeakyReLU(alpha=0.2))
22        model.add(Reshape((7, 7, 128)))
23        model.add(BatchNormalization(momentum=0.8))
24
25        model.add(Conv2DTranspose(128, kernel_size=4, strides
26            =2, padding='same'))
27        model.add(LeakyReLU(alpha=0.2))
28        model.add(BatchNormalization(momentum=0.8))
29
30        model.add(Conv2DTranspose(64, kernel_size=4, strides
31            =2, padding='same'))
32        model.add(LeakyReLU(alpha=0.2))
33        model.add(BatchNormalization(momentum=0.8))
34
35        model.add(Conv2D(1, kernel_size=7, activation='tanh',
36            padding='same'))
37
38        return model
39
40    # input 28*28*1
41    # output 1
42
43    def build_discriminator():
44        model = Sequential()
45
46        model.add(Conv2D(64, kernel_size=3, strides=2,
47            input_shape=(28, 28, 1), padding='same'))

```

```

44     model.add(LeakyReLU(alpha=0.2))
45     model.add(Dropout(0.25))

46
47     model.add(Conv2D(128, kernel_size=3, strides=2,
48                     padding='same'))
49     model.add(LeakyReLU(alpha=0.2))
50     model.add(Dropout(0.25))

51     model.add(Flatten())
52     model.add(Dense(1, activation='sigmoid'))

53
54     return model

55

56 (x_train, _), (_, _) = tf.keras.datasets.mnist.load_data
57 ()
58 x_train = (x_train - 127.5) / 127.5
59 x_train = np.expand_dims(x_train, axis=3)

60 datagen = tf.keras.preprocessing.image.ImageDataGenerator(
61 (
62     width_shift_range=0.1,
63     height_shift_range=0.1,
64 )

65 def calculate_fid(real_images, fake_images):
66     act1 = real_images.reshape((real_images.shape[0], -1)
67                               )
68     mu1, sigma1 = act1.mean(axis=0), np.cov(act1, rowvar=
69                                         False)
70
71     act2 = fake_images.reshape((fake_images.shape[0], -1)
72                               )
73     mu2, sigma2 = act2.mean(axis=0), np.cov(act2, rowvar=
74                                         False)
75
76     ssdiff = np.sum((mu1 - mu2)**2.0)
77     covmean = sqrtm(sigma1.dot(sigma2))
78
79     if np.iscomplexobj(covmean):

```

```

76         covmean = covmean.real
77
78     fid = ssdiff + np.trace(sigma1 + sigma2 - 2.0 *
79                             covmean)
80
81     return fid
82
83
84 def train_gan(epochs=1000, batch_size=64, p_epoch=100):
85     generator = build_generator()
86     discriminator = build_discriminator()
87
88     discriminator.compile(loss='binary_crossentropy',
89                           optimizer=Adam(0.0002, 0.5), metrics=['accuracy'])
90     discriminator.trainable = False
91
92     gan_input = tf.keras.Input(shape=(100,))
93     gan_output = discriminator(generator(gan_input))
94     gan = tf.keras.Model(gan_input, gan_output)
95     gan.compile(loss='binary_crossentropy', optimizer=
96                  Adam(0.0002, 0.5))
97
98     half_batch = int(batch_size / 2)
99
100    d_losses = []
101    g_losses = []
102    d_acc = []
103    fid_scores = []
104
105    start_time = time.time() # Record the start time
106
107    for epoch in range(epochs):
108        # Select a random half batch of real images
109        idx = np.random.randint(0, x_train.shape[0],
110                               half_batch)
111        real_images = x_train[idx]
112
113        real_images_augmented = next(datagen.flow(
114                                      real_images, batch_size=half_batch))

```

```

110     # Generate a half batch of new fake images
111     noise = np.random.normal(0, 1, (half_batch, 100))
112     fake_images = generator.predict(noise)
113
114     # Train the discriminator
115     real_labels = np.ones((half_batch, 1))
116     fake_labels = np.zeros((half_batch, 1))
117
118     d_loss_real = discriminator.train_on_batch(
119         real_images_augmented, real_labels)
120     d_loss_fake = discriminator.train_on_batch(
121         fake_images, fake_labels)
122     d_loss = 0.5 * np.add(d_loss_real, d_loss_fake)
123
124     # Train the generator
125     noise = np.random.normal(0, 1, (batch_size, 100))
126     valid_y = np.ones((batch_size, 1))
127
128     g_loss = gan.train_on_batch(noise, valid_y)
129
130     # Record the losses
131     d_losses.append(d_loss[0])
132     g_losses.append(g_loss)
133     d_acc.append(d_loss[1] * 100)
134
135     # Calculate FID every p_epoch epochs
136     if epoch % p_epoch == 0:
137         noise = np.random.normal(0, 1, (1000, 100))
138         fake_images = generator.predict(noise)
139         fid = calculate_fid(x_train[:1000],
140                             fake_images)
141         fid_scores.append(fid)
142         print(f"Epoch [{epoch}] [D loss: {d_loss[0]}] [G loss: {g_loss}] [FID: {fid}]")
143
144     end_time = time.time()    # Record the end time
145     total_time = end_time - start_time
146     print(f"Total training time: {total_time:.2f} seconds")

```

```

        ")
144     return generator, d_losses, g_losses, d_acc,
           fid_scores
145
146     # Training the GAN with data augmentation and FID
147     # calculation
147     generator, d_losses, g_losses, d_acc, fid_scores =
           train_gan(epochs=1000, batch_size=64, p_epoch=100)

```

Listing A.6: Explore GAN with more convolutional layers 1

```

1 import tensorflow as tf
2 from tensorflow.keras.layers import Dense, Reshape,
3     Flatten, Dropout, LeakyReLU, Conv2D, Conv2DTranspose,
4     BatchNormalization
5 from tensorflow.keras.models import Sequential
6 from tensorflow.keras.optimizers import Adam
7 from scipy.linalg import sqrtm
8 import numpy as np
9 import time
10 import matplotlib.pyplot as plt
11
12 np.random.seed(1000)
13 tf.random.set_seed(1000)
14
15 # input 100
16 # output 28*28*1
17
18 def build_generator():
19     model = Sequential()
20
21     model.add(Dense(7*7*128, input_dim=100))
22     model.add(LeakyReLU(alpha=0.2))
23     model.add(Reshape((7, 7, 128)))
24     model.add(BatchNormalization(momentum=0.8))
25
26     # add 1 convolution layer
27     model.add(Conv2D(128, kernel_size=3, strides=1,
28                     padding='same'))
29     model.add(LeakyReLU(alpha=0.2))

```

```

27     model.add(BatchNormalization(momentum=0.8))

28
29     model.add(Conv2DTranspose(128, kernel_size=4, strides
30         =2, padding='same'))
31     model.add(LeakyReLU(alpha=0.2))
32     model.add(BatchNormalization(momentum=0.8))

33     model.add(Conv2DTranspose(64, kernel_size=4, strides
34         =2, padding='same'))
35     model.add(LeakyReLU(alpha=0.2))
36     model.add(BatchNormalization(momentum=0.8))

37     model.add(Conv2D(1, kernel_size=7, activation='tanh',
38         padding='same'))

39     return model

40
41 # input 28*28*1
42 # output 1

43
44 def build_discriminator():
45     model = Sequential()

46
47     model.add(Conv2D(64, kernel_size=3, strides=2,
48         input_shape=(28, 28, 1), padding='same'))
49     model.add(LeakyReLU(alpha=0.2))
50     model.add(Dropout(0.25))

51     model.add(Conv2D(128, kernel_size=3, strides=2,
52         padding='same'))
53     model.add(LeakyReLU(alpha=0.2))
54     model.add(Dropout(0.25))

55     model.add(Flatten())
56     model.add(Dense(1, activation='sigmoid'))

57
58     return model

59
60 (x_train, _), (_, _) = tf.keras.datasets.mnist.load_data

```

```

    ()

61   x_train = (x_train - 127.5) / 127.5
62   x_train = np.expand_dims(x_train, axis=3)

63
64   def calculate_fid(real_images, fake_images):
65       act1 = real_images.reshape((real_images.shape[0], -1))
66       mu1, sigma1 = act1.mean(axis=0), np.cov(act1, rowvar=False)

67
68       act2 = fake_images.reshape((fake_images.shape[0], -1))
69       mu2, sigma2 = act2.mean(axis=0), np.cov(act2, rowvar=False)

70
71       ssdiff = np.sum((mu1 - mu2)**2.0)
72       covmean = sqrtm(sigma1.dot(sigma2))

73
74       if np.iscomplexobj(covmean):
75           covmean = covmean.real

76
77       fid = ssdiff + np.trace(sigma1 + sigma2 - 2.0 * covmean)

78
79       return fid

80
81   def train_gan(epochs=1000, batch_size=64, p_epoch=100):
82       generator = build_generator()
83       discriminator = build_discriminator()

84
85       discriminator.compile(loss='binary_crossentropy',
86                               optimizer=Adam(0.0002, 0.5), metrics=['accuracy'])
86       discriminator.trainable = False

87
88       gan_input = tf.keras.Input(shape=(100,))
89       gan_output = discriminator(generator(gan_input))
90       gan = tf.keras.Model(gan_input, gan_output)
91       gan.compile(loss='binary_crossentropy', optimizer=
92                               Adam(0.0002, 0.5))

```

```

92
93     half_batch = int(batch_size / 2)
94
95     d_losses = []
96     g_losses = []
97     d_acc = []
98     fid_scores = []
99
100    start_time = time.time() # Record the start time
101
102    for epoch in range(epochs):
103        # Select a random half batch of real images
104        idx = np.random.randint(0, x_train.shape[0],
105                               half_batch)
106        real_images = x_train[idx]
107
108        # Generate a half batch of new fake images
109        noise = np.random.normal(0, 1, (half_batch, 100))
110        fake_images = generator.predict(noise)
111
112        # Train the discriminator
113        real_labels = np.ones((half_batch, 1))
114        fake_labels = np.zeros((half_batch, 1))
115
116        d_loss_real = discriminator.train_on_batch(
117            real_images, real_labels)
118        d_loss_fake = discriminator.train_on_batch(
119            fake_images, fake_labels)
120        d_loss = 0.5 * np.add(d_loss_real, d_loss_fake)
121
122        # Train the generator
123        noise = np.random.normal(0, 1, (batch_size, 100))
124        valid_y = np.ones((batch_size, 1))
125
126        g_loss = gan.train_on_batch(noise, valid_y)
127
128        # Record the losses
129        d_losses.append(d_loss[0])
130        g_losses.append(g_loss)

```

```

128     d_acc.append(d_loss[1] * 100)
129
130     # Calculate FID every p_epoch epochs
131     if epoch % p_epoch == 0:
132         noise = np.random.normal(0, 1, (1000, 100))
133         fake_images = generator.predict(noise)
134         fid = calculate_fid(x_train[:1000],
135                               fake_images)
136         fid_scores.append(fid)
137         print(f"epoch:{d_loss[0]},acc.:{100*d_loss[1]}%[Gloss:{g_loss}]"
138               f"FID:{fid}")
139
140         end_time = time.time()    # Record the end time
141         total_time = end_time - start_time
142         print(f"Total training time:{total_time:.2f} seconds")
143
144     return generator, d_losses, g_losses, d_acc,
145           fid_scores

```

Listing A.7: Explore GAN with more convolutional layers 2

```

1 import tensorflow as tf
2 from tensorflow.keras.layers import Dense, Reshape,
3     Flatten, Dropout, LeakyReLU, Conv2D, Conv2DTranspose,
4     BatchNormalization
5 from tensorflow.keras.models import Sequential
6 from tensorflow.keras.optimizers import Adam
7 from scipy.linalg import sqrtm
8 import numpy as np
9 import time
10 import matplotlib.pyplot as plt
11
12 np.random.seed(1000)
13 tf.random.set_seed(1000)
14
15 # input 100
16 # output 28*28*1
17
18 def build_generator():

```

```

17     model = Sequential()
18
19     model.add(Dense(7*7*128, input_dim=100))
20     model.add(LeakyReLU(alpha=0.2))
21     model.add(Reshape((7, 7, 128)))
22     model.add(BatchNormalization(momentum=0.8))
23
24     # add the first convolution layer
25     model.add(Conv2D(128, kernel_size=3, strides=1,
26                      padding='same'))
27     model.add(LeakyReLU(alpha=0.2))
28     model.add(BatchNormalization(momentum=0.8))
29
30     # add the 2nd convolution layer
31     model.add(Conv2D(128, kernel_size=3, strides=1,
32                      padding='same'))
33     model.add(LeakyReLU(alpha=0.2))
34     model.add(BatchNormalization(momentum=0.8))
35
36     model.add(Conv2DTranspose(128, kernel_size=4, strides
37                            =2, padding='same'))
38     model.add(LeakyReLU(alpha=0.2))
39     model.add(BatchNormalization(momentum=0.8))
40
41     model.add(Conv2DTranspose(64, kernel_size=4, strides
42                            =2, padding='same'))
43     model.add(LeakyReLU(alpha=0.2))
44     model.add(BatchNormalization(momentum=0.8))
45
46     model.add(Conv2D(1, kernel_size=7, activation='tanh',
47                      padding='same'))
48
49     return model
50
51
52     # input 28*28*1
53     # output 1
54
55     def build_discriminator():
56         model = Sequential()

```

```

51
52     model.add(Conv2D(64, kernel_size=3, strides=2,
53                     input_shape=(28, 28, 1), padding='same'))
54     model.add(LeakyReLU(alpha=0.2))
55     model.add(Dropout(0.25))

56
57     model.add(Conv2D(128, kernel_size=3, strides=2,
58                     padding='same'))
59     model.add(LeakyReLU(alpha=0.2))
60     model.add(Dropout(0.25))

61
62     model.add(Flatten())
63     model.add(Dense(1, activation='sigmoid'))

64
65     return model

66
67 (x_train, _), (_, _) = tf.keras.datasets.mnist.load_data()
68 ()
69
70 x_train = (x_train - 127.5) / 127.5
71 x_train = np.expand_dims(x_train, axis=3)

72
73 def calculate_fid(real_images, fake_images):
74     act1 = real_images.reshape((real_images.shape[0], -1))
75
76     mu1, sigma1 = act1.mean(axis=0), np.cov(act1, rowvar=False)

77
78     act2 = fake_images.reshape((fake_images.shape[0], -1))
79
80     mu2, sigma2 = act2.mean(axis=0), np.cov(act2, rowvar=False)

81
82     ssdiff = np.sum((mu1 - mu2)**2.0)
83     covmean = sqrtm(sigma1.dot(sigma2))

84
85     if np.iscomplexobj(covmean):
86         covmean = covmean.real

87
88     fid = ssdiff + np.trace(sigma1 + sigma2 - 2.0 *

```

```

        covmean)

83
84     return fid
85
86 def train_gan(epochs=1000, batch_size=64, p_epoch=100):
87     generator = build_generator()
88     discriminator = build_discriminator()
89
90     discriminator.compile(loss='binary_crossentropy',
91                           optimizer=Adam(0.0002, 0.5), metrics=['accuracy'])
91     discriminator.trainable = False
92
93     gan_input = tf.keras.Input(shape=(100,))
94     gan_output = discriminator(generator(gan_input))
95     gan = tf.keras.Model(gan_input, gan_output)
96     gan.compile(loss='binary_crossentropy', optimizer=
97                  Adam(0.0002, 0.5))
98
99     half_batch = int(batch_size / 2)
100
101    d_losses = []
102    g_losses = []
103    d_acc = []
104    fid_scores = []
105
105    start_time = time.time() # Record the start time
106
107    for epoch in range(epochs):
108        # Select a random half batch of real images
109        idx = np.random.randint(0, x_train.shape[0],
110                               half_batch)
111        real_images = x_train[idx]
112
112        # Generate a half batch of new fake images
113        noise = np.random.normal(0, 1, (half_batch, 100))
114        fake_images = generator.predict(noise)
115
116        # Train the discriminator
117        real_labels = np.ones((half_batch, 1))

```

```

118         fake_labels = np.zeros((half_batch, 1))
119
120         d_loss_real = discriminator.train_on_batch(
121             real_images, real_labels)
122         d_loss_fake = discriminator.train_on_batch(
123             fake_images, fake_labels)
124         d_loss = 0.5 * np.add(d_loss_real, d_loss_fake)
125
126         # Train the generator
127         noise = np.random.normal(0, 1, (batch_size, 100))
128         valid_y = np.ones((batch_size, 1))
129
130         g_loss = gan.train_on_batch(noise, valid_y)
131
132         # Record the losses
133         d_losses.append(d_loss[0])
134         g_losses.append(g_loss)
135         d_acc.append(d_loss[1] * 100)
136
137         # Calculate FID every p_epoch epochs
138         if epoch % p_epoch == 0:
139             noise = np.random.normal(0, 1, (1000, 100))
140             fake_images = generator.predict(noise)
141             fid = calculate_fid(x_train[:1000],
142                 fake_images)
143             fid_scores.append(fid)
144             print(f"epoch:{d_loss[0]},acc.:{100*d_loss[1]}%,[Gloss:{g_loss}][FID:{fid}]")
145
146             end_time = time.time() # Record the end time
147             total_time = end_time - start_time
148             print(f"Total training time:{total_time:.2f} seconds")
149
150             return generator, d_losses, g_losses, d_acc,
151                 fid_scores

```

Listing A.8: Explore GAN with more convolutional layers 3

```
1 import tensorflow as tf
```

```

2   from tensorflow.keras.layers import Dense, Reshape,
3       Flatten, Dropout, LeakyReLU, Conv2D, Conv2DTranspose,
4       BatchNormalization
5   from tensorflow.keras.models import Sequential
6   from tensorflow.keras.optimizers import Adam
7   from scipy.linalg import sqrtm
8   import numpy as np
9   import time
10  import matplotlib.pyplot as plt
11
12
13  # input 100
14  # output 28*28*1
15
16  def build_generator():
17      model = Sequential()
18
19      model.add(Dense(7*7*128, input_dim=100))
20      model.add(LeakyReLU(alpha=0.2))
21      model.add(Reshape((7, 7, 128)))
22      model.add(BatchNormalization(momentum=0.8))
23
24      # add the 1st convolution layer
25      model.add(Conv2D(128, kernel_size=3, strides=1,
26                      padding='same'))
27      model.add(LeakyReLU(alpha=0.2))
28      model.add(BatchNormalization(momentum=0.8))
29
30      # add the 2nd convolution layer
31      model.add(Conv2D(128, kernel_size=3, strides=1,
32                      padding='same'))
33      model.add(LeakyReLU(alpha=0.2))
34      model.add(BatchNormalization(momentum=0.8))
35
36      # add the 3rd convolution layer
37      model.add(Conv2D(128, kernel_size=3, strides=1,
38                      padding='same'))

```

```

36     model.add(LeakyReLU(alpha=0.2))
37     model.add(BatchNormalization(momentum=0.8))

38
39     model.add(Conv2DTranspose(128, kernel_size=4, strides
40         =2, padding='same'))
41     model.add(LeakyReLU(alpha=0.2))
42     model.add(BatchNormalization(momentum=0.8))

43
44     model.add(Conv2DTranspose(64, kernel_size=4, strides
45         =2, padding='same'))
46     model.add(LeakyReLU(alpha=0.2))
47     model.add(BatchNormalization(momentum=0.8))

48
49     model.add(Conv2D(1, kernel_size=7, activation='tanh',
50         padding='same'))

51
52     return model

53
54 # input 28*28*1
55 # output 1

56
57 def build_discriminator():
58     model = Sequential()

59
60     model.add(Conv2D(64, kernel_size=3, strides=2,
61         input_shape=(28, 28, 1), padding='same'))
62     model.add(LeakyReLU(alpha=0.2))
63     model.add(Dropout(0.25))

64
65     model.add(Conv2D(128, kernel_size=3, strides=2,
66         padding='same'))
67     model.add(LeakyReLU(alpha=0.2))
68     model.add(Dropout(0.25))

69
70     model.add(Flatten())
71     model.add(Dense(1, activation='sigmoid'))

72
73     return model

```

```

70     (x_train, _), (_, _) = tf.keras.datasets.mnist.load_data()
71     ()
72     x_train = (x_train - 127.5) / 127.5
73     x_train = np.expand_dims(x_train, axis=3)
74
75     def calculate_fid(real_images, fake_images):
76         act1 = real_images.reshape((real_images.shape[0], -1))
77         )
78         mu1, sigma1 = act1.mean(axis=0), np.cov(act1, rowvar=
79             False)
80
81         act2 = fake_images.reshape((fake_images.shape[0], -1))
82         )
83         mu2, sigma2 = act2.mean(axis=0), np.cov(act2, rowvar=
84             False)
85
86         ssdiff = np.sum((mu1 - mu2)**2.0)
87         covmean = sqrtm(sigma1.dot(sigma2))
88
89         if np.iscomplexobj(covmean):
90             covmean = covmean.real
91
92         fid = ssdiff + np.trace(sigma1 + sigma2 - 2.0 *
93             covmean)
94
95         return fid
96
97
98     def train_gan(epochs=1000, batch_size=64, p_epoch=100):
99         generator = build_generator()
100        discriminator = build_discriminator()
101
102        discriminator.compile(loss='binary_crossentropy',
103            optimizer=Adam(0.0002, 0.5), metrics=['accuracy'])
104        discriminator.trainable = False
105
106        gan_input = tf.keras.Input(shape=(100,))
107        gan_output = discriminator(generator(gan_input))
108        gan = tf.keras.Model(gan_input, gan_output)
109        gan.compile(loss='binary_crossentropy', optimizer=

```

```

          Adam(0.0002, 0.5))

102
103     half_batch = int(batch_size / 2)
104
105     d_losses = []
106     g_losses = []
107     d_acc = []
108     fid_scores = []
109
110     start_time = time.time() # Record the start time
111
112     for epoch in range(epochs):
113         # Select a random half batch of real images
114         idx = np.random.randint(0, x_train.shape[0],
115                               half_batch)
116         real_images = x_train[idx]
117
118         # Generate a half batch of new fake images
119         noise = np.random.normal(0, 1, (half_batch, 100))
120         fake_images = generator.predict(noise)
121
122         # Train the discriminator
123         real_labels = np.ones((half_batch, 1))
124         fake_labels = np.zeros((half_batch, 1))
125
126         d_loss_real = discriminator.train_on_batch(
127             real_images, real_labels)
128         d_loss_fake = discriminator.train_on_batch(
129             fake_images, fake_labels)
130         d_loss = 0.5 * np.add(d_loss_real, d_loss_fake)
131
132         # Train the generator
133         noise = np.random.normal(0, 1, (batch_size, 100))
134         valid_y = np.ones((batch_size, 1))
135
136         g_loss = gan.train_on_batch(noise, valid_y)

          # Record the losses
          d_losses.append(d_loss[0])

```

```

137     g_losses.append(g_loss)
138     d_acc.append(d_loss[1] * 100)
139
140     # Calculate FID every p_epoch epochs
141     if epoch % p_epoch == 0:
142         noise = np.random.normal(0, 1, (1000, 100))
143         fake_images = generator.predict(noise)
144         fid = calculate_fid(x_train[:1000],
145                             fake_images)
146         fid_scores.append(fid)
147         print(f'{epoch}[Dloss:{d_loss[0]},acc.:'
148               f'{100*d_loss[1]}%][Gloss:{g_loss}]'
149               f'[FID:{fid}]]')
150
151     end_time = time.time() # Record the end time
152     total_time = end_time - start_time
153     print(f'Total training time: {total_time:.2f} seconds')
154
155     return generator, d_losses, g_losses, d_acc,
156         fid_scores

```

Listing A.9: Explore GAN with more convolutional layers 4

```

1 import tensorflow as tf
2 from tensorflow.keras.layers import Dense, Reshape,
3     Flatten, Dropout, LeakyReLU, Conv2D, Conv2DTranspose,
4     BatchNormalization
5 from tensorflow.keras.models import Sequential
6 from tensorflow.keras.optimizers import Adam
7 from scipy.linalg import sqrtm
8 import numpy as np
9 import time
10 import matplotlib.pyplot as plt
11
12 np.random.seed(1000)
13 tf.random.set_seed(1000)
14
15 # input 100
16 # output 28*28*1

```

```

16 def build_generator():
17     model = Sequential()
18
19     model.add(Dense(7*7*128, input_dim=100))
20     model.add(LeakyReLU(alpha=0.2))
21     model.add(Reshape((7, 7, 128)))
22     model.add(BatchNormalization(momentum=0.8))
23
24     # add the 1st convolution layer
25     model.add(Conv2D(128, kernel_size=3, strides=1,
26                      padding='same'))
27     model.add(LeakyReLU(alpha=0.2))
28     model.add(BatchNormalization(momentum=0.8))
29
30     # add the 2nd convolution layer
31     model.add(Conv2D(128, kernel_size=3, strides=1,
32                      padding='same'))
33     model.add(LeakyReLU(alpha=0.2))
34     model.add(BatchNormalization(momentum=0.8))
35
36     # add the 3rd convolution layer
37     model.add(Conv2D(128, kernel_size=3, strides=1,
38                      padding='same'))
39     model.add(LeakyReLU(alpha=0.2))
40     model.add(BatchNormalization(momentum=0.8))
41
42
43     model.add(Conv2DTranspose(64, kernel_size=4, strides
44                             =2, padding='same'))
45     model.add(LeakyReLU(alpha=0.2))
46     model.add(BatchNormalization(momentum=0.8))
47
48     model.add(Conv2D(1, kernel_size=7, activation='tanh',
49                      padding='same'))

```

```

49         return model
50
51     # input 28*28*1
52     # output 1
53
54     def build_discriminator():
55         model = Sequential()
56
57         model.add(Conv2D(64, kernel_size=3, strides=2,
58                         input_shape=(28, 28, 1), padding='same'))
59         model.add(LeakyReLU(alpha=0.2))
60         model.add(Dropout(0.25))
61
62         model.add(Conv2D(128, kernel_size=3, strides=2,
63                         padding='same'))
64         model.add(LeakyReLU(alpha=0.2))
65         model.add(Dropout(0.25))
66
67         # add the 1st convolution layer
68         model.add(Conv2D(256, kernel_size=3, strides=2,
69                         padding='same'))
70         model.add(LeakyReLU(alpha=0.2))
71         model.add(Dropout(0.25))
72
73         model.add(Flatten())
74         model.add(Dense(1, activation='sigmoid'))
75
76         return model
77
78
79     (x_train, _), (_, _) = tf.keras.datasets.mnist.load_data()
80     ()
81     x_train = (x_train - 127.5) / 127.5
82     x_train = np.expand_dims(x_train, axis=3)
83
84     def calculate_fid(real_images, fake_images):
85         act1 = real_images.reshape((real_images.shape[0], -1))
86
87         mu1, sigma1 = act1.mean(axis=0), np.cov(act1, rowvar=False)

```

```

82
83     act2 = fake_images.reshape((fake_images.shape[0], -1)
84         )
85     mu2, sigma2 = act2.mean(axis=0), np.cov(act2, rowvar=
86         False)
87
88     ssdiff = np.sum((mu1 - mu2)**2.0)
89     covmean = sqrtm(sigma1.dot(sigma2))
90
91     if np.iscomplexobj(covmean):
92         covmean = covmean.real
93
94     fid = ssdiff + np.trace(sigma1 + sigma2 - 2.0 *
95         covmean)
96
97     return fid
98
99
100    def train_gan(epochs=1000, batch_size=64, p_epoch=100):
101        generator = build_generator()
102        discriminator = build_discriminator()
103
104        discriminator.compile(loss='binary_crossentropy',
105            optimizer=Adam(0.0002, 0.5), metrics=['accuracy'])
106        discriminator.trainable = False
107
108        gan_input = tf.keras.Input(shape=(100,))
109        gan_output = discriminator(generator(gan_input))
110        gan = tf.keras.Model(gan_input, gan_output)
111        gan.compile(loss='binary_crossentropy', optimizer=
112            Adam(0.0002, 0.5))
113
114        half_batch = int(batch_size / 2)
115
116        d_losses = []
117        g_losses = []
118        d_acc = []
119        fid_scores = []
120
121        start_time = time.time() # Record the start time

```

```

116
117     for epoch in range(epochs):
118         # Select a random half batch of real images
119         idx = np.random.randint(0, x_train.shape[0],
120                               half_batch)
121         real_images = x_train[idx]
122
123         # Generate a half batch of new fake images
124         noise = np.random.normal(0, 1, (half_batch, 100))
125         fake_images = generator.predict(noise)
126
127         # Train the discriminator
128         real_labels = np.ones((half_batch, 1))
129         fake_labels = np.zeros((half_batch, 1))
130
131         d_loss_real = discriminator.train_on_batch(
132             real_images, real_labels)
133         d_loss_fake = discriminator.train_on_batch(
134             fake_images, fake_labels)
135         d_loss = 0.5 * np.add(d_loss_real, d_loss_fake)
136
137         # Train the generator
138         noise = np.random.normal(0, 1, (batch_size, 100))
139         valid_y = np.ones((batch_size, 1))
140
141         g_loss = gan.train_on_batch(noise, valid_y)
142
143         # Record the losses
144         d_losses.append(d_loss[0])
145         g_losses.append(g_loss)
146         d_acc.append(d_loss[1] * 100)
147
148         # Calculate FID every p_epoch epochs
149         if epoch % p_epoch == 0:
150             noise = np.random.normal(0, 1, (1000, 100))

```

```
151         print(f'{epoch} [D_loss:{d_loss[0]}, acc.: {  
152             100*d_loss[1]}%] [G_loss:{g_loss}] [  
153             FID:{fid}]')  
  
154  
155     end_time = time.time() # Record the end time  
156     total_time = end_time - start_time  
157     print(f'Total training time:{total_time:.2f} seconds')  
158  
159     return generator, d_losses, g_losses, d_acc,  
160         fid_scores
```

Listing A.10: Explore GAN with more convolutional layers 5

```
1 import tensorflow as tf
2 from tensorflow.keras.layers import Dense, Reshape,
3     Flatten, Dropout, LeakyReLU, Conv2D, Conv2DTranspose,
4     BatchNormalization
5
6 from tensorflow.keras.models import Sequential
7 from tensorflow.keras.optimizers import Adam
8 from scipy.linalg import sqrtm
9
10 import numpy as np
11 import time
12
13 import matplotlib.pyplot as plt
14
15
16 np.random.seed(1000)
17 tf.random.set_seed(1000)
18
19
20 # input 100
21 # output 28*28*1
22
23
24 def build_generator():
25     model = Sequential()
26
27
28     model.add(Dense(7*7*128, input_dim=100))
29     model.add(LeakyReLU(alpha=0.2))
30     model.add(Reshape((7, 7, 128)))
31     model.add(BatchNormalization(momentum=0.8))
32
33
34     # add the 1st convolution layer
35     model.add(Conv2D(128, kernel_size=3, strides=1,
```

```

26         padding='same'))
27     model.add(LeakyReLU(alpha=0.2))
28     model.add(BatchNormalization(momentum=0.8))

29     # add the 2nd convolution layer
30     model.add(Conv2D(128, kernel_size=3, strides=1,
31                      padding='same'))
32     model.add(LeakyReLU(alpha=0.2))
33     model.add(BatchNormalization(momentum=0.8))

34     # add the 3rd convolution layer
35     model.add(Conv2D(128, kernel_size=3, strides=1,
36                      padding='same'))
37     model.add(LeakyReLU(alpha=0.2))
38     model.add(BatchNormalization(momentum=0.8))

39     model.add(Conv2DTranspose(128, kernel_size=4, strides
40                             =2, padding='same'))
41     model.add(LeakyReLU(alpha=0.2))
42     model.add(BatchNormalization(momentum=0.8))

43     model.add(Conv2DTranspose(64, kernel_size=4, strides
44                             =2, padding='same'))
45     model.add(LeakyReLU(alpha=0.2))
46     model.add(BatchNormalization(momentum=0.8))

47     model.add(Conv2D(1, kernel_size=7, activation='tanh',
48                      padding='same'))

49     return model

50
51 # input 28*28*1
52 # output 1

53
54 def build_discriminator():
55     model = Sequential()

56
57     model.add(Conv2D(64, kernel_size=3, strides=2,
58                      input_shape=(28, 28, 1), padding='same'))

```

```

58     model.add(LeakyReLU(alpha=0.2))
59     model.add(Dropout(0.25))
60
61     model.add(Conv2D(128, kernel_size=3, strides=2,
62                     padding='same'))
63     model.add(LeakyReLU(alpha=0.2))
64     model.add(Dropout(0.25))
65
66     model.add(Conv2D(256, kernel_size=3, strides=2,
67                     padding='same'))
68     model.add(LeakyReLU(alpha=0.2))
69     model.add(Dropout(0.25))
70
71     # add the 1st convolution layer
72     model.add(Conv2D(512, kernel_size=3, strides=2,
73                     padding='same'))
74     model.add(LeakyReLU(alpha=0.2))
75     model.add(Dropout(0.25))
76
77     # add the 2nd convolution layer
78     model.add(Conv2D(1024, kernel_size=3, strides=2,
79                     padding='same'))
80     model.add(LeakyReLU(alpha=0.2))
81     model.add(Dropout(0.25))
82
83     model.add(Flatten())
84     model.add(Dense(1, activation='sigmoid'))
85
86     return model
87
88
89
90
(x_train, _), (_, _) = tf.keras.datasets.mnist.load_data()
()
x_train = (x_train - 127.5) / 127.5
x_train = np.expand_dims(x_train, axis=3)

def calculate_fid(real_images, fake_images):
    act1 = real_images.reshape((real_images.shape[0], -1))
    mu1, sigma1 = act1.mean(axis=0), np.cov(act1, rowvar=

```

```

    False)

91
92     act2 = fake_images.reshape((fake_images.shape[0], -1)
93                               )
93     mu2, sigma2 = act2.mean(axis=0), np.cov(act2, rowvar=
94                               False)

94
95     ssdiff = np.sum((mu1 - mu2)**2.0)
96     covmean = sqrtm(sigma1.dot(sigma2))

97
98     if np.iscomplexobj(covmean):
99         covmean = covmean.real
100
101    fid = ssdiff + np.trace(sigma1 + sigma2 - 2.0 *
102                             covmean)
103
104
105    return fid
106
107
108
109    def train_gan(epochs=1000, batch_size=64, p_epoch=100):
110        generator = build_generator()
111        discriminator = build_discriminator()

112
113        discriminator.compile(loss='binary_crossentropy',
114                               optimizer=Adam(0.0002, 0.5), metrics=['accuracy'])
115        discriminator.trainable = False

116
117        gan_input = tf.keras.Input(shape=(100,))
118        gan_output = discriminator(generator(gan_input))
119        gan = tf.keras.Model(gan_input, gan_output)
120        gan.compile(loss='binary_crossentropy', optimizer=
121                     Adam(0.0002, 0.5))

122
123        half_batch = int(batch_size / 2)

124
125        d_losses = []
126        g_losses = []
127        d_acc = []
128        fid_scores = []

```

```

124     start_time = time.time() # Record the start time
125
126     for epoch in range(epochs):
127         # Select a random half batch of real images
128         idx = np.random.randint(0, x_train.shape[0],
129                               half_batch)
130         real_images = x_train[idx]
131
132         # Generate a half batch of new fake images
133         noise = np.random.normal(0, 1, (half_batch, 100))
134         fake_images = generator.predict(noise)
135
136         # Train the discriminator
137         real_labels = np.ones((half_batch, 1))
138         fake_labels = np.zeros((half_batch, 1))
139
140         d_loss_real = discriminator.train_on_batch(
141             real_images, real_labels)
142         d_loss_fake = discriminator.train_on_batch(
143             fake_images, fake_labels)
144         d_loss = 0.5 * np.add(d_loss_real, d_loss_fake)
145
146
147         # Train the generator
148         noise = np.random.normal(0, 1, (batch_size, 100))
149         valid_y = np.ones((batch_size, 1))
150
151         g_loss = gan.train_on_batch(noise, valid_y)
152
153
154         # Record the losses
155         d_losses.append(d_loss[0])
156         g_losses.append(g_loss)
157         d_acc.append(d_loss[1] * 100)
158
159
160         # Calculate FID every p_epoch epochs
161         if epoch % p_epoch == 0:
162             noise = np.random.normal(0, 1, (1000, 100))
163             fake_images = generator.predict(noise)
164             fid = calculate_fid(x_train[:1000],
165                               fake_images)

```

```

159         fid_scores.append(fid)
160         print(f"[epoch]: [D_loss: {d_loss[0]}, acc.: {100*d_loss[1]}%][G_loss: {g_loss}][FID: {fid}]")
161
162     end_time = time.time() # Record the end time
163     total_time = end_time - start_time
164     print(f"Total training time: {total_time:.2f} seconds")
165
166     return generator, d_losses, g_losses, d_acc,
167     fid_scores

```

Listing A.11: Explore GAN with more convolutional layers 6

```

1 def build_generator():
2     model = Sequential()
3
4     model.add(Dense(7*7*128, input_dim=100))
5     model.add(LeakyReLU(alpha=0.2))
6     model.add(Reshape((7, 7, 128)))
7     model.add(BatchNormalization(momentum=0.8))
8
9     # add the 1st convolution layer
10    model.add(Conv2D(128, kernel_size=3, strides=1, padding='same'))
11    model.add(LeakyReLU(alpha=0.2))
12    model.add(BatchNormalization(momentum=0.8))
13
14    # add the 2nd convolution layer
15    model.add(Conv2D(128, kernel_size=3, strides=1, padding='same'))
16    model.add(LeakyReLU(alpha=0.2))
17    model.add(BatchNormalization(momentum=0.8))
18
19    # add the 3rd convolution layer
20    model.add(Conv2D(128, kernel_size=3, strides=1, padding='same'))
21    model.add(LeakyReLU(alpha=0.2))
22    model.add(BatchNormalization(momentum=0.8))
23

```

```

24     model.add(Conv2DTranspose(128, kernel_size=4, strides=2,
25         padding='same'))
26     model.add(LeakyReLU(alpha=0.2))
27     model.add(BatchNormalization(momentum=0.8))
28
29     model.add(Conv2DTranspose(64, kernel_size=4, strides=2,
30         padding='same'))
31     model.add(LeakyReLU(alpha=0.2))
32     model.add(BatchNormalization(momentum=0.8))
33
34     model.add(Conv2D(1, kernel_size=7, activation='tanh',
35         padding='same'))
36
37     return model
38
39
40 def build_discriminator():
41     model = Sequential()
42
43     model.add(Conv2D(64, kernel_size=3, strides=2,
44         input_shape=(28, 28, 1), padding='same'))
45     model.add(LeakyReLU(alpha=0.2))
46     model.add(Dropout(0.25))
47
48     model.add(Conv2D(128, kernel_size=3, strides=2, padding='same'))
49     model.add(LeakyReLU(alpha=0.2))
50     model.add(Dropout(0.25))
51
52     model.add(Conv2D(256, kernel_size=3, strides=2, padding='same'))
53     model.add(LeakyReLU(alpha=0.2))
54     model.add(Dropout(0.25))
55
56     # add the 1st convolution layer
57     model.add(Conv2D(512, kernel_size=3, strides=2, padding='same'))
58     model.add(LeakyReLU(alpha=0.2))
59     model.add(Dropout(0.25))

```

```

56     # add the 2nd convolution layer
57     model.add(Conv2D(1024, kernel_size=3, strides=2, padding=
58                     'same'))
59     model.add(LeakyReLU(alpha=0.2))
60     model.add(Dropout(0.25))

61     # add the 3rd convolution layer
62     model.add(Conv2D(2048, kernel_size=3, strides=2, padding=
63                     'same'))
64     model.add(LeakyReLU(alpha=0.2))
65     model.add(Dropout(0.25))

66     model.add(Flatten())
67     model.add(Dense(1, activation='sigmoid'))

68
69     return model

70
71
72
73
74 import tensorflow as tf
75 from tensorflow.keras.layers import Dense, Reshape, Flatten,
76     Dropout, LeakyReLU, Conv2D, Conv2DTranspose,
77     BatchNormalization
78 from tensorflow.keras.models import Sequential
79 from tensorflow.keras.optimizers import Adam
80 from scipy.linalg import sqrtm
81 import numpy as np
82 import time
83 import matplotlib.pyplot as plt

84 np.random.seed(1000)
85 tf.random.set_seed(1000)

86 (x_train, _), (_, _) = tf.keras.datasets.mnist.load_data()
87 x_train = (x_train - 127.5) / 127.5
88 x_train = np.expand_dims(x_train, axis=3)

89 def calculate_fid(real_images, fake_images):

```

```

91     act1 = real_images.reshape((real_images.shape[0], -1))
92     mu1, sigma1 = act1.mean(axis=0), np.cov(act1, rowvar=
93         False)
94
95     act2 = fake_images.reshape((fake_images.shape[0], -1))
96     mu2, sigma2 = act2.mean(axis=0), np.cov(act2, rowvar=
97         False)
98
99     ssdiff = np.sum((mu1 - mu2)**2.0)
100    covmean = sqrtm(sigma1.dot(sigma2))
101
102    if np.iscomplexobj(covmean):
103        covmean = covmean.real
104
105    fid = ssdiff + np.trace(sigma1 + sigma2 - 2.0 * covmean)
106
107    return fid
108
109
110
111
112
113
114
115
116
117
118
119
120
121
122
123
124
125
def train_gan(epochs=1000, batch_size=64, p_epoch=100):
    generator = build_generator()
    discriminator = build_discriminator()

    discriminator.compile(loss='binary_crossentropy',
                           optimizer=Adam(0.0002, 0.5), metrics=['accuracy'])
    discriminator.trainable = False

    gan_input = tf.keras.Input(shape=(100,))
    gan_output = discriminator(generator(gan_input))
    gan = tf.keras.Model(gan_input, gan_output)
    gan.compile(loss='binary_crossentropy', optimizer=Adam
                (0.0002, 0.5))

    half_batch = int(batch_size / 2)

    d_losses = []
    g_losses = []
    d_acc = []
    fid_scores = []

```

```

126     start_time = time.time() # Record the start time
127
128     for epoch in range(epochs):
129         # Select a random half batch of real images
130         idx = np.random.randint(0, x_train.shape[0],
131                               half_batch)
132         real_images = x_train[idx]
133
134         # Generate a half batch of new fake images
135         noise = np.random.normal(0, 1, (half_batch, 100))
136         fake_images = generator.predict(noise)
137
138         # Train the discriminator
139         real_labels = np.ones((half_batch, 1))
140         fake_labels = np.zeros((half_batch, 1))
141
142         d_loss_real = discriminator.train_on_batch(
143             real_images, real_labels)
144         d_loss_fake = discriminator.train_on_batch(
145             fake_images, fake_labels)
146         d_loss = 0.5 * np.add(d_loss_real, d_loss_fake)
147
148         # Train the generator
149         noise = np.random.normal(0, 1, (batch_size, 100))
150         valid_y = np.ones((batch_size, 1))
151
152         g_loss = gan.train_on_batch(noise, valid_y)
153
154         # Record the losses
155         d_losses.append(d_loss[0])
156         g_losses.append(g_loss)
157         d_acc.append(d_loss[1] * 100)
158
159         # Calculate FID every p_epoch epochs
160         if epoch % p_epoch == 0:
161             noise = np.random.normal(0, 1, (1000, 100))
162             fake_images = generator.predict(noise)
163             fid = calculate_fid(x_train[:1000], fake_images)
164             fid_scores.append(fid)

```

```

162     print(f"epoch){D_loss:[d_loss[0]},acc.:100_
163         *d_loss[1]}%)[G_loss:[g_loss]][FID:[fid]]
164     ")
165
166     end_time = time.time() # Record the end time
167     total_time = end_time - start_time
168     print(f"Total training time:{total_time:.2f}seconds")
169     return generator, d_losses, g_losses, d_acc, fid_scores

```

Listing A.12: Apply new dataset

```

1
2     import os
3     import time
4     import numpy as np
5     import tensorflow as tf
6     import matplotlib.pyplot as plt
7     from tensorflow.keras.models import Sequential, Model
8     from tensorflow.keras.layers import Dense, Reshape,
9         BatchNormalization, LeakyReLU, Conv2D, Conv2DTranspose
10        , Flatten, Dropout, Input
11        from tensorflow.keras.optimizers import Adam
12        from tensorflow.keras.preprocessing.image import load_img
13        , img_to_array
14
15
16
17
18
19
20
21
22
23
24
25

```

```

11
12     def load_images_as_rgb_matrices(path, size):
13         images = []
14         for img_name in os.listdir(path):
15             img_path = os.path.join(path, img_name)
16             img = load_img(img_path, target_size=size)
17             img_array = img_to_array(img)
18             images.append(img_array)
19         images = np.array(images)
20         images = (images - 127.5) / 127.5 # Normalize images to
21             [-1, 1]
22         return images
23
24
25     from google.colab import drive

```

```

26     drive.mount('/content/drive')
27
28     cat_path = "/content/drive/My_Drive/gan/afhq/train/cat"
29     size = (128, 128)
30     x_train = load_images_as_rgb_matrices(cat_path, size)
31
32     # generator
33     # input 100
34     # output 128*128*3
35
36     def build_generator():
37         model = Sequential()
38
39         # Increase the dimension
40         model.add(Dense(16*16*256, input_dim=100))
41         model.add(LeakyReLU(alpha=0.2))
42         model.add(Reshape((16, 16, 256)))
43         model.add(BatchNormalization(momentum=0.8))
44
45         model.add(Conv2DTranspose(256, kernel_size=4, strides
46             =2, padding='same')) # 32x32
47         model.add(LeakyReLU(alpha=0.2))
48         model.add(BatchNormalization(momentum=0.8))
49
50         model.add(Conv2DTranspose(128, kernel_size=4, strides
51             =2, padding='same')) # 64x64
52         model.add(LeakyReLU(alpha=0.2))
53         model.add(BatchNormalization(momentum=0.8))
54
55         model.add(Conv2DTranspose(64, kernel_size=4, strides
56             =2, padding='same')) # 128x128
57         model.add(LeakyReLU(alpha=0.2))
58         model.add(BatchNormalization(momentum=0.8))
59
60         model.add(Conv2D(3, kernel_size=7, activation='tanh',
61             padding='same')) # 128x128x3
62
63     return model

```

```

61     # discriminator
62     # input 128*128*3
63     # output 1
64
65     def build_discriminator():
66         model = Sequential()
67
68         model.add(Conv2D(64, kernel_size=3, strides=2,
69                         input_shape=(128, 128, 3), padding='same'))
70         model.add(LeakyReLU(alpha=0.2))
71         model.add(Dropout(0.25))
72
73         model.add(Conv2D(128, kernel_size=3, strides=2,
74                         padding='same'))
75         model.add(LeakyReLU(alpha=0.2))
76         model.add(Dropout(0.25))
77
78         model.add(Conv2D(256, kernel_size=3, strides=2,
79                         padding='same'))
80         model.add(LeakyReLU(alpha=0.2))
81         model.add(Dropout(0.25))
82
83         model.add(Flatten())
84         model.add(Dense(1, activation='sigmoid'))
85
86
87     return model
88
89
90     # initial GAN
91     def build_gan(generator, discriminator):
92         discriminator.trainable = False
93         gan_input = Input(shape=(100,))
94         x = generator(gan_input)
95         gan_output = discriminator(x)
96         gan = Model(gan_input, gan_output)

```

```

96     gan.compile(loss='binary_crossentropy', optimizer=tf.
97                   keras.optimizers.legacy.Adam(0.0002, 0.5))
98     return gan
99
100    # define training step
101    def train(generator, discriminator, gan, x_train, epochs,
102              batch_size=128):
103        valid = np.ones((batch_size, 1))
104        fake = np.zeros((batch_size, 1))
105
106
107
108        for epoch in range(epochs):
109            start_time = time.time()
110
111
112            idx = np.random.randint(0, x_train.shape[0],
113                                   batch_size)
114            imgs = x_train[idx]
115
116            noise = np.random.normal(0, 1, (batch_size, 100))
117            gen_imgs = generator.predict(noise)
118
119            d_loss_real = discriminator.train_on_batch(imgs,
120                                                       valid)
121            d_loss_fake = discriminator.train_on_batch(
122                gen_imgs, fake)
123            d_loss = 0.5 * np.add(d_loss_real, d_loss_fake)
124
125
126            noise = np.random.normal(0, 1, (batch_size, 100))
127            g_loss = gan.train_on_batch(noise, valid)
128
129            end_time = time.time()
130            epoch_time = end_time - start_time
131
132            print(f"epoch [{d_loss[0]}] | D accuracy : {100*d_loss[1]}% [G loss: {g_loss}] [Epoch time: {epoch_time:.2f} seconds]")
133
134
135
136
137
138
139
140
141
142
143
144
145
146
147
148
149
150
151
152
153
154
155
156
157
158
159
160
161
162
163
164
165
166
167
168
169
170
171
172
173
174
175
176
177
178
179
180
181
182
183
184
185
186
187
188
189
190
191
192
193
194
195
196
197
198
199
200
201
202
203
204
205
206
207
208
209
210
211
212
213
214
215
216
217
218
219
220
221
222
223
224
225
226
227
228
229
230
231
232
233
234
235
236
237
238
239
240
241
242
243
244
245
246
247
248
249
250
251
252
253
254
255
256
257
258
259
260
261
262
263
264
265
266
267
268
269
270
271
272
273
274
275
276
277
278
279
280
281
282
283
284
285
286
287
288
289
290
291
292
293
294
295
296
297
298
299
300
301
302
303
304
305
306
307
308
309
310
311
312
313
314
315
316
317
318
319
320
321
322
323
324
325
326
327
328
329
330
331
332
333
334
335
336
337
338
339
340
341
342
343
344
345
346
347
348
349
350
351
352
353
354
355
356
357
358
359
360
361
362
363
364
365
366
367
368
369
370
371
372
373
374
375
376
377
378
379
380
381
382
383
384
385
386
387
388
389
390
391
392
393
394
395
396
397
398
399
400
401
402
403
404
405
406
407
408
409
410
411
412
413
414
415
416
417
418
419
420
421
422
423
424
425
426
427
428
429
430
431
432
433
434
435
436
437
438
439
440
441
442
443
444
445
446
447
448
449
450
451
452
453
454
455
456
457
458
459
460
461
462
463
464
465
466
467
468
469
470
471
472
473
474
475
476
477
478
479
480
481
482
483
484
485
486
487
488
489
490
491
492
493
494
495
496
497
498
499
500
501
502
503
504
505
506
507
508
509
510
511
512
513
514
515
516
517
518
519
520
521
522
523
524
525
526
527
528
529
530
531
532
533
534
535
536
537
538
539
540
541
542
543
544
545
546
547
548
549
550
551
552
553
554
555
556
557
558
559
560
561
562
563
564
565
566
567
568
569
570
571
572
573
574
575
576
577
578
579
580
581
582
583
584
585
586
587
588
589
589

```

```
128 generator = build_generator()
129 discriminator = build_discriminator()
130 discriminator.compile(loss='binary_crossentropy',
131     optimizer=tf.keras.optimizers.legacy.Adam(0.0002, 0.5)
132     , metrics=['accuracy'])
133 gan = build_gan(generator, discriminator)
134
135 # gen images
136 def show_generated_images(generator, num_images=25, dim
137     =(5, 5), figsize=(10, 10)):
138     noise = np.random.normal(0, 1, (num_images, 100))
139     gen_imgs = generator.predict(noise)
140     gen_imgs = 0.5 * gen_imgs + 0.5
141
142     plt.figure(figsize=figsize)
143     for i in range(num_images):
144         plt.subplot(dim[0], dim[1], i+1)
145         plt.imshow(gen_imgs[i])
146         plt.axis('off')
147     plt.tight_layout()
148     plt.show()
149
150 show_generated_images(generator)
```

Bibliography

- [1] I. Gulrajani, F. Ahmed, M. Arjovsky, V. Dumoulin, and A. Courville. Improved training of wasserstein gans. 2017.
- [2] Y. Jiang. Applications of generative adversarial networks in materials science. *Materials Genome Engineering Advances*, 2, 2024.
- [3] Y. Xin, E. Walia, and P. Babyn. Generative adversarial network in medical imaging: a review. *Medical Image Analysis*, 58:101552, 2019.
- [4] S. Kazeminia, C. Baur, A. Kuijper, B. Ginneken, N. Navab, S. Albarqouni, and A. Mukhopadhyay. Gans for medical image analysis. *Artificial Intelligence in Medicine*, 109:101938, 2020.
- [5] M. Frid-Adar, I. Diamant, E. Klang, M. M. Amitai, J. Goldberger, and H. Greenspan. Gan-based synthetic medical image augmentation for increased cnn performance in liver lesion classification. *Neurocomputing*, 321:321–331, 2018.
- [6] R. Abdal, Y. Qin, and P. Wonka. Image2stylegan: how to embed images into the stylegan latent space? *2019 IEEE/CVF International Conference on Computer Vision (ICCV)*, 2019.
- [7] L. Han, Y. Huang, and T. Zhang. Candidates vs. noises estimation for large multi-class classification problem. 2017.
- [8] B. Liu, E. Rosenfeld, P. Ravikumar, and A. Risteski. Analyzing and improving the optimization landscape of noise-contrastive estimation. 2021.

- [9] M. Labeau and A. Allauzen. An experimental analysis of noise-contrastive estimation: the noise distribution matters. 2017.
- [10] B. Damavandi, S. Kumar, N. Shazeer, and A. Bruguier. Nn-grams: unifying neural network and n-gram language models for speech recognition. 2016.
- [11] Y. Song. How to train your energy-based models. 2021.
- [12] D. Kingma and M. Welling. An introduction to variational autoencoders. *Foundations and Trends® in Machine Learning*, 12:307–392, 2019.
- [13] B. Kiran, D. Thomas, and R. Parakkal. An overview of deep learning based methods for unsupervised and semi-supervised anomaly detection in videos. *Journal of Imaging*, 4:36, 2018.
- [14] Y. Varolgunes, T. Bereau, and J. Rudzinski. Interpretable embeddings from molecular simulations using gaussian mixture variational autoencoders. *Machine Learning Science and Technology*, 1:015012, 2020.
- [15] S. Portillo. Dimensionality reduction of sdss spectra with variational autoencoders. 2020.
- [16] C. Guo, J. Zhou, H. Chen, N. Ying, J. Zhang, and D. Zhou. Variational autoencoder with optimizing gaussian mixture model priors. *Ieee Access*, 8:43992–44005, 2020.
- [17] P. Munjal, A. Paul, and N. Krishnan. Implicit discriminator in variational autoencoder. 2019.
- [18] X. Bie, L. Girin, S. Leglaive, T. Hueber, and X. Alameda-Pineda. A benchmark of dynamical variational autoencoders applied to speech spectrogram modeling. 2021.
- [19] P. Dhariwal. Diffusion models beat gans on image synthesis. 2021.
- [20] Z. Kong. Diffwave: a versatile diffusion model for audio synthesis. 2020.

- [21] Q. Zhang, M. Tao, and Y. Chen. Gddim: generalized denoising diffusion implicit models. 2022.
- [22] S. Liu, D. Su, and D. Yu. Diffgan-tts: high-fidelity and efficient text-to-speech with denoising diffusion gans. 2022.
- [23] J. Wu, H. Fang, Y. Yang, and Y. Xu. Medsegdiff: medical image segmentation with diffusion probabilistic model. 2022.
- [24] K. Pandey, A. Mukherjee, P. Rai, and A. Kumar. Diffusevae: efficient, controllable and high-fidelity generation from low-dimensional latents. 2022.
- [25] M. Fathallah, M. Sakr, and S. El-etriby. Stabilizing and improving training of generative adversarial networks through identity blocks and modified loss function. *Ieee Access*, 11:43276–43285, 2023.
- [26] J. Mu, C. Chen, W. Zhu, S. Li, and Y. Zhou. Taming mode collapse in generative adversarial networks using cooperative realness discriminators. *Iet Image Processing*, 16:2240–2262, 2022.
- [27] Y. Zou. Auto-encoding generative adversarial networks towards mode collapse reduction and feature representation enhancement. *Entropy*, 25:1657, 2023.
- [28] Y. Gong, Z. Ming, J. Yang, M. Xie, and X. Ma. Distribution constraining for combating mode collapse in generative adversarial networks. *Journal of Electronic Imaging*, 32, 2023.
- [29] A. Dieng. Prescribed generative adversarial networks. 2019.
- [30] J. Dubinski, K. Deja, S. Wenzel, P. Rokita, and T. Trzciński. Selectively increasing the diversity of gan-generated samples. 2022.
- [31] Y. Chen, Q. Gao, and W. Xiao. Inferential wasserstein generative adversarial networks. *Journal of the Royal Statistical Society Series B (Statistical Methodology)*, 84:83–113, 2021.

- [32] H. Zhu. Comparison of deep convolutional gan and progressive gan for facial image generation. *Applied and Computational Engineering*, 18:165–171, 2023.
- [33] H. Zhao, T. Li, Y. Xiao, and Y. Wang. Improving multi-agent generative adversarial nets with variational latent representation. *Entropy*, 22:1055, 2020.
- [34] J. Ho and T. Salimans. Classifier-free diffusion guidance. 2022.
- [35] J. Song, C. Meng, and S. Ermon. Denoising diffusion implicit models. 2020.
- [36] K. Preechakul, N. Chatthee, S. Wizadwongsa, and S. Suwajanakorn. Diffusion autoencoders: toward a meaningful and decodable representation. 2021.
- [37] E. Luhman. Knowledge distillation in iterative generative models for improved sampling speed. 2021.
- [38] A. Creswell, T. White, V. Dumoulin, K. Arulkumaran, B. Sengupta, and A. A. Bharath. Generative adversarial networks: an overview. *IEEE Signal Processing Magazine*, 35:53–65, 2018.
- [39] I. Goodfellow, J. Pouget-Abadie, M. Mirza, B. Xu, D. Warde-Farley, S. Ozair, A. Courville, and Y. Bengio. Generative adversarial networks. *Communications of the ACM*, 63:139–144, 2020.
- [40] A. Kazerouni, E. Aghdam, M. Heidari, R. Azad, M. Fayyaz, I. Haci-haliloglu, and D. Merhof. Diffusion models for medical image analysis: a comprehensive survey. 2022.
- [41] Y. Song. Score-based generative modeling through stochastic differential equations. 2020.
- [42] A. Lugmayr, M. Danelljan, A. Romero, F. Yu, R. Timofte, and L. Gool. Repaint: inpainting using denoising diffusion probabilistic models. 2022.

- [43] Y. Saito, S. Takamichi, and H. Saruwatari. Statistical parametric speech synthesis incorporating generative adversarial networks. *Ieee/Acm Transactions on Audio Speech and Language Processing*, 26:84–96, 2018.
- [44] J. Parikh, T. Rumbell, X. Butova, T. Myachina, J. Acero, S. Khamzin, O. Solovyova, J. Kozloski, A. Khokhlova, and V. Gurev. Generative adversarial networks for construction of virtual populations of mechanistic models: simulations to study omecamtiv mecarbil action. *Journal of Pharmacokinetics and Pharmacodynamics*, 49:51–64, 2021.
- [45] T. Miyato. Cgans with projection discriminator. 2018.
- [46] Ian J. Goodfellow, Jean Pouget-Abadie, Mehdi Mirza, Bing Xu, David Warde-Farley, Sherjil Ozair, Aaron Courville, and Yoshua Bengio. Generative adversarial networks. 2014.
- [47] T. Kynkäanniemi, T. Karras, M. Aittala, T. Aila, and J. Lehtinen. The role of imagenet classes in fréchet inception distance. 2022.
- [48] Y. Xu, T. Wu, J. Charlton, and K. Bennett. Gan training acceleration using fréchet descriptor-based coreset. *Applied Sciences*, 12:7599, 2022.
- [49] R. Xu, J. Wang, J. Liu, F. Ni, and B. Cao. Thermal infrared face image data enhancement method based on deep learning. 2023.
- [50] D. Berthelot, T. Schumm, and L. Metz.Began: boundary equilibrium generative adversarial networks. 2017.
- [51] H. Hyungrok, T. Jun, and D. Kim. Unbalanced gans: pre-training the generator of generative adversarial network using variational autoencoder. 2020.