

Generating images using a Deep Convolutional GANs



Author: Kai Deng

Supervisor: Serhiy Yanchuk

A thesis submitted in partial fulfilment of the
requirements for the degree of
MSc Mathematical Modelling and Machine Learning

School of Mathematical Sciences,
University College Cork,
Ireland

September 2024

Declaration of Authorship

This report is wholly the work of the author, except where explicitly stated otherwise. The source of any material which was not created by the author has been clearly cited.

Date: 10/09/2024

Signature: Kai Deng

Acknowledgements

I would like to express my deepest gratitude to my supervisor, Dr Serhiy Yanchuk, for his invaluable guidance, continuous support, and insightful feedback throughout the course of my research. His expertise and encouragement have been instrumental in the successful completion of this thesis.

Abstract

Generative Adversarial Networks (GANs) have become a pivotal tool in the field of generative modeling, offering a novel approach to generating realistic data through the adversarial training of a generator and a discriminator. This thesis investigates the performance of GANs in generating high-quality images, with a particular focus on the challenges associated with model structure, training stability, and evaluation metrics.

In this study, I employed the Animal Faces-HQ (AFHQ) dataset, consisting of 16,130 high-resolution images, to train a standard GAN model aimed at generating realistic cat images. The research explored the impact of various architectural choices, including the use of convolutional layers versus dense layers, as well as the effects of data augmentation techniques on model performance. The experiments revealed that convolutional architectures significantly outperformed dense ones, highlighting the importance of spatial feature extraction in image generation tasks.

The study also examined the role of data augmentation in GAN training, finding that while augmentation can introduce useful variability into the training data, it may also lead to optimization challenges if not carefully applied. The final model, trained on downscaled images due to hardware limitations, was able to produce cat images, demonstrating the effectiveness of GANs in image synthesis.

The findings contribute to a deeper understanding of how GAN architecture and training strategies influence the quality of generated images, offering insights for future research and practical applications in fields requiring high-quality image generation. The study concludes by suggesting directions for further exploration, including the potential benefits of more advanced GAN variants and higher-resolution training data.

Contents

| | | |
|----------|--|-----------|
| 1 | Introduction | 5 |
| 2 | Historical Models of Image Generation | 7 |
| 2.1 | Deep Boltzmann Machines | 7 |
| 2.2 | Variational Autoencoders | 8 |
| 2.3 | Noise Contrastive Estimation | 9 |
| 3 | Theoretical Background | 11 |
| 3.1 | Objective Function of GAN | 11 |
| 3.2 | Training Process of GAN | 14 |
| 3.2.1 | Distribution Changes During GAN Training | 15 |
| 3.2.2 | Mathematical Formulation Changes During GAN Training | 16 |
| 3.3 | Evaluating GAN Performance | 19 |
| 3.4 | Limitations of Accuracy in Evaluating GANs | 20 |
| 4 | Results | 22 |
| 4.1 | Standard GAN Versus Other GAN Realizations | 22 |
| 4.2 | GAN With Convolutional or Dense layers | 23 |
| 4.3 | Exploring Layer Depth | 27 |
| 4.4 | Impact of Data Augmentation on Model Performance | 28 |
| 4.5 | Applying the Model Animal Faces-HQ Dataset | 34 |
| 5 | Discussion | 35 |
| 5.1 | Summary of Key Findings | 35 |

| | | |
|----------|---------------------------------------|-----------|
| 5.2 | Impact of Model Structure | 35 |
| 5.3 | Data Augmentation Effects | 36 |
| 5.4 | Performance on New Dataset | 36 |
| 5.5 | Limitations and Future Work | 36 |
| 5.6 | Conclusion | 37 |
| A | Code | 38 |

List of Figures

| | | |
|-----|---|----|
| 3.1 | Comparison of $\log(1 - D)$ and $-\log(D)$ | 14 |
| 3.2 | Diagram of GAN’s training process. The green curve represents the distribution of generated samples. Initially, the generated samples may differ significantly from the real samples. As training progresses, the generated samples’ distribution gradually approaches the real samples. The black dots represent the distribution of real samples, which remains unchanged throughout the training process and represents the target distribution. The blue dashed line represents the discriminator’s output probability distribution. At the beginning of training, the discriminator can easily distinguish between real and generated data, resulting in a strong classification boundary. As training progresses and the generated data becomes more realistic, the discriminator’s ability to differentiate between the two distributions weakens. Eventually, the discriminator’s output approaches 0.5, indicating it can no longer effectively distinguish between real and generated data. The lines labeled x and z below represent the distribution of samples in the latent space. During GAN training, samples from the latent space z are mapped to the data space x through the generator. | 15 |
| 3.3 | GAN Training Accuracy Over Epochs | 20 |
| 3.4 | Generated Images from GAN | 21 |
| 4.1 | Generator Architecture with Dense and Convolutional Layers . | 24 |

| | | |
|------|--|----|
| 4.2 | Discriminator Architecture with Dense and Convolutional Layers | 25 |
| 4.3 | Comparison of GAN Performance at 3000 Epochs | 25 |
| 4.4 | FID Scores Across 3000 Epochs | 26 |
| 4.5 | Generated Images from Different GAN Architectures | 28 |
| 4.6 | The FID scores for standard GAN without data augmentation | 29 |
| 4.7 | The FID scores for standard GAN with data augmentation . . | 30 |
| 4.8 | The FID scores for standard GAN with data augmentation but remove flip | 32 |
| 4.9 | The FID scores for standard GAN with data augmentation but remove flip and rotation | 33 |
| 4.10 | Cat Faces Generated by GAN | 34 |

Chapter 1

Introduction

Generative Adversarial Networks (GANs) have emerged as a powerful class of generative models that revolutionize generative modeling by framing it as a game between two networks: a generator network that produces synthetic data from noise and a discriminator network that distinguishes between the generated data and real data [1]. These networks, introduced in 2014, have found applications in various fields, including materials science, radiology, and computer vision [2], [3], [4]. For example, CycleGAN has been effectively applied in the medical field, notably in medical imaging tasks. It has enhanced liver lesion classification through GAN-based synthetic medical image augmentation, surpassing traditional data augmentation methods in sensitivity and specificity [5]. Abdal et al. (2019) demonstrated the effectiveness of StyleGAN in tasks such as image deformation and style transfer [6]. GANs have gained significant attention in the computer vision community due to their ability to generate data without explicitly modeling the probability density function [3].

Despite these advances, training GANs remains a challenging task due to issues such as mode collapse, instability during training, and the requirement for substantial computational resources. Additionally, evaluating the performance of GANs is not straightforward, as traditional metrics such as accuracy do not adequately capture the quality or diversity of the generated data. This complexity has led to the development of various GAN architec-

tures and training techniques aimed at improving the stability and quality of generated images.

The purpose of this thesis is to contribute to a better understanding of the factors influencing the performance of GANs, particularly in the context of generating realistic images. To achieve this, I will utilize the Animal Faces-HQ (AFHQ) dataset, which comprises 16,130 high-quality images with a resolution of 512×512 pixels, to train a standard GAN model specifically designed for generating realistic cat pictures. The high resolution of these images presents both an opportunity and a challenge, as it requires careful consideration of model structure and training strategies to avoid issues such as overfitting or insufficient model capacity.

This thesis is structured as follows: Chapter 2 provides an overview of historical models of image generation, including Deep Boltzmann Machines, Variational Autoencoders, and Noise Contrastive Estimation. Chapter 3 discusses the theoretical background of GANs, focusing on their objective functions, the training process, and methods for evaluating their performance. Chapter 4 presents the experimental work conducted, including model selection, architecture exploration, and the impact of data augmentation, followed by the application of the GAN model to the AFHQ dataset. Finally, Chapter 5 discusses the results, highlights the implications of the findings, and suggests directions for future research.

Chapter 2

Historical Models of Image Generation

This chapter reviews key historical models of image generation, including Deep Boltzmann Machines (DBMs), Variational Autoencoders (VAEs), and Noise Contrastive Estimation (NCE). Each of these models has contributed to the development of generative modeling by introducing different approaches to learning data distributions.

2.1 Deep Boltzmann Machines

Deep Boltzmann Machines (DBMs) are sophisticated generative models that excel in learning intricate data representations in an unsupervised manner. Unlike traditional feedforward networks, DBMs are structured with multiple layers of stochastic, binary latent variables, where each layer captures progressively abstract features of the input data. This architecture is characterized by fully connected layers, with no direct connections between units within the same layer, which enables DBMs to effectively model complex dependencies inherent in the data [7]. The undirected nature of the connections in DBMs allows for the propagation of uncertainties across layers, enhancing their capability in feature learning and unsupervised pre-training [7][8].

However, the training of DBMs presents significant challenges, primarily

due to their reliance on Markov Chain Monte Carlo (MCMC) methods, such as Gibbs sampling, for approximate inference. These methods are essential for estimating gradients and the intractable partition function, which is vital for updating model parameters. The computational overhead introduced by MCMC methods often results in slow convergence and increased complexity, particularly when applied to large-scale datasets [9]. The requirement for numerous sampling steps to achieve equilibrium further complicates the training process, making DBMs less accessible compared to other generative models, such as Variational Autoencoders (VAEs).

To mitigate these training difficulties, DBMs often utilize layer-wise pre-training, which initializes model parameters before the entire network is fine-tuned. This approach has been shown to enhance the efficiency of the learning process by providing a sensible initialization for the weights and variational inference [10]. Despite these strategies, the inherent complexity associated with MCMC methods continues to pose challenges in scaling DBMs effectively, particularly in comparison to more straightforward generative models like VAEs [9]. Recent advancements have explored alternative training methodologies, including the use of specialized hardware systems to improve the efficiency of sampling tasks, thereby addressing some of the computational hurdles associated with traditional training approaches [9].

2.2 Variational Autoencoders

Variational Autoencoders (VAEs) are generative models in machine learning that excel in tasks like image generation and data representation by learning the underlying structure of data through a latent space.

The architecture of VAEs consists of an encoder that maps input data into a latent space, where the latent variables are typically assumed to follow a Gaussian distribution. This assumption simplifies the learning process, as it allows for the use of techniques such as the reparameterization trick, which enables backpropagation through stochastic layers [11]. The decoder then reconstructs the input data from the latent variables, ensuring that the model captures the essential features of the data distribution. The loss

function of a VAE combines reconstruction loss with a Kullback-Leibler (KL) divergence term, which regularizes the latent space and encourages the model to learn a smooth and continuous representation [12].

In comparison to DBMs, VAEs offer several advantages, including simpler training dynamics and better scalability. DBMs often require complex sampling methods, such as Gibbs sampling, which can be computationally intensive and slow to converge. In contrast, VAEs can be trained using standard gradient descent methods, making them more accessible for practitioners [11][13]. Furthermore, the structured latent space of VAEs not only allows for efficient sampling but also supports various applications, such as image generation, anomaly detection, and data imputation, where the ability to interpolate between data points is crucial [14][15].

However, VAEs face limitations in modeling discrete data as they require back-propagation through hidden units, which poses challenges in handling such data types effectively [16]. Despite this limitation, VAEs have been extensively used to represent high-dimensional complex data by learning a low-dimensional latent space in an unsupervised manner [17].

2.3 Noise Contrastive Estimation

Noise-Contrastive Estimation (NCE) is a statistical method employed in machine learning, particularly for estimating parameters of unnormalized probabilistic models. The fundamental principle of NCE is to recast the maximum likelihood estimation problem into a binary classification task. In this framework, the model is trained to differentiate between actual data samples and artificially generated noise samples, which are drawn from a known distribution [18]. This transformation simplifies the estimation process and circumvents the computational challenges associated with calculating the intractable partition function, a common hurdle in traditional maximum likelihood estimation methods for unnormalized models [19].

The methodology of NCE involves the introduction of noise samples, which serve as a baseline for comparison against real data. The model is trained to assign higher probabilities to genuine data samples while assign-

ing lower probabilities to noise samples. This approach effectively enables the model to learn the underlying data distribution without the necessity of explicit normalization [20]. The binary classification framework utilized in NCE allows for the application of standard logistic regression techniques, enhancing both the tractability and computational efficiency of the training process [21]. Moreover, NCE has been shown to be particularly advantageous in large-scale models, such as energy-based models and word embeddings, where the computational burden of normalization can be substantial [22].

Despite its advantages, NCE is not without limitations. The choice of noise distribution is critical to the performance of the model; poorly selected noise distributions can lead to suboptimal parameter estimates and slow convergence rates [19]. Additionally, while NCE simplifies the training of unnormalized models, it may not perform as effectively in scenarios where the noise distribution is challenging to define or when dealing with highly complex data distributions. The empirical observations regarding the importance of noise distribution have been formalized in various studies, underscoring the need for careful consideration in its selection [19].

Chapter 3

Theoretical Background

This chapter covers the theoretical foundations of Generative Adversarial Networks (GANs), including their objective function, the training dynamics between the generator and discriminator, and the use of metrics like Fréchet Inception Distance (FID) to evaluate performance.

3.1 Objective Function of GAN

Generative Adversarial Networks (GANs) consist of a generator G and a discriminator D , both implemented using artificial neural networks. The parametrization of GANs involves defining the network structure of these components and initializing their weights and biases [23]. The success of GANs relies on balancing the training of these two networks, where the G aims to produce samples $G(z)$ that mimic real data distributions $p_{data}(x)$ by input noise z , while the D learns to differentiate between real and generated samples [24]. The training process involves iteratively updating the weights and biases of the networks through adversarial training, where the generator tries to deceive the discriminator, and the discriminator aims to accurately classify samples [25]. The objective function for GAN has the following form:

$$V(D, G) = \mathbb{E}_{x \sim p_{data}(x)}[\log D(x)] + \mathbb{E}_{z \sim p_z(z)}[\log(1 - D(G(z)))] \quad (3.1)$$

It describes the game process between the Generator (G) and the Discriminator (D). Specifically, the formula defines a minimax game between the Discriminator and the Generator.

- $z \sim p_z(z)$ represents the distribution from which the noise z is sampled. The purpose of using a noise distribution is to provide random input for the generator, which transforms this noise into more complex, high-dimensional data.

Example: If we assume $p_z(z) = \mathcal{N}(0, 1)$, then z could be a vector of random values sampled independently from $\mathcal{N}(0, 1)$. The generator would then take this random vector z and map it to a synthetic image, such as a cat image, that resembles data from the real world.

- $x \sim p_{data}(x)$ represent that sample x is drawn from the true data distribution p_{data} .
- $D(G(z))$ represent the output of the discriminator for the data $G(z)$ generated by the generator. Indicating the probability that the discriminator believes that the generated data comes from the real data distribution.
- $D(x)$ represent the output of the discriminator for the real data x . Indicating the probability that the discriminator believes that the real data x comes from the real data distribution.
- $\mathbb{E}_x \sim p_{data}(x)[\log D(x)]$ represent the average value of $\log D(x)$ for all samples x on the true data distribution $p_{data}(x)$.
- $\mathbb{E}x \sim p_z(z)[\log(1-D(x))]$ represent the average value of $\log(1-D(G(z)))$ for all samples z on the noise distribution $p_z(z)$.

The following formulas:

$$\max_D V(D, G) = \mathbb{E}_{x \sim p_{data}(x)}[\log D(x)] + \mathbb{E}_{z \sim p_z(z)}[\log(1 - D(G(z)))] \quad (3.2)$$

$$\min_G V(G) = \mathbb{E}_{z \sim p_z(z)} [\log(1 - D(G(z)))] \quad (3.3)$$

explains the continuous balance between the generator and discriminator. As the generator improves its ability to create more realistic data, the discriminator adjusts to better distinguish real data from generated data. This interaction drives the GAN model to produce outputs that resemble real data more closely.

The formula (3.2) explains the goal of the discriminator, which is to maximize the value of $V(D, G)$. The first component measures how likely the discriminator is to accurately identify real data, while the second component reflects the discriminator's ability to correctly recognize generated data as fake. Conversely, the generator's objective is to minimize $V(D, G)$ by producing data that makes the second term larger, aiming to trick the discriminator into believing the generated data is genuine.

The formula (3.3) shows the generator tries to minimize $V(G)$. This means that the generator tries to generate realistic data $G(z)$ so that the discriminator cannot distinguish them from real data.

In real training process, the formula (3.3) will replace by (3.4):

$$\min_G V(D, G) = \mathbb{E}_{z \sim p_z(z)} [-\log(D(G(z)))] \quad (3.4)$$

The reason for replacing $\log(1 - D(G(z)))$ with $-\log(D(G(z)))$ in practice is to address the vanishing gradient problem. As seen in the figure 3.1, $\log(1 - D)$ (blue curve) rapidly decays towards negative infinity when D approaches 1, resulting in very small gradients for the generator when the discriminator is strong. This significantly slows down the generator's learning process. In contrast, $-\log(D)$ (red curve) provides larger gradients even when D is close to 0, allowing the generator to continue learning effectively even against a strong discriminator. Therefore, the modified formula ensures that the generator receives significant gradient updates and avoids stagnation during training.

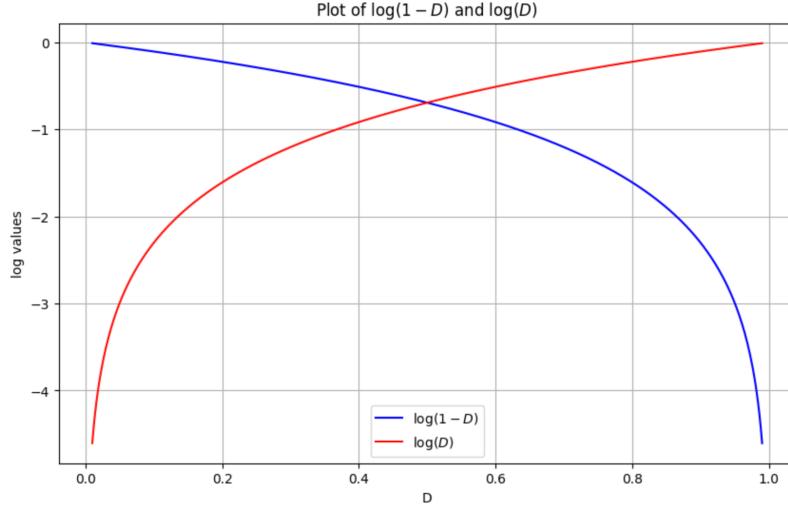


Figure 3.1: Comparison of $\log(1 - D)$ and $-\log(D)$.

The GAN network parameters play a crucial role in determining the quality and diversity of the generated samples $G(x)$ [23]. The optimization process in GANs is a min-max game. The generator aims to minimize the loss function, attempting to generate data that can fool the discriminator, while the discriminator aims to maximize its ability to correctly classify real and generated data. This min-max dynamic is critical for balancing the training of both networks. [24]. Maintaining this balance during training is essential for achieving high-quality sample generation [23]. The D 's role is to provide feedback to the generator by acting as a critic, guiding the G towards producing more realistic samples [25].

3.2 Training Process of GAN

The training process of Generative Adversarial Networks (GANs) is inherently adversarial. The generator endeavors to create increasingly realistic fake samples to deceive the discriminator, while the discriminator seeks to more accurately distinguish between real and fake samples. To understand how this process works, in the following sections, I will describe the GAN training process from the perspectives of data distribution and mathematical formulation.

3.2.1 Distribution Changes During GAN Training

A simaple diagram show how distribution change in GAN training.

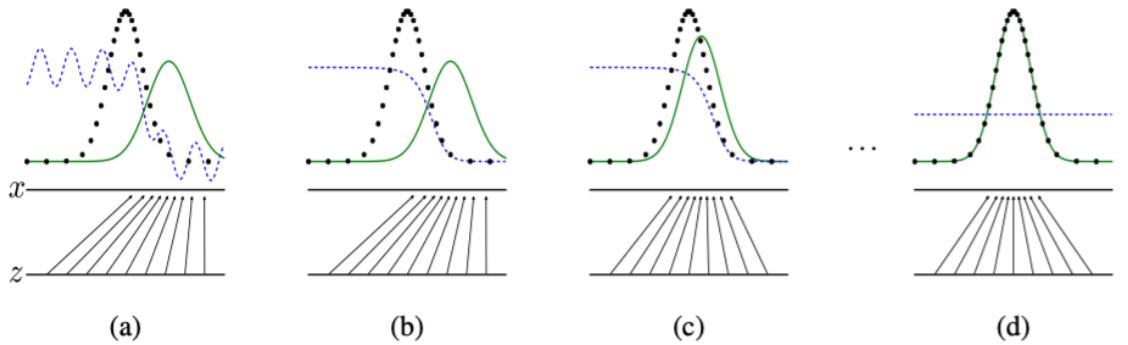


Figure 3.2: Diagram of GAN’s training process. The green curve represents the distribution of generated samples. Initially, the generated samples may differ significantly from the real samples. As training progresses, the generated samples’ distribution gradually approaches the real samples. The black dots represent the distribution of real samples, which remains unchanged throughout the training process and represents the target distribution. The blue dashed line represents the discriminator’s output probability distribution. At the beginning of training, the discriminator can easily distinguish between real and generated data, resulting in a strong classification boundary. As training progresses and the generated data becomes more realistic, the discriminator’s ability to differentiate between the two distributions weakens. Eventually, the discriminator’s output approaches 0.5, indicating it can no longer effectively distinguish between real and generated data. The lines labeled x and z below represent the distribution of samples in the latent space. During GAN training, samples from the latent space z are mapped to the data space x through the generator.

Source: [26]

During the training process of the GAN model, the generator and the discriminator fight against each other. The generator tries to produce realistic samples to fool the discriminator, while the discriminator strives to distinguish real samples from generated samples. Initially, the generator generates samples with poor quality and large errors, similar to the blue sine wave in Figure (a). However, as training proceeds and the generator continues to improve, the error in generated samples gradually decreases, similar

to the blue lines in figures (b) and (c) becoming flat. At the same time, the distribution of real samples and generated samples gradually becomes consistent, and finally reaches the optimal state in Figure (d). At this moment $p_{\text{data}}(x) = p_g(x)$, the samples generated by the generator are almost the same as the real samples, and the error is minimized. During the entire process, the sample distribution gradually converges from the initial noise and discrete state to a state that highly coincides with the real distribution, reflecting a significant improvement in the quality of the generated samples.

3.2.2 Mathematical Formulation Changes During GAN Training

1. Problem Setup

In a Generative Adversarial Network (GAN), the objective is to train a generator G and a discriminator D to generate data that resembles the real data distribution. The objective function to be maximized is given by:

$$V(D, G) = \mathbb{E}_{x \sim p_{\text{data}}(x)}[\log D(x)] + \mathbb{E}_{z \sim p_z(z)}[\log(1 - D(G(z)))] \quad (3.5)$$

2. Rewriting the Objective Function

First, the objective function is rewritten in integral form:

$$V(D, G) = \int p_{\text{data}}(x) \log D(x) dx + \int p_z(z) \log(1 - D(G(z))) dz. \quad (3.6)$$

By changing variables $x' = G(z)$, the generated data is represented as x' , which corresponds to samples produced by the generator G . This allows the second term to be rewritten as an integral over the generated data distribution $p_g(x)$. The integral now reflects the contribution of the generated samples to the objective function.

$$\int p_g(x) \log(1 - D(x)) dx. \quad (3.7)$$

Thus, the objective function becomes:

$$V(D, G) = \int [p_{\text{data}}(x) \log D(x) + p_g(x) \log(1 - D(x))] dx. \quad (3.8)$$

3. Deriving the Optimal Discriminator

To find the optimal discriminator D^* , it needs to take the derivative of the objective function with respect to $D(x)$ and set it to zero.

Let:

$$f(D(x)) = p_{\text{data}}(x) \log D(x) + p_g(x) \log(1 - D(x)). \quad (3.9)$$

Taking the derivative with respect to $D(x)$:

$$\frac{d}{dD(x)} f(D(x)) = \frac{p_{\text{data}}(x)}{D(x)} - \frac{p_g(x)}{1 - D(x)}. \quad (3.10)$$

Setting the derivative to zero:

$$\frac{p_{\text{data}}(x)}{D(x)} = \frac{p_g(x)}{1 - D(x)} \quad (3.11)$$

Solving equation (3.11):

$$D(x) = \frac{p_{\text{data}}(x)}{p_{\text{data}}(x) + p_g(x)}. \quad (3.12)$$

4. Optimal Discriminator Formula

Therefore, the optimal discriminator D^* is given by:

$$D^*(x) = \frac{p_{\text{data}}(x)}{p_{\text{data}}(x) + p_g(x)}. \quad (3.13)$$

- When $p_{\text{data}}(x)$ is much larger than $p_g(x)$, $D^*(x) \approx 1$, indicating that the data point is almost certainly from the real data.
- When $p_{\text{data}}(x)$ is much smaller than $p_g(x)$, $D^*(x) \approx 0$, indicating that the data point is almost certainly from the generated data.

- When $p_{\text{data}}(x)$ is close to $p_g(x)$, $D^*(x) \approx 0.5$, indicating that the discriminator cannot confidently determine whether the data point is real or generated, giving each a 50% probability.

5. Verifying the Optimal Discriminator

To verify that this D^* maximizes the objective function, substitute D^* back into the objective function:

$$V(D^*, G) = \int \left[p_{\text{data}}(x) \log \left(\frac{p_{\text{data}}(x)}{p_{\text{data}}(x) + p_g(x)} \right) + p_g(x) \log \left(1 - \frac{p_{\text{data}}(x)}{p_{\text{data}}(x) + p_g(x)} \right) \right] dx. \quad (3.14)$$

Since:

$$1 - D^*(x) = 1 - \frac{p_{\text{data}}(x)}{p_{\text{data}}(x) + p_g(x)} = \frac{p_g(x)}{p_{\text{data}}(x) + p_g(x)}, \quad (3.15)$$

substituting this in:

$$V(D^*, G) = \int \left[p_{\text{data}}(x) \log \left(\frac{p_{\text{data}}(x)}{p_{\text{data}}(x) + p_g(x)} \right) + p_g(x) \log \left(\frac{p_g(x)}{p_{\text{data}}(x) + p_g(x)} \right) \right] dx. \quad (3.16)$$

This objective function represents the negative of the cross-entropy, which is maximized when $D(x) = D^*(x)$. When $D(x) = 0.5$, the discriminator cannot distinguish between real and generated data, indicating that the generator has produced samples that closely resemble the real data. Maximizing the negative cross-entropy aligns the generated data distribution with the real data distribution.

6. Conclusion

Through the above derivation, it has shown that given the generator G , the optimal form of the discriminator D is:

$$D^*(x) = \frac{p_{\text{data}}(x)}{p_{\text{data}}(x) + p_g(x)}. \quad (3.17)$$

This demonstrates that the optimal discriminator D^* outputs the probability that the input data comes from the real data distribution. This formula

provides a theoretical foundation for training GANs, guiding the updates to the generator G so that its generated data gradually approaches the real data distribution.

3.3 Evaluating GAN Performance

Fréchet Inception Distance (FID) is a metric commonly used in Generative Adversarial Network (GAN) models to quantify the dissimilarity between two image distributions [27]. It measures the distance between the distributions of real images and generated images, providing a numerical assessment of the quality of generated images. FID has gained prominence in evaluating the performance of GANs due to its ability to capture both the quality and diversity of generated images [28]. The following is the objective function for FID:

$$\text{FID} = \|\mu_r - \mu_g\|^2 + \text{Tr}(\Sigma_r + \Sigma_g - 2(\Sigma_r \Sigma_g)^{1/2}) \quad (3.18)$$

$$\mu_r = \frac{1}{N} \sum_{i=1}^N f(x_i), \quad \Sigma_r = \frac{1}{N} \sum_{i=1}^N (f(x_i) - \mu_r)(f(x_i) - \mu_r)^T \quad (3.19)$$

$$\mu_g = \frac{1}{M} \sum_{i=1}^M f(G(z_i)), \quad \Sigma_g = \frac{1}{M} \sum_{i=1}^M (f(G(z_i)) - \mu_g)(f(G(z_i)) - \mu_g)^T \quad (3.20)$$

- μ_r and μ_g represent the feature means of the real and generated images, respectively.
- Σ_r and Σ_g represent the feature covariance matrices of the real and generated images, respectively.
- Tr represents the trace (the sum of the diagonal elements of the matrix).
- f represents the feature extraction function, which extracts feature vectors from images using the Inception network. These feature vectors are

used to compute the mean and covariance for both real and generated images.

A lower FID value indicates that the distribution of the generated images is closer to that of real images, reflecting higher quality and diversity in the generated images [29]. Specifically, a FID value below 10 is considered to represent very high-quality generated images, while values between 10 and 50 indicate good quality, and values above 50 suggest average or poor quality [29].

3.4 Limitations of Accuracy in Evaluating GANs

Accuracy is not suitable for evaluating GANs because accuracy is an indicator of classification tasks, which is used to measure the prediction accuracy of the model in classification tasks, and cannot measure the quality and diversity of generated data. In generation tasks, there is no clear "correct answer" and the generated data has no "real label", so it is impossible to directly compare the correspondence between the generated data and a real sample.

The following is a result for a standard GANs model with high accuracy but generate low quality images.

```
1900 [D loss: 9.02800047697383e-06 | D accuracy: 100.0] [G loss: 0.0005307616665959358] [FID: -1.49196970922936e+87]
2/2 [=====] - 0s 28ms/step
1901 [D loss: 7.511995590903098e-06 | D accuracy: 100.0] [G loss: 0.0006589822005480528] [Epoch time: 0.51 seconds]
2/2 [=====] - 0s 29ms/step
1902 [D loss: 9.099765065911924e-06 | D accuracy: 100.0] [G loss: 0.0008528590551577508] [Epoch time: 0.48 seconds]
2/2 [=====] - 0s 29ms/step
1903 [D loss: 7.715634183114162e-06 | D accuracy: 100.0] [G loss: 0.0008156942203640938] [Epoch time: 0.48 seconds]
2/2 [=====] - 0s 29ms/step
1904 [D loss: 7.386817742371932e-06 | D accuracy: 100.0] [G loss: 0.0005978870904073119] [Epoch time: 0.49 seconds]
2/2 [=====] - 0s 29ms/step
1905 [D loss: 1.5066923879203387e-05 | D accuracy: 100.0] [G loss: 0.0014342099893838167] [Epoch time: 0.48 seconds]
2/2 [=====] - 0s 29ms/step
1906 [D loss: 1.9851730939990375e-05 | D accuracy: 100.0] [G loss: 0.0007973920437507331] [Epoch time: 0.49 seconds]
2/2 [=====] - 0s 29ms/step
1907 [D loss: 1.692915657258709e-05 | D accuracy: 100.0] [G loss: 0.0009673223830759525] [Epoch time: 0.49 seconds]
2/2 [=====] - 0s 29ms/step
1908 [D loss: 1.5433080079874255e-05 | D accuracy: 100.0] [G loss: 0.0010331417433917522] [Epoch time: 0.49 seconds]
2/2 [=====] - 0s 29ms/step
1909 [D loss: 3.74487547105673e-06 | D accuracy: 100.0] [G loss: 0.0008501751581206918] [Epoch time: 0.52 seconds]
```

Figure 3.3: GAN Training Accuracy Over Epochs

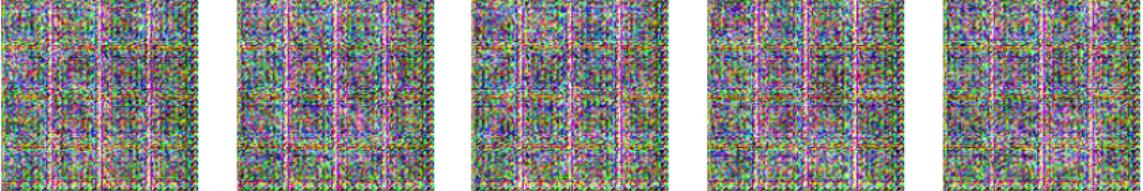
```
▶ generate_images(generator, 10, 100)
⇒ 1/1 [=====] - 1s 844ms/step

```

Figure 3.4: Generated Images from GAN

FID (Fréchet Inception Distance) quantifies the difference between generated data and real data by comparing their distribution in the Inception network feature space. FID takes into account the overall distribution of generated data and real data, and can reflect the quality and diversity of generated data. A low FID value indicates that the distribution of generated data is very close to the distribution of real data, that is, the generated data is both realistic and diverse, so FID is a more suitable indicator for evaluating the performance of generative models.

Chapter 4

Results

This chapter details the numerical processes undertaken to evaluate the performance of various GAN architectures. It includes model selection, structural adjustments, exploration of layer depth, and the impact of data augmentation on model performance. Additionally, the chapter discusses the application of the chosen model to a new dataset, providing insights into the practical effectiveness of the models in generating high-quality images.

4.1 Standard GAN Versus Other GAN Realizations

Approximately a decade has passed since Goodfellow introduced Generative Adversarial Networks (GANs), during which numerous variants of GAN models have been developed. For the purpose of this study, I have selected seven distinct GAN models for examination and minimal implementation: Standard GANs, Conditional GANs, Auxiliary Classifier GANs, Cycle GANs, Domain Transfer Network GANs, Coupled GANs, and Style GANs. Based on their foundational nature, extensive research and documentation, training and implementation efficiency, and flexibility and versatility, Standard GANs were selected for this study. The following outlines the reasons for this choice.

1. Foundational Nature: Standard GANs, introduced by Ian Goodfel-

low et al. in 2014, serve as the foundational model for all subsequent GAN variants. Understanding the principles and mechanics of Standard GANs is crucial for comprehending more complex versions like Conditional GANs, Cycle GANs, and Style GANs. By focusing on the Standard GAN, this study lays a solid foundation for exploring more advanced models.

2. Widely Studied and Well-Documented: Standard GANs have been extensively researched, with a vast amount of literature available. This wealth of resources provides a robust theoretical background and a variety of implementation strategies, facilitating a more thorough and well-supported analysis. This also means that there is ample precedent for common challenges and solutions, making it easier to troubleshoot and refine the model during the study.
3. Training and Implementation Efficiency: Compared to more complex GAN variants, Standard GANs typically require less computational power and shorter training times, making them more accessible for experimentation and analysis. This efficiency allows for multiple experiments and parameter tuning within the constraints of the study, leading to more reliable and reproducible results.
4. Flexibility and Versatility: Standard GANs are highly versatile and can be adapted to a wide range of tasks and datasets. This flexibility makes them an excellent choice for a detailed study that may involve exploring various applications or extending the model to new domains.

4.2 GAN With Convolutional or Dense layers

On the MNIST dataset for 3000 epochs, I compared the generated images from two GAN architectures: one using dense layers and the other using convolutional layers (CNN). It was observed that the GAN with the CNN

architecture outperformed the dense layer architecture in terms of image quality, producing clearer and more realistic images.

As shown in Figures 4.1 and 4.2, the generator and discriminator architectures for both dense and convolutional layers are compared. Figure 4.1a illustrates the generator with dense layers, which consists of a series of fully connected layers, while Figure 4.1b shows the generator with convolutional layers, where convolutional and transpose convolutional layers are used to capture spatial features more effectively. The same comparison applies to the discriminator architectures shown in Figure 4.2. The dense-layer discriminator (Figure 4.2a) is built with fully connected layers, whereas the convolutional discriminator (Figure 4.2b) leverages convolutional layers to better identify patterns in the data.

| Layer (type) | Output Shape | Param # |
|--|-------------------|------------|
| dense_30 (Dense) | (None, 256) | 25856 |
| leaky_re_lu_22 (LeakyReLU) | (None, 256) | 0 |
| batch_normalization_18 (Ba tchNormalization) | (None, 256) | 1024 |
| dense_31 (Dense) | (None, 512) | 131584 |
| leaky_re_lu_23 (LeakyReLU) | (None, 512) | 0 |
| batch_normalization_19 (Ba tchNormalization) | (None, 512) | 2048 |
| dense_32 (Dense) | (None, 1024) | 525312 |
| leaky_re_lu_24 (LeakyReLU) | (None, 1024) | 0 |
| batch_normalization_20 (Ba tchNormalization) | (None, 1024) | 4096 |
| dense_33 (Dense) | (None, 784) | 803600 |
| reshape_6 (Reshape) | (None, 28, 28, 1) | 0 |
| <hr/> | | |
| Total params: | 1493520 | (5.78 MB) |
| Trainable params: | 1489936 | (5.68 MB) |
| Non-trainable params: | 3584 | (14.00 KB) |

(a) Generator with dense layer

| Layer (type) | Output Shape | Param # |
|--|---------------------|-----------|
| dense_38 (Dense) | (None, 6272) | 633472 |
| leaky_re_lu_28 (LeakyReLU) | (None, 6272) | 0 |
| reshape_8 (Reshape) | (None, 7, 7, 128) | 0 |
| batch_normalization_24 (Ba tchNormalization) | (None, 7, 7, 128) | 512 |
| conv2d_transpose (Conv2DTr anspose) | (None, 14, 14, 128) | 262272 |
| leaky_re_lu_29 (LeakyReLU) | (None, 14, 14, 128) | 0 |
| batch_normalization_25 (Ba tchNormalization) | (None, 14, 14, 128) | 512 |
| conv2d_transpose_1 (Conv2D Transpose) | (None, 28, 28, 64) | 131136 |
| leaky_re_lu_30 (LeakyReLU) | (None, 28, 28, 64) | 0 |
| batch_normalization_26 (Ba tchNormalization) | (None, 28, 28, 64) | 256 |
| conv2d (Conv2D) | (None, 28, 28, 1) | 3137 |
| <hr/> | | |
| Total params: | 1031297 | (3.93 MB) |
| Trainable params: | 1030657 | (3.93 MB) |
| Non-trainable params: | 640 | (2.50 KB) |

(b) Generator with convolution layer

Figure 4.1: Generator Architecture with Dense and Convolutional Layers

| Layer (type) | Output Shape | Param # |
|----------------------------|--------------|---------|
| flatten_1 (Flatten) | (None, 784) | 0 |
| dense_23 (Dense) | (None, 512) | 401920 |
| leaky_re_lu_17 (LeakyReLU) | (None, 512) | 0 |
| dropout_2 (Dropout) | (None, 512) | 0 |
| dense_24 (Dense) | (None, 256) | 131328 |
| leaky_re_lu_18 (LeakyReLU) | (None, 256) | 0 |
| dropout_3 (Dropout) | (None, 256) | 0 |
| dense_25 (Dense) | (None, 1) | 257 |

Total params: 533505 (2.04 MB)
Trainable params: 533505 (2.04 MB)
Non-trainable params: 0 (0.00 Byte)

(a) Discriminator with dense layer

| Layer (type) | Output Shape | Param # |
|----------------------------|--------------------|---------|
| conv2d_2 (Conv2D) | (None, 14, 14, 64) | 640 |
| leaky_re_lu_39 (LeakyReLU) | (None, 14, 14, 64) | 0 |
| dropout_6 (Dropout) | (None, 14, 14, 64) | 0 |
| conv2d_3 (Conv2D) | (None, 7, 7, 128) | 73856 |
| leaky_re_lu_40 (LeakyReLU) | (None, 7, 7, 128) | 0 |
| dropout_7 (Dropout) | (None, 7, 7, 128) | 0 |
| flatten_3 (Flatten) | (None, 6272) | 0 |
| dense_47 (Dense) | (None, 1) | 6273 |

Total params: 80769 (315.50 KB)
Trainable params: 80769 (315.50 KB)
Non-trainable params: 0 (0.00 Byte)

(b) Discriminator with convolution layer

Figure 4.2: Discriminator Architecture with Dense and Convolutional Layers

Notably, the dense-layer models contain significantly more parameters compared to the convolutional models. The generator with dense layers has approximately 1.5 million parameters, whereas the convolutional generator has only about 1 million parameters. Similarly, the dense-layer discriminator has over 500,000 parameters, while the convolutional discriminator has only 80,000. Despite the larger parameter size, the dense-layer architectures were less effective in producing high-quality images, highlighting the advantage of convolutional layers in capturing spatial dependencies in the data.

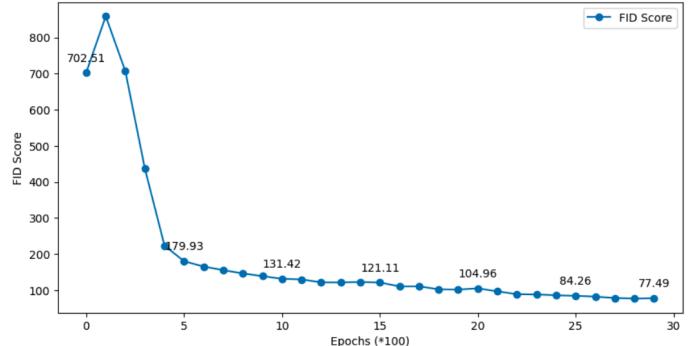


(a) Images generated by GAN with dense layers

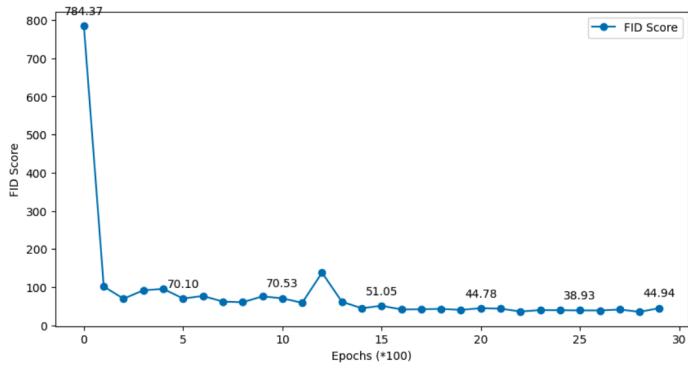


(b) Images generated by GAN with convolution layers

Figure 4.3: Comparison of GAN Performance at 3000 Epochs



(a) FID score for dense layer



(b) FID score for convolutional layer

Figure 4.4: FID Scores Across 3000 Epochs

Based on the results, the GAN model with convolutional layers demonstrated superior performance in generating high-quality images compared to the dense layer model. As shown in the generated images (Figure 4.3), the convolutional GAN produced more coherent and realistic outputs, whereas the dense layer GAN struggled to generate clear and recognizable digits. Additionally, the FID scores, which measure the dissimilarity between generated and real images, further validate this observation. The FID scores for the convolutional GAN (Figure 4.4b) are consistently lower across training epochs compared to the dense layer model (Figure 4.4a), indicating better quality in the generated images.

Given these results, I selected the standard GAN model with convolutional layers for further experiments, as it not only generated higher-quality

images but also used fewer parameters, making it a more efficient and effective choice.

4.3 Exploring Layer Depth

In this section, I explored the impact of changing convolutional layers in the generator and discriminator of a GAN.

Firstly, I trained the basic convolutional GAN model three times and recorded the FID scores. Secondly, based on the basic structure, I gradually added convolutional layers to the generator, from one to three layers, while keeping the discriminator structure fixed. I trained the model three times for each configuration and recorded the FID scores.

Finally, using the base GAN with three additional convolutional layers in the generator, I gradually added convolutional layers to the discriminator, from one to three layers, while keeping the generator structure fixed. I trained the model three times for each configuration and recorded the FID scores.

Upon comparing, I found that increasing the number of layers in either the generator or the discriminator alone worsened the model's performance. This imbalance disrupts the dynamic equilibrium between the generator and discriminator in the GAN framework. However, simultaneously increasing the layers in both the generator and the discriminator improved the model's performance. This finding aligns with the concept of maintaining a balance between the generator and discriminator during training, as highlighted in the literature [30]. The importance of this balance is crucial for the effective operation of GANs, ensuring stable learning and improved image quality [31].

Table 4.1: Average FID Scores for Different GAN Architectures (Lower is Better)

| | | Layers in D | 3 | 4 | 5 | 6 |
|---|--|-------------|-------|-------|-------|---|
| | | Layers in G | 3 | 4 | 5 | 6 |
| 3 | | 68.03 | 72.45 | 78.56 | 85.67 | |
| 4 | | 78.95 | 70.89 | 74.23 | 79.45 | |
| 5 | | 99.17 | 85.36 | 37.09 | 65.54 | |
| 6 | | 194.66 | 71.54 | 55.21 | 35.78 | |

Note: A FID value below 10 is considered to represent very high-quality generated images, while values between 10 and 50 indicate good quality, and values above 50 suggest average or poor quality.



(a) Images generated by basic structure




(b) Images generated by both generator and discriminator are added 3 layers

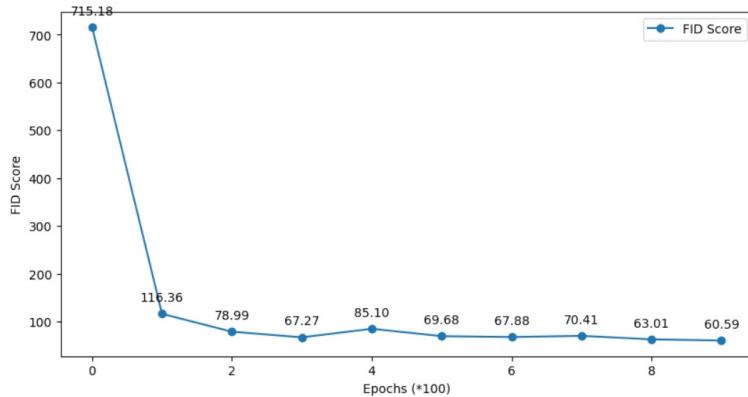
Figure 4.5: Generated Images from Different GAN Architectures

4.4 Impact of Data Augmentation on Model Performance

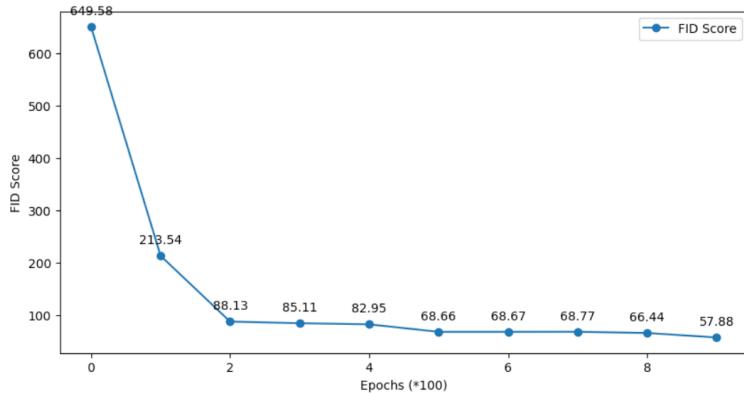
I conducted three training runs of the standard GAN model, both with and without data augmentation, and compared the FID scores. The results indicate that data augmentation generally led to lower optimization performance.



(a) Standard GAN without data augmentation 1

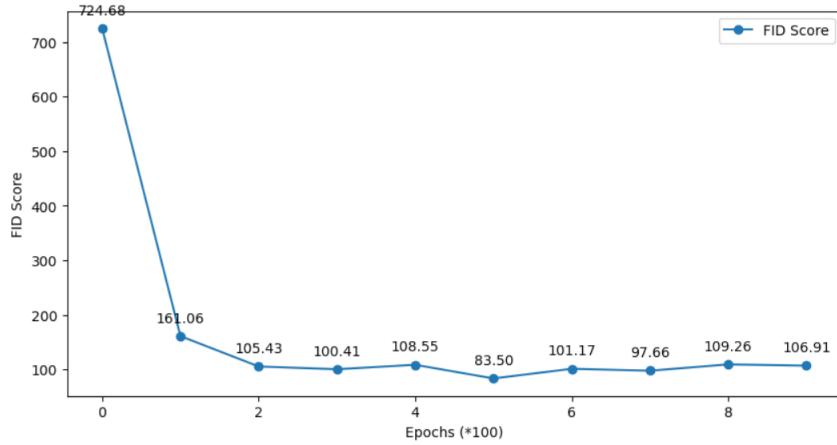


(b) Standard GAN without data augmentation 2




(c) Standard GAN without data augmentation 3


Figure 4.6: The FID scores for standard GAN without data augmentation



(a) Standard GAN with data augmentation 1



(b) Standard GAN with data augmentation 2



(c) Standard GAN with data augmentation 3

Figure 4.7: The FID scores for standard GAN with data augmentation

Typically, a model with data augmentation is expected to perform better than one without it. However, in my experiments, the results were different from what I anticipated.

I believe this discrepancy might be due to the fact that certain data augmentation techniques disturb the distribution of the real dataset, which in turn misleads the generator to train in the wrong direction.

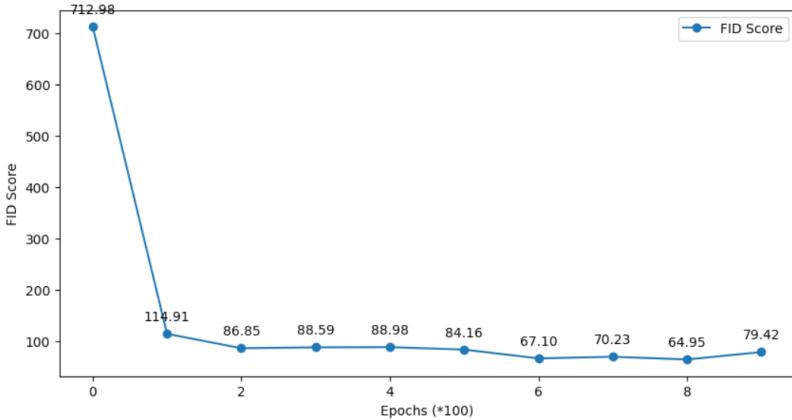
Based on this assumption, I revisited the code and found that I had used four parameters in the data augmentation process: rotation range, width shift range, height shift range, and horizontal flip.

First, I removed the horizontal flip. The results showed that the FID score improved compared to the model with all augmentations but was still worse than the model without any data augmentation.


Next, I removed the rotation augmentation. The FID score dropped slightly compared to the model that only had the horizontal flip removed.

Since width shift and height shift are very similar operations, I did not explore them separately.


In conclusion, the results indicate that data augmentation techniques such as rotation range, width shift range, height shift range, and horizontal flip can disturb the distribution of the real dataset, leading to a decline in model performance. This suggests that while data augmentation typically enhances a model's generalization capabilities, if not carefully chosen or if overused, it can have a detrimental effect on model training.



(a) Standard GAN with data augmentation but remove flip 1

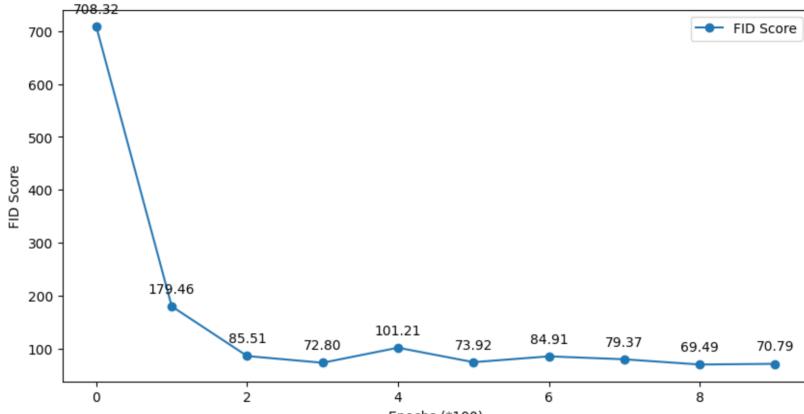


(b) Standard GAN with data augmentation but remove flip 2

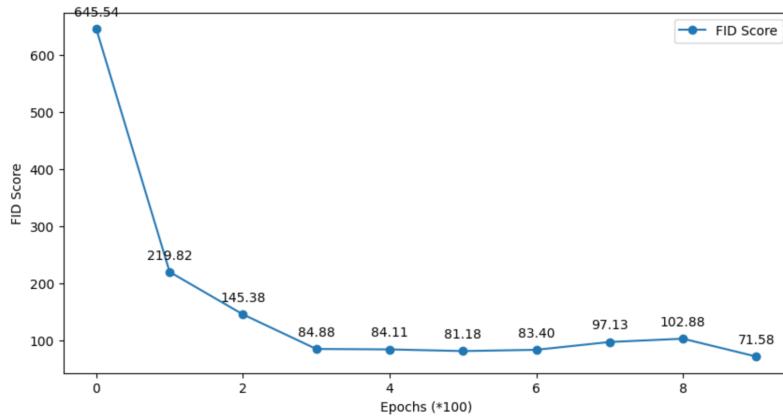


(c) Standard GAN with data augmentation but remove flip 3

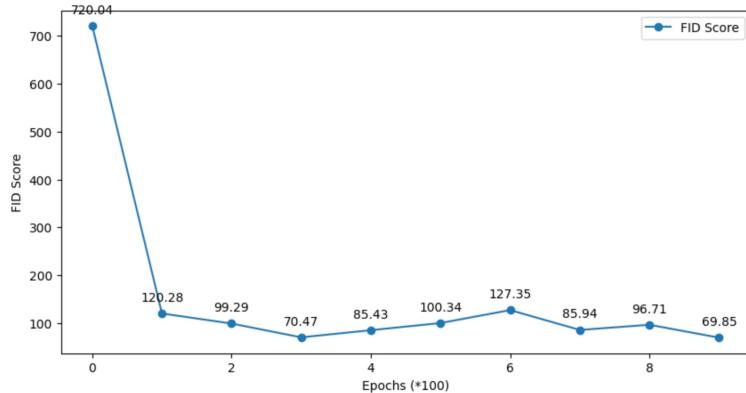
Figure 4.8: The FID scores for standard GAN with data augmentation but remove flip



(a) Standard GAN with data augmentation but remove flip and rotation 1



(b) Standard GAN with data augmentation but remove flip and rotation 2



(c) Standard GAN with data augmentation but remove flip and rotation 3

Figure 4.9: The FID scores for standard GAN with data augmentation but remove flip and rotation

4.5 Applying the Model Animal Faces-HQ Dataset

In this phase, I utilized the Animal Faces-HQ (AFHQ) dataset, which consists of 16,130 high-quality images at a resolution of 512x512 pixels, to train a standard GAN model focused on generating realistic cat images. Due to the high resolution of the original images, which could potentially cause GPU crashes, I downscaled the images to 128x128 pixels. The model was then trained for 4000 epochs. Below are the images generated by the model's generator.



Figure 4.10: Cat Faces Generated by GAN

Chapter 5

Discussion

5.1 Summary of Key Findings

In this study, I conducted a series of experiments to evaluate the performance of different GAN architectures and explored various factors influencing the quality of generated images. Specifically, I compared the effects of using dense layers versus convolutional layers in a standard GAN, examined the impact of varying the number of convolutional layers in the generator and discriminator, analyzed the effects of data augmentation on GAN training, and ultimately applied a GAN model to a new dataset to generate high-quality cat face images.

5.2 Impact of Model Structure

The experiments revealed that GAN models utilizing convolutional layers significantly outperformed those using dense layers in terms of image quality. This finding aligns with the well-established ability of convolutional layers to capture spatial features in images. Further experiments demonstrated that increasing the number of layers in either the generator or the discriminator individually led to decreased model performance. This imbalance disrupts the dynamic equilibrium between the generator and discriminator in the GAN framework, underscoring the importance of maintaining this balance.

during GAN training.

5.3 Data Augmentation Effects

The updated experimental results indicate that data augmentation generally led to lower optimization performance. Although data augmentation typically enhances a model’s generalization ability, in this study, it may have introduced excessive noise, interfering with the GAN training process. This suggests that the careful selection of augmentation methods and parameters is crucial to avoid destabilizing the model and hindering convergence. Therefore, when training GANs, the use of data augmentation should be approached with caution to prevent negative impacts on model performance.

5.4 Performance on New Dataset

When applying the model to the high-resolution Animal Faces-HQ dataset, I successfully trained a standard GAN model to generate realistic cat images. This result validates the effectiveness of standard GANs in handling high-resolution images and demonstrates the potential of GANs in generating realistic images. However, due to hardware constraints, I had to downscale the images from 512x512 pixels to 128x128 pixels, which may have limited the final quality of the generated images.

5.5 Limitations and Future Work

Despite the positive outcomes, this study has some limitations. Firstly, due to limited computational resources, I was unable to train the model on higher-resolution images, which may have affected the quality of the generated images. Secondly, my experiments primarily focused on standard GANs and their simple variants; future research could explore more advanced GAN variants such as StyleGAN or BigGAN. Additionally, in the area of data augmentation, further research is needed to design more effective augmenta-

tion strategies that can improve GAN training efficiency and the quality of generated images.

In future work, I plan to optimize the data augmentation strategies and attempt training on higher-resolution images. Furthermore, exploring different types of GAN models, such as Conditional GANs or Adaptive GANs, to enhance the diversity and quality of generated images could be valuable.

5.6 Conclusion

Overall, this study provides a comprehensive analysis of GAN performance through a series of experiments, demonstrating its potential in generating high-quality images. Despite some challenges, GANs remain a powerful tool in the field of generative models, and future improvements and optimizations are expected to further enhance their applicability.

Appendix A

Code

Listing A.1: GAN model with dense layers

```
1 import tensorflow as tf
2 from tensorflow.keras.layers import Dense, Reshape, Flatten,
3     Dropout, LeakyReLU, BatchNormalization
4 from tensorflow.keras.models import Sequential
5 from tensorflow.keras.optimizers import Adam
6 from scipy.linalg import sqrtm
7 import numpy as np
8 import time
9 import matplotlib.pyplot as plt
10
11 np.random.seed(1000)
12 tf.random.set_seed(1000)
13
14 # input 100
15 # output 28*28*1
16
17 def build_generator():
18     model = Sequential()
19
20     # increase the dimension
21     model.add(Dense(256, input_dim=100))
22     model.add(LeakyReLU(alpha=0.2))
23     model.add(BatchNormalization(momentum=0.8))
```

```

24     model.add(Dense(512))
25     model.add(LeakyReLU(alpha=0.2))
26     model.add(BatchNormalization(momentum=0.8))
27
28     model.add(Dense(1024))
29     model.add(LeakyReLU(alpha=0.2))
30     model.add(BatchNormalization(momentum=0.8))
31
32     model.add(Dense(28*28, activation='tanh'))
33     model.add(Reshape((28, 28, 1)))
34
35     return model
36
37 # input 28*28*1
38 # output 1
39
40 def build_discriminator():
41     model = Sequential()
42
43     model.add(Flatten(input_shape=(28, 28, 1)))
44
45     model.add(Dense(512))
46     model.add(LeakyReLU(alpha=0.2))
47     model.add(Dropout(0.25))
48
49     model.add(Dense(256))
50     model.add(LeakyReLU(alpha=0.2))
51     model.add(Dropout(0.25))
52
53     model.add(Dense(1, activation='sigmoid'))
54
55     return model
56
57 # Load and preprocess the MNIST dataset
58 (x_train, _), (_, _) = tf.keras.datasets.mnist.load_data()
59 x_train = (x_train - 127.5) / 127.5
60 x_train = np.expand_dims(x_train, axis=3)
61
62 # Calculate FID function

```

```

63 def calculate_fid(real_images, fake_images):
64     act1 = real_images.reshape((real_images.shape[0], -1))
65     mu1, sigma1 = act1.mean(axis=0), np.cov(act1, rowvar=
66                               False)
67
68     act2 = fake_images.reshape((fake_images.shape[0], -1))
69     mu2, sigma2 = act2.mean(axis=0), np.cov(act2, rowvar=
70                               False)
71
72     ssdiff = np.sum((mu1 - mu2)**2.0)
73     covmean = sqrtm(sigma1.dot(sigma2))
74
75     if np.iscomplexobj(covmean):
76         covmean = covmean.real
77
78     fid = ssdiff + np.trace(sigma1 + sigma2 - 2.0 * covmean)
79
80     return fid
81
82
83
84 def train_gan(epochs=1000, batch_size=64, p_epoch=100):
85     generator = build_generator()
86     discriminator = build_discriminator()
87
88     discriminator.compile(loss='binary_crossentropy',
89                           optimizer=Adam(0.0002, 0.5), metrics=['accuracy'])
90     discriminator.trainable = False
91
92     gan_input = tf.keras.Input(shape=(100,))
93     gan_output = discriminator(generator(gan_input))
94     gan = tf.keras.Model(gan_input, gan_output)
95     gan.compile(loss='binary_crossentropy', optimizer=Adam
96                  (0.0002, 0.5))
97
98     half_batch = int(batch_size / 2)
99
100    d_losses = []
101    g_losses = []
102    d_acc = []
103    fid_scores = [] # List to store FID scores

```

```

98
99     start_time = time.time() # Record the start time
100
101    for epoch in range(epochs):
102        # Select a random half batch of real images
103        idx = np.random.randint(0, x_train.shape[0],
104                                half_batch)
105        real_images = x_train[idx]
106
107        # Generate a half batch of new fake images
108        noise = np.random.normal(0, 1, (half_batch, 100))
109        fake_images = generator.predict(noise)
110
111        # Train the discriminator
112        real_labels = np.ones((half_batch, 1))
113        fake_labels = np.zeros((half_batch, 1))
114
115        d_loss_real = discriminator.train_on_batch(
116            real_images, real_labels)
117        d_loss_fake = discriminator.train_on_batch(
118            fake_images, fake_labels)
119        d_loss = 0.5 * np.add(d_loss_real, d_loss_fake)
120
121
122        # Train the generator
123        noise = np.random.normal(0, 1, (batch_size, 100))
124        valid_y = np.ones((batch_size, 1))
125
126        g_loss = gan.train_on_batch(noise, valid_y)
127
128
129        # Record the losses
130        d_losses.append(d_loss[0])
131        g_losses.append(g_loss)
132        d_acc.append(d_loss[1] * 100)
133
134
135        # Calculate and print FID every p_epoch epochs
136        if epoch % p_epoch == 0:
137            noise = np.random.normal(0, 1, (1000, 100))
138            fake_images = generator.predict(noise)
139            fid = calculate_fid(x_train[:1000], fake_images)

```

```

134         fid_scores.append(fid)
135         print(f"{{epoch}} [{D_loss}: {{d_loss[0]}}, {acc}: {{100 * d_loss[1]}%}] [{G_loss}: {{g_loss}}] [FID: {{fid}}]")
136     )
137
138     end_time = time.time() # Record the end time
139     total_time = end_time - start_time
140     print(f"Total training time: {{total_time:.2f}} seconds")
141     return generator, d_losses, g_losses, d_acc, fid_scores

```

Listing A.2: GAN model with Convolutional layers

```

1 import tensorflow as tf
2 from tensorflow.keras.layers import Dense, Reshape,
3     Flatten, Dropout, LeakyReLU, Conv2D, Conv2DTranspose,
4     BatchNormalization
5 from tensorflow.keras.models import Sequential
6 from tensorflow.keras.optimizers import Adam
7 from scipy.linalg import sqrtm
8 import numpy as np
9 import time
10 import matplotlib.pyplot as plt
11
12 np.random.seed(1000)
13 tf.random.set_seed(1000)
14
15 # input 100
16 # output 28*28*1
17
18 def build_generator():
19     model = Sequential()
20
21     # increase the dimension
22     model.add(Dense(7*7*128, input_dim=100))
23     model.add(LeakyReLU(alpha=0.2))
24     model.add(Reshape((7, 7, 128)))
25     model.add(BatchNormalization(momentum=0.8))
26
27     model.add(Conv2DTranspose(128, kernel_size=4, strides
28         =2, padding='same'))

```

```

26     model.add(LeakyReLU(alpha=0.2))
27     model.add(BatchNormalization(momentum=0.8))

28
29     model.add(Conv2DTranspose(64, kernel_size=4, strides
30         =2, padding='same'))
31     model.add(LeakyReLU(alpha=0.2))
32     model.add(BatchNormalization(momentum=0.8))

33     model.add(Conv2D(1, kernel_size=7, activation='tanh',
34         padding='same'))

35     return model

36
37 # input 28*28*1
38 # output 1

39
40 def build_discriminator():
41     model = Sequential()

42
43     model.add(Conv2D(64, kernel_size=3, strides=2,
44         input_shape=(28, 28, 1), padding='same'))
45     model.add(LeakyReLU(alpha=0.2))
46     model.add(Dropout(0.25))

47     model.add(Conv2D(128, kernel_size=3, strides=2,
48         padding='same'))
49     model.add(LeakyReLU(alpha=0.2))
50     model.add(Dropout(0.25))

51     model.add(Flatten())
52     model.add(Dense(1, activation='sigmoid'))

53
54     return model

55
56 # Load and preprocess the MNIST dataset
57 (x_train, _), (_, _) = tf.keras.datasets.mnist.load_data()
58
59     x_train = (x_train - 127.5) / 127.5
60     x_train = np.expand_dims(x_train, axis=3)

```

```

60
61     # Calculate FID function
62
63     def calculate_fid(real_images, fake_images):
64         act1 = real_images.reshape((real_images.shape[0], -1))
65             )
66         mu1, sigma1 = act1.mean(axis=0), np.cov(act1, rowvar=
67             False)
68
69         act2 = fake_images.reshape((fake_images.shape[0], -1))
70             )
71         mu2, sigma2 = act2.mean(axis=0), np.cov(act2, rowvar=
72             False)
73
74         ssdiff = np.sum((mu1 - mu2)**2.0)
75         covmean = sqrtm(sigma1.dot(sigma2))
76
77         if np.iscomplexobj(covmean):
78             covmean = covmean.real
79
80         fid = ssdiff + np.trace(sigma1 + sigma2 - 2.0 *
81             covmean)
82
83         return fid
84
85
86     def train_gan(epochs=1000, batch_size=64, p_epoch=100):
87         generator = build_generator()
88         discriminator = build_discriminator()
89
90         discriminator.compile(loss='binary_crossentropy',
91             optimizer=Adam(0.0002, 0.5), metrics=['accuracy'])
92         discriminator.trainable = False
93
94         gan_input = tf.keras.Input(shape=(100,))
95         gan_output = discriminator(generator(gan_input))
96         gan = tf.keras.Model(gan_input, gan_output)
97         gan.compile(loss='binary_crossentropy', optimizer=
98             Adam(0.0002, 0.5))
99
100
101     # using half batches for the discriminator ensures

```

```

92         balanced and efficient training,
# better memory management, and more stable training
93         dynamics in GANs.
94
95         half_batch = int(batch_size / 2)
96
97
98         d_losses = []
99         g_losses = []
100        d_acc = []
101        fid_scores = [] # List to store FID scores
102
103        start_time = time.time() # Record the start time
104
105        for epoch in range(epochs):
106            # Select a random half batch of real images
107            idx = np.random.randint(0, x_train.shape[0],
108                                   half_batch)
109            real_images = x_train[idx]
110
111            # Generate a half batch of new fake images
112            noise = np.random.normal(0, 1, (half_batch, 100))
113            fake_images = generator.predict(noise)
114
115            # Train the discriminator
116            real_labels = np.ones((half_batch, 1))
117            fake_labels = np.zeros((half_batch, 1))
118
119            d_loss_real = discriminator.train_on_batch(
120                real_images, real_labels)
121            d_loss_fake = discriminator.train_on_batch(
122                fake_images, fake_labels)
123            d_loss = 0.5 * np.add(d_loss_real, d_loss_fake)
124
125            # Train the generator
126            noise = np.random.normal(0, 1, (batch_size, 100))
127            valid_y = np.ones((batch_size, 1))
128
129            g_loss = gan.train_on_batch(noise, valid_y)
130
131            # Record the losses

```

```

126     d_losses.append(d_loss[0])
127     g_losses.append(g_loss)
128     d_acc.append(d_loss[1] * 100)
129
130     # Calculate FID every p_epoch epochs
131     if epoch % p_epoch == 0:
132         noise = np.random.normal(0, 1, (1000, 100))
133         fake_images = generator.predict(noise)
134         fid = calculate_fid(x_train[:1000],
135                               fake_images)
136         fid_scores.append(fid)
137         print(f'{epoch}[D_loss:{d_loss[0]}, acc.: {100*d_loss[1]}%][G_loss:{g_loss}][FID:{fid}]')
138
139     end_time = time.time() # Record the end time
140     total_time = end_time - start_time
141     print(f'Total training time: {total_time:.2f} seconds')
142
143     return generator, d_losses, g_losses, d_acc,
144           fid_scores

```

Listing A.3: Explore data augmentation 1

```

1 import tensorflow as tf
2 from tensorflow.keras.layers import Dense, Reshape,
3     Flatten, Dropout, LeakyReLU, Conv2D, Conv2DTranspose,
4     BatchNormalization
5 from tensorflow.keras.models import Sequential
6 from tensorflow.keras.optimizers import Adam
7 from scipy.linalg import sqrtm
8 import numpy as np
9 import time
10 import matplotlib.pyplot as plt
11
12 np.random.seed(1000)
13 tf.random.set_seed(1000)
14
# input 100
# output 28*28*1

```

```

15
16     def build_generator():
17         model = Sequential()
18
19         # increase the dimension
20         model.add(Dense(7*7*128, input_dim=100))
21         model.add(LeakyReLU(alpha=0.2))
22         model.add(Reshape((7, 7, 128)))
23         model.add(BatchNormalization(momentum=0.8))
24
25         model.add(Conv2DTranspose(128, kernel_size=4, strides
26             =2, padding='same'))
27         model.add(LeakyReLU(alpha=0.2))
28         model.add(BatchNormalization(momentum=0.8))
29
30         model.add(Conv2DTranspose(64, kernel_size=4, strides
31             =2, padding='same'))
32         model.add(LeakyReLU(alpha=0.2))
33         model.add(BatchNormalization(momentum=0.8))
34
35         model.add(Conv2D(1, kernel_size=7, activation='tanh',
36             padding='same'))
37
38     return model
39
40     # input 28*28*1
41     # output 1
42
43     def build_discriminator():
44         model = Sequential()
45
46         model.add(Conv2D(64, kernel_size=3, strides=2,
47             input_shape=(28, 28, 1), padding='same'))
48         model.add(LeakyReLU(alpha=0.2))

```

```

49     model.add(Dropout(0.25))
50
51     model.add(Flatten())
52     model.add(Dense(1, activation='sigmoid'))
53
54     return model
55
56 (x_train, _), (_, _) = tf.keras.datasets.mnist.load_data()
57
58 x_train = (x_train - 127.5) / 127.5
59 x_train = np.expand_dims(x_train, axis=3)
60
61 datagen = tf.keras.preprocessing.image.ImageDataGenerator(
62
63     rotation_range=10,
64     width_shift_range=0.1,
65     height_shift_range=0.1,
66     horizontal_flip=True
67 )
68
69 def calculate_fid(real_images, fake_images):
70     act1 = real_images.reshape((real_images.shape[0], -1))
71
72     mu1, sigma1 = act1.mean(axis=0), np.cov(act1, rowvar=False)
73
74     act2 = fake_images.reshape((fake_images.shape[0], -1))
75
76     mu2, sigma2 = act2.mean(axis=0), np.cov(act2, rowvar=False)
77
78     ssdiff = np.sum((mu1 - mu2)**2.0)
79     covmean = sqrtm(sigma1.dot(sigma2))
80
81
82     if np.iscomplexobj(covmean):
83         covmean = covmean.real
84
85     fid = ssdiff + np.trace(sigma1 + sigma2 - 2.0 *
86         covmean)

```

```

81
82     return fid
83
84 def train_gan(epochs=1000, batch_size=64, p_epoch=100):
85     generator = build_generator()
86     discriminator = build_discriminator()
87
88     discriminator.compile(loss='binary_crossentropy',
89                           optimizer=Adam(0.0002, 0.5), metrics=['accuracy'])
90     discriminator.trainable = False
91
92     gan_input = tf.keras.Input(shape=(100,))
93     gan_output = discriminator(generator(gan_input))
94     gan = tf.keras.Model(gan_input, gan_output)
95     gan.compile(loss='binary_crossentropy', optimizer=
96                  Adam(0.0002, 0.5))
97
98     half_batch = int(batch_size / 2)
99
100    d_losses = []
101    g_losses = []
102    d_acc = []
103    fid_scores = []
104
105    start_time = time.time() # Record the start time
106
107    for epoch in range(epochs):
108        # Select a random half batch of real images
109        idx = np.random.randint(0, x_train.shape[0],
110                               half_batch)
111        real_images = x_train[idx]
112
113        real_images_augmented = next(datagen.flow(
114            real_images, batch_size=half_batch))
115
116        # Generate a half batch of new fake images
117        noise = np.random.normal(0, 1, (half_batch, 100))
118        fake_images = generator.predict(noise)

```

```

116     # Train the discriminator
117     real_labels = np.ones((half_batch, 1))
118     fake_labels = np.zeros((half_batch, 1))
119
120     d_loss_real = discriminator.train_on_batch(
121         real_images_augmented, real_labels)
122     d_loss_fake = discriminator.train_on_batch(
123         fake_images, fake_labels)
124     d_loss = 0.5 * np.add(d_loss_real, d_loss_fake)
125
126     # Train the generator
127     noise = np.random.normal(0, 1, (batch_size, 100))
128     valid_y = np.ones((batch_size, 1))
129
130     g_loss = gan.train_on_batch(noise, valid_y)
131
132     # Record the losses
133     d_losses.append(d_loss[0])
134     g_losses.append(g_loss)
135     d_acc.append(d_loss[1] * 100)
136
137     # Calculate FID every p_epoch epochs
138     if epoch % p_epoch == 0:
139         noise = np.random.normal(0, 1, (1000, 100))
140         fake_images = generator.predict(noise)
141         fid = calculate_fid(x_train[:1000],
142                             fake_images)
143         fid_scores.append(fid)
144         print(f"epoch:{d_loss[0]},acc.:{100*d_loss[1]}%[Gloss:{g_loss}][FID:{fid}]")
145
146         end_time = time.time() # Record the end time
147         total_time = end_time - start_time
148         print(f"Total training time:{total_time:.2f}seconds")
149
150         return generator, d_losses, g_losses, d_acc,
151             fid_scores

```

```

148     # Training the GAN with data augmentation and FID
149     # calculation
150     generator, d_losses, g_losses, d_acc, fid_scores =
151         train_gan(epochs=1000, batch_size=64, p_epoch=100)

```

Listing A.4: Explore data augmentation 2

```

1 import tensorflow as tf
2 from tensorflow.keras.layers import Dense, Reshape,
3     Flatten, Dropout, LeakyReLU, Conv2D, Conv2DTranspose,
4     BatchNormalization
5 from tensorflow.keras.models import Sequential
6 from tensorflow.keras.optimizers import Adam
7 from scipy.linalg import sqrtm
8 import numpy as np
9 import time
10 import matplotlib.pyplot as plt
11
12 np.random.seed(1000)
13 tf.random.set_seed(1000)
14
15 # input 100
16 # output 28*28*1
17
18 def build_generator():
19     model = Sequential()
20
21     # increase the dimension
22     model.add(Dense(7*7*128, input_dim=100))
23     model.add(LeakyReLU(alpha=0.2))
24     model.add(Reshape((7, 7, 128)))
25     model.add(BatchNormalization(momentum=0.8))
26
27     model.add(Conv2DTranspose(128, kernel_size=4, strides
28         =2, padding='same'))
29     model.add(LeakyReLU(alpha=0.2))
30     model.add(BatchNormalization(momentum=0.8))
31
32     model.add(Conv2DTranspose(64, kernel_size=4, strides
33         =2, padding='same'))

```

```

30         model.add(LeakyReLU(alpha=0.2))
31         model.add(BatchNormalization(momentum=0.8))

32
33     model.add(Conv2D(1, kernel_size=7, activation='tanh',
34                      padding='same'))

35     return model

36
37 # input 28*28*1
38 # output 1

39
40 def build_discriminator():
41     model = Sequential()

42
43     model.add(Conv2D(64, kernel_size=3, strides=2,
44                      input_shape=(28, 28, 1), padding='same'))
45     model.add(LeakyReLU(alpha=0.2))
46     model.add(Dropout(0.25))

47     model.add(Conv2D(128, kernel_size=3, strides=2,
48                      padding='same'))
49     model.add(LeakyReLU(alpha=0.2))
50     model.add(Dropout(0.25))

51     model.add(Flatten())
52     model.add(Dense(1, activation='sigmoid'))

53
54     return model

55
56 (x_train, _), (_, _) = tf.keras.datasets.mnist.load_data()
57
58 x_train = (x_train - 127.5) / 127.5
59 x_train = np.expand_dims(x_train, axis=3)

60 datagen = tf.keras.preprocessing.image.ImageDataGenerator(
61
62     rotation_range=10,
63     width_shift_range=0.1,
64     height_shift_range=0.1,

```

```

64
65
66     def calculate_fid(real_images, fake_images):
67         act1 = real_images.reshape((real_images.shape[0], -1)
68                                   )
69         mu1, sigma1 = act1.mean(axis=0), np.cov(act1, rowvar=
70                                       False)
71
72         act2 = fake_images.reshape((fake_images.shape[0], -1)
73                                   )
74         mu2, sigma2 = act2.mean(axis=0), np.cov(act2, rowvar=
75                                       False)
76
77         ssdiff = np.sum((mu1 - mu2)**2.0)
78         covmean = sqrtm(sigma1.dot(sigma2))
79
80         if np.iscomplexobj(covmean):
81             covmean = covmean.real
82
83         fid = ssdiff + np.trace(sigma1 + sigma2 - 2.0 *
84                               covmean)
85
86
87         return fid
88
89
90     def train_gan(epochs=1000, batch_size=64, p_epoch=100):
91         generator = build_generator()
92         discriminator = build_discriminator()
93
94         discriminator.compile(loss='binary_crossentropy',
95                               optimizer=Adam(0.0002, 0.5), metrics=['accuracy'])
96         discriminator.trainable = False
97
98         gan_input = tf.keras.Input(shape=(100,))
99         gan_output = discriminator(generator(gan_input))
100        gan = tf.keras.Model(gan_input, gan_output)
101        gan.compile(loss='binary_crossentropy', optimizer=
102                      Adam(0.0002, 0.5))
103
104        half_batch = int(batch_size / 2)

```

```

96
97     d_losses = []
98     g_losses = []
99     d_acc = []
100    fid_scores = []
101
102    start_time = time.time() # Record the start time
103
104    for epoch in range(epochs):
105        # Select a random half batch of real images
106        idx = np.random.randint(0, x_train.shape[0],
107                               half_batch)
108        real_images = x_train[idx]
109
110        real_images_augmented = next(datagen.flow(
111            real_images, batch_size=half_batch))
112
113        # Generate a half batch of new fake images
114        noise = np.random.normal(0, 1, (half_batch, 100))
115        fake_images = generator.predict(noise)
116
117        # Train the discriminator
118        real_labels = np.ones((half_batch, 1))
119        fake_labels = np.zeros((half_batch, 1))
120
121        d_loss_real = discriminator.train_on_batch(
122            real_images_augmented, real_labels)
123        d_loss_fake = discriminator.train_on_batch(
124            fake_images, fake_labels)
125        d_loss = 0.5 * np.add(d_loss_real, d_loss_fake)
126
127        # Train the generator
128        noise = np.random.normal(0, 1, (batch_size, 100))
129        valid_y = np.ones((batch_size, 1))
130
131        g_loss = gan.train_on_batch(noise, valid_y)
132
133        # Record the losses
134        d_losses.append(d_loss[0])

```

```

131         g_losses.append(g_loss)
132         d_acc.append(d_loss[1] * 100)
133
134     # Calculate FID every p_epoch epochs
135     if epoch % p_epoch == 0:
136         noise = np.random.normal(0, 1, (1000, 100))
137         fake_images = generator.predict(noise)
138         fid = calculate_fid(x_train[:1000],
139                             fake_images)
140         fid_scores.append(fid)
141         print(f'{epoch}[Dloss:{d_loss[0]},acc.:'
142               f'{100*d_loss[1]}%][Gloss:{g_loss}]'
143               f'[FID:{fid}]]')
144
145     end_time = time.time() # Record the end time
146     total_time = end_time - start_time
147     print(f'Total training time: {total_time:.2f} seconds')
148
149     return generator, d_losses, g_losses, d_acc,
150           fid_scores
151
152
153 # Training the GAN with data augmentation and FID
154 # calculation
155 generator, d_losses, g_losses, d_acc, fid_scores =
156   train_gan(epochs=1000, batch_size=64, p_epoch=100)

```

Listing A.5: Explore data augmentation 3

```

1 import tensorflow as tf
2 from tensorflow.keras.layers import Dense, Reshape,
3   Flatten, Dropout, LeakyReLU, Conv2D, Conv2DTranspose,
4   BatchNormalization
5 from tensorflow.keras.models import Sequential
6 from tensorflow.keras.optimizers import Adam
7 from scipy.linalg import sqrtm
8 import numpy as np
9 import time
10 import matplotlib.pyplot as plt
11
12 np.random.seed(1000)

```

```

11  tf.random.set_seed(1000)
12
13  # input 100
14  # output 28*28*1
15
16  def build_generator():
17      model = Sequential()
18
19      # increase the dimension
20      model.add(Dense(7*7*128, input_dim=100))
21      model.add(LeakyReLU(alpha=0.2))
22      model.add(Reshape((7, 7, 128)))
23      model.add(BatchNormalization(momentum=0.8))
24
25      model.add(Conv2DTranspose(128, kernel_size=4, strides
26          =2, padding='same'))
27      model.add(LeakyReLU(alpha=0.2))
28      model.add(BatchNormalization(momentum=0.8))
29
30      model.add(Conv2DTranspose(64, kernel_size=4, strides
31          =2, padding='same'))
32      model.add(LeakyReLU(alpha=0.2))
33      model.add(BatchNormalization(momentum=0.8))
34
35      model.add(Conv2D(1, kernel_size=7, activation='tanh',
36          padding='same'))
37
38      return model
39
40  # input 28*28*1
41  # output 1
42
43  def build_discriminator():
44      model = Sequential()
45
46      model.add(Conv2D(64, kernel_size=3, strides=2,
47          input_shape=(28, 28, 1), padding='same'))
48      model.add(LeakyReLU(alpha=0.2))
49      model.add(Dropout(0.25))

```

```

46
47     model.add(Conv2D(128, kernel_size=3, strides=2,
48                     padding='same'))
49     model.add(LeakyReLU(alpha=0.2))
50     model.add(Dropout(0.25))

51     model.add(Flatten())
52     model.add(Dense(1, activation='sigmoid'))

53
54     return model

55

56 (x_train, _), (_, _) = tf.keras.datasets.mnist.load_data()
57
58 x_train = (x_train - 127.5) / 127.5
59 x_train = np.expand_dims(x_train, axis=3)

60 datagen = tf.keras.preprocessing.image.ImageDataGenerator(
61
62     width_shift_range=0.1,
63     height_shift_range=0.1,
64 )
65
66 def calculate_fid(real_images, fake_images):
67     act1 = real_images.reshape((real_images.shape[0], -1))
68
69     mu1, sigma1 = act1.mean(axis=0), np.cov(act1, rowvar=False)

70     act2 = fake_images.reshape((fake_images.shape[0], -1))
71
72     mu2, sigma2 = act2.mean(axis=0), np.cov(act2, rowvar=False)

73     ssdiff = np.sum((mu1 - mu2)**2.0)
74     covmean = sqrtm(sigma1.dot(sigma2))

75     if np.iscomplexobj(covmean):
76         covmean = covmean.real
77

```

```

78         fid = ssdiff + np.trace(sigma1 + sigma2 - 2.0 *
79             covmean)
80
81     return fid
82
83 def train_gan(epochs=1000, batch_size=64, p_epoch=100):
84     generator = build_generator()
85     discriminator = build_discriminator()
86
87     discriminator.compile(loss='binary_crossentropy',
88                             optimizer=Adam(0.0002, 0.5), metrics=['accuracy'])
89     discriminator.trainable = False
90
91     gan_input = tf.keras.Input(shape=(100,))
92     gan_output = discriminator(generator(gan_input))
93     gan = tf.keras.Model(gan_input, gan_output)
94     gan.compile(loss='binary_crossentropy', optimizer=
95                  Adam(0.0002, 0.5))
96
97     half_batch = int(batch_size / 2)
98
99     d_losses = []
100    g_losses = []
101    d_acc = []
102    fid_scores = []
103
104    start_time = time.time() # Record the start time
105
106    for epoch in range(epochs):
107        # Select a random half batch of real images
108        idx = np.random.randint(0, x_train.shape[0],
109                               half_batch)
110        real_images = x_train[idx]
111
112        real_images_augmented = next(datagen.flow(
113            real_images, batch_size=half_batch))
114
115        # Generate a half batch of new fake images
116        noise = np.random.normal(0, 1, (half_batch, 100))

```

```

112     fake_images = generator.predict(noise)
113
114     # Train the discriminator
115     real_labels = np.ones((half_batch, 1))
116     fake_labels = np.zeros((half_batch, 1))
117
118     d_loss_real = discriminator.train_on_batch(
119         real_images_augmented, real_labels)
120     d_loss_fake = discriminator.train_on_batch(
121         fake_images, fake_labels)
122     d_loss = 0.5 * np.add(d_loss_real, d_loss_fake)
123
124     # Train the generator
125     noise = np.random.normal(0, 1, (batch_size, 100))
126     valid_y = np.ones((batch_size, 1))
127
128     g_loss = gan.train_on_batch(noise, valid_y)
129
130     # Record the losses
131     d_losses.append(d_loss[0])
132     g_losses.append(g_loss)
133     d_acc.append(d_loss[1] * 100)
134
135     # Calculate FID every p_epoch epochs
136     if epoch % p_epoch == 0:
137         noise = np.random.normal(0, 1, (1000, 100))
138         fake_images = generator.predict(noise)
139         fid = calculate_fid(x_train[:1000],
140                             fake_images)
141         fid_scores.append(fid)
142         print(f"epoch:[D_loss:{d_loss[0]},acc.:"
143               f"{100*d_loss[1]}%][G_loss:{g_loss}]"
144               f" FID:{fid}]")
145
146         end_time = time.time() # Record the end time
147         total_time = end_time - start_time
148         print(f"Total training time:{total_time:.2f} seconds"
149             ")
150
151     return generator, d_losses, g_losses, d_acc,

```

```
    fid_scores

145
146     # Training the GAN with data augmentation and FID
           calculation
147     generator, d_losses, g_losses, d_acc, fid_scores =
           train_gan(epochs=1000, batch_size=64, p_epoch=100)
```

Listing A.6: Explore GAN with more convolutional layers 1

```
1 import tensorflow as tf
2 from tensorflow.keras.layers import Dense, Reshape,
3     Flatten, Dropout, LeakyReLU, Conv2D, Conv2DTranspose,
4     BatchNormalization
5 from tensorflow.keras.models import Sequential
6 from tensorflow.keras.optimizers import Adam
7 from scipy.linalg import sqrtm
8 import numpy as np
9 import time
10 import matplotlib.pyplot as plt
11
12 np.random.seed(1000)
13 tf.random.set_seed(1000)
14
15 # input 100
16 # output 28*28*1
17
18 def build_generator():
19     model = Sequential()
20
21     model.add(Dense(7*7*128, input_dim=100))
22     model.add(LeakyReLU(alpha=0.2))
23     model.add(Reshape((7, 7, 128)))
24     model.add(BatchNormalization(momentum=0.8))
25
26     # add 1 convolution layer
27     model.add(Conv2D(128, kernel_size=3, strides=1,
28                     padding='same'))
29     model.add(LeakyReLU(alpha=0.2))
30     model.add(BatchNormalization(momentum=0.8))
```

```

29         model.add(Conv2DTranspose(128, kernel_size=4, strides
30             =2, padding='same'))
31         model.add(LeakyReLU(alpha=0.2))
32         model.add(BatchNormalization(momentum=0.8))

33         model.add(Conv2DTranspose(64, kernel_size=4, strides
34             =2, padding='same'))
35         model.add(LeakyReLU(alpha=0.2))
36         model.add(BatchNormalization(momentum=0.8))

37         model.add(Conv2D(1, kernel_size=7, activation='tanh',
38             padding='same'))

39     return model

40
41 # input 28*28*1
42 # output 1

43
44 def build_discriminator():
45     model = Sequential()

46
47     model.add(Conv2D(64, kernel_size=3, strides=2,
48         input_shape=(28, 28, 1), padding='same'))
49     model.add(LeakyReLU(alpha=0.2))
50     model.add(Dropout(0.25))

51     model.add(Conv2D(128, kernel_size=3, strides=2,
52         padding='same'))
53     model.add(LeakyReLU(alpha=0.2))
54     model.add(Dropout(0.25))

55     model.add(Flatten())
56     model.add(Dense(1, activation='sigmoid'))

57
58     return model

59
60 (x_train, _), (_, _) = tf.keras.datasets.mnist.load_data()
61 ()
62 x_train = (x_train - 127.5) / 127.5

```

```

62     x_train = np.expand_dims(x_train, axis=3)
63
64     def calculate_fid(real_images, fake_images):
65         act1 = real_images.reshape((real_images.shape[0], -1))
66             )
67         mu1, sigma1 = act1.mean(axis=0), np.cov(act1, rowvar=
68             False)
69
70         act2 = fake_images.reshape((fake_images.shape[0], -1))
71             )
72         mu2, sigma2 = act2.mean(axis=0), np.cov(act2, rowvar=
73             False)
74
75         ssdiff = np.sum((mu1 - mu2)**2.0)
76         covmean = sqrtm(sigma1.dot(sigma2))
77
78         if np.iscomplexobj(covmean):
79             covmean = covmean.real
80
81         fid = ssdiff + np.trace(sigma1 + sigma2 - 2.0 *
82             covmean)
83
84
85         return fid
86
87
88     def train_gan(epochs=1000, batch_size=64, p_epoch=100):
89         generator = build_generator()
90         discriminator = build_discriminator()
91
92         discriminator.compile(loss='binary_crossentropy',
93             optimizer=Adam(0.0002, 0.5), metrics=['accuracy'])
94         discriminator.trainable = False
95
96         gan_input = tf.keras.Input(shape=(100,))
97         gan_output = discriminator(generator(gan_input))
98         gan = tf.keras.Model(gan_input, gan_output)
99         gan.compile(loss='binary_crossentropy', optimizer=
100             Adam(0.0002, 0.5))
101
102
103         half_batch = int(batch_size / 2)

```

```

94
95     d_losses = []
96     g_losses = []
97     d_acc = []
98     fid_scores = []
99
100    start_time = time.time() # Record the start time
101
102    for epoch in range(epochs):
103        # Select a random half batch of real images
104        idx = np.random.randint(0, x_train.shape[0],
105                               half_batch)
106        real_images = x_train[idx]
107
108        # Generate a half batch of new fake images
109        noise = np.random.normal(0, 1, (half_batch, 100))
110        fake_images = generator.predict(noise)
111
112        # Train the discriminator
113        real_labels = np.ones((half_batch, 1))
114        fake_labels = np.zeros((half_batch, 1))
115
116        d_loss_real = discriminator.train_on_batch(
117            real_images, real_labels)
118        d_loss_fake = discriminator.train_on_batch(
119            fake_images, fake_labels)
120        d_loss = 0.5 * np.add(d_loss_real, d_loss_fake)
121
122        # Train the generator
123        noise = np.random.normal(0, 1, (batch_size, 100))
124        valid_y = np.ones((batch_size, 1))
125
126        g_loss = gan.train_on_batch(noise, valid_y)
127
128        # Record the losses
129        d_losses.append(d_loss[0])
        g_losses.append(g_loss)
        d_acc.append(d_loss[1] * 100)

```

```

130     # Calculate FID every p_epoch epochs
131     if epoch % p_epoch == 0:
132         noise = np.random.normal(0, 1, (1000, 100))
133         fake_images = generator.predict(noise)
134         fid = calculate_fid(x_train[:1000],
135                             fake_images)
136         fid_scores.append(fid)
137         print(f'{epoch}[D_loss:{d_loss[0]}, acc.: {100*d_loss[1]}%][G_loss:{g_loss}][FID:{fid}]')
138
139         end_time = time.time() # Record the end time
140         total_time = end_time - start_time
141         print(f'Total training time:{total_time:.2f} seconds')
142
143         return generator, d_losses, g_losses, d_acc,
144             fid_scores

```

Listing A.7: Explore GAN with more convolutional layers 2

```

1 import tensorflow as tf
2 from tensorflow.keras.layers import Dense, Reshape,
3     Flatten, Dropout, LeakyReLU, Conv2D, Conv2DTranspose,
4     BatchNormalization
5 from tensorflow.keras.models import Sequential
6 from tensorflow.keras.optimizers import Adam
7 from scipy.linalg import sqrtm
8 import numpy as np
9 import time
10 import matplotlib.pyplot as plt
11
12 np.random.seed(1000)
13 tf.random.set_seed(1000)
14
15 # input 100
16 # output 28*28*1
17
18 def build_generator():
19     model = Sequential()

```

```

19     model.add(Dense(7*7*128, input_dim=100))
20     model.add(LeakyReLU(alpha=0.2))
21     model.add(Reshape((7, 7, 128)))
22     model.add(BatchNormalization(momentum=0.8))

23
24     # add the first convolution layer
25     model.add(Conv2D(128, kernel_size=3, strides=1,
26                      padding='same'))
27     model.add(LeakyReLU(alpha=0.2))
28     model.add(BatchNormalization(momentum=0.8))

29
30     # add the 2nd convolution layer
31     model.add(Conv2D(128, kernel_size=3, strides=1,
32                      padding='same'))
33     model.add(LeakyReLU(alpha=0.2))
34     model.add(BatchNormalization(momentum=0.8))

35
36     model.add(Conv2DTranspose(128, kernel_size=4, strides
37                            =2, padding='same'))
38     model.add(LeakyReLU(alpha=0.2))
39     model.add(BatchNormalization(momentum=0.8))

40
41     model.add(Conv2DTranspose(64, kernel_size=4, strides
42                            =2, padding='same'))
43     model.add(LeakyReLU(alpha=0.2))
44     model.add(BatchNormalization(momentum=0.8))

45
46     model.add(Conv2D(1, kernel_size=7, activation='tanh',
47                      padding='same'))

48
49     return model

50
51
52     # input 28*28*1
53     # output 1

54
55     def build_discriminator():
56         model = Sequential()

57
58         model.add(Conv2D(64, kernel_size=3, strides=2,

```

```

    input_shape=(28, 28, 1), padding='same'))
model.add(LeakyReLU(alpha=0.2))
model.add(Dropout(0.25))

56
model.add(Conv2D(128, kernel_size=3, strides=2,
    padding='same'))
model.add(LeakyReLU(alpha=0.2))
model.add(Dropout(0.25))

59
model.add(Flatten())
model.add(Dense(1, activation='sigmoid'))

62
return model

64

65 (x_train, _), (_, _) = tf.keras.datasets.mnist.load_data()
()
66 x_train = (x_train - 127.5) / 127.5
67 x_train = np.expand_dims(x_train, axis=3)

68

69 def calculate_fid(real_images, fake_images):
70     act1 = real_images.reshape((real_images.shape[0], -1)
71         )
72     mu1, sigma1 = act1.mean(axis=0), np.cov(act1, rowvar=
73         False)

74     act2 = fake_images.reshape((fake_images.shape[0], -1)
75         )
76     mu2, sigma2 = act2.mean(axis=0), np.cov(act2, rowvar=
77         False)

78     ssdiff = np.sum((mu1 - mu2)**2.0)
79     covmean = sqrtm(sigma1.dot(sigma2))

80
81     if np.iscomplexobj(covmean):
82         covmean = covmean.real

83     fid = ssdiff + np.trace(sigma1 + sigma2 - 2.0 *
        covmean)

```

```

84     return fid
85
86 def train_gan(epochs=1000, batch_size=64, p_epoch=100):
87     generator = build_generator()
88     discriminator = build_discriminator()
89
90     discriminator.compile(loss='binary_crossentropy',
91                           optimizer=Adam(0.0002, 0.5), metrics=['accuracy'])
92     discriminator.trainable = False
93
94     gan_input = tf.keras.Input(shape=(100,))
95     gan_output = discriminator(generator(gan_input))
96     gan = tf.keras.Model(gan_input, gan_output)
97     gan.compile(loss='binary_crossentropy', optimizer=
98                  Adam(0.0002, 0.5))
99
100    half_batch = int(batch_size / 2)
101
102    d_losses = []
103    g_losses = []
104    d_acc = []
105    fid_scores = []
106
107    start_time = time.time() # Record the start time
108
109    for epoch in range(epochs):
110        # Select a random half batch of real images
111        idx = np.random.randint(0, x_train.shape[0],
112                               half_batch)
113        real_images = x_train[idx]
114
115        # Generate a half batch of new fake images
116        noise = np.random.normal(0, 1, (half_batch, 100))
117        fake_images = generator.predict(noise)
118
119        # Train the discriminator
120        real_labels = np.ones((half_batch, 1))
121        fake_labels = np.zeros((half_batch, 1))

```

```

120     d_loss_real = discriminator.train_on_batch(
121         real_images, real_labels)
122     d_loss_fake = discriminator.train_on_batch(
123         fake_images, fake_labels)
124     d_loss = 0.5 * np.add(d_loss_real, d_loss_fake)
125
126     # Train the generator
127     noise = np.random.normal(0, 1, (batch_size, 100))
128     valid_y = np.ones((batch_size, 1))
129
130     g_loss = gan.train_on_batch(noise, valid_y)
131
132     # Record the losses
133     d_losses.append(d_loss[0])
134     g_losses.append(g_loss)
135     d_acc.append(d_loss[1] * 100)
136
137     # Calculate FID every p_epoch epochs
138     if epoch % p_epoch == 0:
139         noise = np.random.normal(0, 1, (1000, 100))
140         fake_images = generator.predict(noise)
141         fid = calculate_fid(x_train[:1000],
142                             fake_images)
143         fid_scores.append(fid)
144         print(f"Epoch [{d_loss[0]}]{acc.:0{100*d_loss[1]}%}[G loss:{g_loss}][FID:{fid}]")
145
146     end_time = time.time() # Record the end time
147     total_time = end_time - start_time
148     print(f"Total training time: {total_time:.2f} seconds")
149
150     return generator, d_losses, g_losses, d_acc,
151           fid_scores

```

Listing A.8: Explore GAN with more convolutional layers 3

```

1 import tensorflow as tf
2 from tensorflow.keras.layers import Dense, Reshape,
3     Flatten, Dropout, LeakyReLU, Conv2D, Conv2DTranspose,

```

```

    BatchNormalization
3   from tensorflow.keras.models import Sequential
4   from tensorflow.keras.optimizers import Adam
5   from scipy.linalg import sqrtm
6   import numpy as np
7   import time
8   import matplotlib.pyplot as plt
9
10  np.random.seed(1000)
11  tf.random.set_seed(1000)
12
13  # input 100
14  # output 28*28*1
15
16  def build_generator():
17      model = Sequential()
18
19      model.add(Dense(7*7*128, input_dim=100))
20      model.add(LeakyReLU(alpha=0.2))
21      model.add(Reshape((7, 7, 128)))
22      model.add(BatchNormalization(momentum=0.8))
23
24      # add the 1st convolution layer
25      model.add(Conv2D(128, kernel_size=3, strides=1,
26                      padding='same'))
27      model.add(LeakyReLU(alpha=0.2))
28      model.add(BatchNormalization(momentum=0.8))
29
30      # add the 2nd convolution layer
31      model.add(Conv2D(128, kernel_size=3, strides=1,
32                      padding='same'))
33      model.add(LeakyReLU(alpha=0.2))
34      model.add(BatchNormalization(momentum=0.8))
35
36      # add the 3rd convolution layer
37      model.add(Conv2D(128, kernel_size=3, strides=1,
                      padding='same'))
38      model.add(LeakyReLU(alpha=0.2))
39      model.add(BatchNormalization(momentum=0.8))

```

```

38
39     model.add(Conv2DTranspose(128, kernel_size=4, strides
40         =2, padding='same'))
41     model.add(LeakyReLU(alpha=0.2))
42     model.add(BatchNormalization(momentum=0.8))

43
44     model.add(Conv2DTranspose(64, kernel_size=4, strides
45         =2, padding='same'))
46     model.add(LeakyReLU(alpha=0.2))
47     model.add(BatchNormalization(momentum=0.8))

48
49     return model

50

51 # input 28*28*1
52 # output 1

53

54 def build_discriminator():
55     model = Sequential()

56

57     model.add(Conv2D(64, kernel_size=3, strides=2,
58         input_shape=(28, 28, 1), padding='same'))
59     model.add(LeakyReLU(alpha=0.2))
60     model.add(Dropout(0.25))

61
62     model.add(Conv2D(128, kernel_size=3, strides=2,
63         padding='same'))
64     model.add(LeakyReLU(alpha=0.2))
65     model.add(Dropout(0.25))

66
67     model.add(Flatten())
68     model.add(Dense(1, activation='sigmoid'))

69

70     return model

71

72 (x_train, _), (_, _) = tf.keras.datasets.mnist.load_data()

```

```

71     x_train = (x_train - 127.5) / 127.5
72     x_train = np.expand_dims(x_train, axis=3)
73
74     def calculate_fid(real_images, fake_images):
75         act1 = real_images.reshape((real_images.shape[0], -1))
76         mu1, sigma1 = act1.mean(axis=0), np.cov(act1, rowvar=False)
77
78         act2 = fake_images.reshape((fake_images.shape[0], -1))
79         mu2, sigma2 = act2.mean(axis=0), np.cov(act2, rowvar=False)
80
81         ssdiff = np.sum((mu1 - mu2)**2.0)
82         covmean = sqrtm(sigma1.dot(sigma2))
83
84         if np.iscomplexobj(covmean):
85             covmean = covmean.real
86
87         fid = ssdiff + np.trace(sigma1 + sigma2 - 2.0 * covmean)
88
89     return fid
90
91     def train_gan(epochs=1000, batch_size=64, p_epoch=100):
92         generator = build_generator()
93         discriminator = build_discriminator()
94
95         discriminator.compile(loss='binary_crossentropy',
96                               optimizer=Adam(0.0002, 0.5), metrics=['accuracy'])
96         discriminator.trainable = False
97
98         gan_input = tf.keras.Input(shape=(100,))
99         gan_output = discriminator(generator(gan_input))
100        gan = tf.keras.Model(gan_input, gan_output)
101        gan.compile(loss='binary_crossentropy', optimizer=
102                               Adam(0.0002, 0.5))

```

```

103     half_batch = int(batch_size / 2)
104
105     d_losses = []
106     g_losses = []
107     d_acc = []
108     fid_scores = []
109
110     start_time = time.time() # Record the start time
111
112     for epoch in range(epochs):
113         # Select a random half batch of real images
114         idx = np.random.randint(0, x_train.shape[0],
115                               half_batch)
116         real_images = x_train[idx]
117
118         # Generate a half batch of new fake images
119         noise = np.random.normal(0, 1, (half_batch, 100))
120         fake_images = generator.predict(noise)
121
122         # Train the discriminator
123         real_labels = np.ones((half_batch, 1))
124         fake_labels = np.zeros((half_batch, 1))
125
126         d_loss_real = discriminator.train_on_batch(
127             real_images, real_labels)
128         d_loss_fake = discriminator.train_on_batch(
129             fake_images, fake_labels)
130         d_loss = 0.5 * np.add(d_loss_real, d_loss_fake)
131
132
133         # Train the generator
134         noise = np.random.normal(0, 1, (batch_size, 100))
135         valid_y = np.ones((batch_size, 1))
136
137         g_loss = gan.train_on_batch(noise, valid_y)
138
139         # Record the losses
140         d_losses.append(d_loss[0])
141         g_losses.append(g_loss)
142         d_acc.append(d_loss[1] * 100)

```

```

139
140     # Calculate FID every p_epoch epochs
141     if epoch % p_epoch == 0:
142         noise = np.random.normal(0, 1, (1000, 100))
143         fake_images = generator.predict(noise)
144         fid = calculate_fid(x_train[:1000],
145                             fake_images)
146         fid_scores.append(fid)
147         print(f"epoch:{D_loss:[d_loss[0]},acc.:{100*d_loss[1]}%[G_loss:[g_loss]]"
148               FID:{fid})")
149
150         end_time = time.time() # Record the end time
151         total_time = end_time - start_time
152         print(f"Total training time:{total_time:.2f} seconds")
153
154     return generator, d_losses, g_losses, d_acc,
155         fid_scores

```

Listing A.9: Explore GAN with more convolutional layers 4

```

1 import tensorflow as tf
2 from tensorflow.keras.layers import Dense, Reshape,
3     Flatten, Dropout, LeakyReLU, Conv2D, Conv2DTranspose,
4     BatchNormalization
5 from tensorflow.keras.models import Sequential
6 from tensorflow.keras.optimizers import Adam
7 from scipy.linalg import sqrtm
8 import numpy as np
9 import time
10 import matplotlib.pyplot as plt
11
12 np.random.seed(1000)
13 tf.random.set_seed(1000)
14
15 # input 100
16 # output 28*28*1
17
18 def build_generator():
19     model = Sequential()

```

```

18
19     model.add(Dense(7*7*128, input_dim=100))
20     model.add(LeakyReLU(alpha=0.2))
21     model.add(Reshape((7, 7, 128)))
22     model.add(BatchNormalization(momentum=0.8))

23
24     # add the 1st convolution layer
25     model.add(Conv2D(128, kernel_size=3, strides=1,
26                      padding='same'))
27     model.add(LeakyReLU(alpha=0.2))
28     model.add(BatchNormalization(momentum=0.8))

29
30     # add the 2nd convolution layer
31     model.add(Conv2D(128, kernel_size=3, strides=1,
32                      padding='same'))
33     model.add(LeakyReLU(alpha=0.2))
34     model.add(BatchNormalization(momentum=0.8))

35
36     # add the 3rd convolution layer
37     model.add(Conv2D(128, kernel_size=3, strides=1,
38                      padding='same'))
39     model.add(LeakyReLU(alpha=0.2))
40     model.add(BatchNormalization(momentum=0.8))

41
42     model.add(Conv2DTranspose(128, kernel_size=4, strides
43                             =2, padding='same'))
44     model.add(LeakyReLU(alpha=0.2))
45     model.add(BatchNormalization(momentum=0.8))

46
47     model.add(Conv2DTranspose(64, kernel_size=4, strides
48                             =2, padding='same'))
49     model.add(LeakyReLU(alpha=0.2))
50     model.add(BatchNormalization(momentum=0.8))

51
52     model.add(Conv2D(1, kernel_size=7, activation='tanh',
53                      padding='same'))

54
55     return model

```

```

51     # input 28*28*1
52     # output 1
53
54     def build_discriminator():
55         model = Sequential()
56
57         model.add(Conv2D(64, kernel_size=3, strides=2,
58                         input_shape=(28, 28, 1), padding='same'))
59         model.add(LeakyReLU(alpha=0.2))
60         model.add(Dropout(0.25))
61
62         model.add(Conv2D(128, kernel_size=3, strides=2,
63                         padding='same'))
64         model.add(LeakyReLU(alpha=0.2))
65         model.add(Dropout(0.25))
66
67         # add the 1st convolution layer
68         model.add(Conv2D(256, kernel_size=3, strides=2,
69                         padding='same'))
70         model.add(LeakyReLU(alpha=0.2))
71         model.add(Dropout(0.25))
72
73         model.add(Flatten())
74         model.add(Dense(1, activation='sigmoid'))
75
76         return model
77
78
79     (x_train, _), (_, _) = tf.keras.datasets.mnist.load_data()
80     ()
81     x_train = (x_train - 127.5) / 127.5
82     x_train = np.expand_dims(x_train, axis=3)
83
84
85     def calculate_fid(real_images, fake_images):
86         act1 = real_images.reshape((real_images.shape[0], -1))
87
88         mu1, sigma1 = act1.mean(axis=0), np.cov(act1, rowvar=False)
89
90
91         act2 = fake_images.reshape((fake_images.shape[0], -1))

```

```

        )
84     mu2, sigma2 = act2.mean(axis=0), np.cov(act2, rowvar=
85         False)
86
87     ssdiff = np.sum((mu1 - mu2)**2.0)
88     covmean = sqrtm(sigma1.dot(sigma2))
89
90     if np.iscomplexobj(covmean):
91         covmean = covmean.real
92
93     fid = ssdiff + np.trace(sigma1 + sigma2 - 2.0 *
94         covmean)
95
96     return fid
97
98
99
100    def train_gan(epochs=1000, batch_size=64, p_epoch=100):
101        generator = build_generator()
102        discriminator = build_discriminator()
103
104        discriminator.compile(loss='binary_crossentropy',
105            optimizer=Adam(0.0002, 0.5), metrics=['accuracy'])
106        discriminator.trainable = False
107
108        gan_input = tf.keras.Input(shape=(100,))
109        gan_output = discriminator(generator(gan_input))
110        gan = tf.keras.Model(gan_input, gan_output)
111        gan.compile(loss='binary_crossentropy', optimizer=
112            Adam(0.0002, 0.5))
113
114        half_batch = int(batch_size / 2)
115
116        d_losses = []
117        g_losses = []
118        d_acc = []
119        fid_scores = []
120
121        start_time = time.time() # Record the start time
122
123        for epoch in range(epochs):

```

```

118     # Select a random half batch of real images
119     idx = np.random.randint(0, x_train.shape[0],
120                             half_batch)
121     real_images = x_train[idx]
122
123     # Generate a half batch of new fake images
124     noise = np.random.normal(0, 1, (half_batch, 100))
125     fake_images = generator.predict(noise)
126
127     # Train the discriminator
128     real_labels = np.ones((half_batch, 1))
129     fake_labels = np.zeros((half_batch, 1))
130
131     d_loss_real = discriminator.train_on_batch(
132         real_images, real_labels)
133     d_loss_fake = discriminator.train_on_batch(
134         fake_images, fake_labels)
135     d_loss = 0.5 * np.add(d_loss_real, d_loss_fake)
136
137     # Train the generator
138     noise = np.random.normal(0, 1, (batch_size, 100))
139     valid_y = np.ones((batch_size, 1))
140
141     g_loss = gan.train_on_batch(noise, valid_y)
142
143     # Record the losses
144     d_losses.append(d_loss[0])
145     g_losses.append(g_loss)
146     d_acc.append(d_loss[1] * 100)
147
148     # Calculate FID every p_epoch epochs
149     if epoch % p_epoch == 0:
150         noise = np.random.normal(0, 1, (1000, 100))
151         fake_images = generator.predict(noise)
152         fid = calculate_fid(x_train[:1000],
153                             fake_images)
154         fid_scores.append(fid)
155         print(f"epoch:[D_loss:{d_loss[0]},acc.:{100*d_loss[1]}%][G_loss:{g_loss}]")

```

```

152
153     FID:{fid}])")
154
155     end_time = time.time() # Record the end time
156     total_time = end_time - start_time
157     print(f"Total training time:{total_time:.2f} seconds
158 ")
159
160     return generator, d_losses, g_losses, d_acc,
161         fid_scores

```

Listing A.10: Explore GAN with more convolutional layers 5

```

1 import tensorflow as tf
2 from tensorflow.keras.layers import Dense, Reshape,
3     Flatten, Dropout, LeakyReLU, Conv2D, Conv2DTranspose,
4     BatchNormalization
5 from tensorflow.keras.models import Sequential
6 from tensorflow.keras.optimizers import Adam
7 from scipy.linalg import sqrtm
8 import numpy as np
9 import time
10 import matplotlib.pyplot as plt
11
12
13 np.random.seed(1000)
14 tf.random.set_seed(1000)
15
16
17 # input 100
18 # output 28*28*1
19
20
21 def build_generator():
22     model = Sequential()
23
24     model.add(Dense(7*7*128, input_dim=100))
25     model.add(LeakyReLU(alpha=0.2))
26     model.add(Reshape((7, 7, 128)))
27     model.add(BatchNormalization(momentum=0.8))
28
29
30     # add the 1st convolution layer
31     model.add(Conv2D(128, kernel_size=3, strides=1,
32                     padding='same'))
33     model.add(LeakyReLU(alpha=0.2))

```

```

27     model.add(BatchNormalization(momentum=0.8))

28
29     # add the 2nd convolution layer
30     model.add(Conv2D(128, kernel_size=3, strides=1,
31                     padding='same'))
32     model.add(LeakyReLU(alpha=0.2))
33     model.add(BatchNormalization(momentum=0.8))

34
35     # add the 3rd convolution layer
36     model.add(Conv2D(128, kernel_size=3, strides=1,
37                     padding='same'))
38     model.add(LeakyReLU(alpha=0.2))
39     model.add(BatchNormalization(momentum=0.8))

40
41     model.add(Conv2DTranspose(128, kernel_size=4, strides
42                             =2, padding='same'))
43     model.add(LeakyReLU(alpha=0.2))
44     model.add(BatchNormalization(momentum=0.8))

45
46     model.add(Conv2DTranspose(64, kernel_size=4, strides
47                             =2, padding='same'))
48     model.add(LeakyReLU(alpha=0.2))
49     model.add(BatchNormalization(momentum=0.8))

50
51     model.add(Conv2D(1, kernel_size=7, activation='tanh',
52                     padding='same'))

53
54     return model

55
56
57
58
59

```

```

60
61     model.add(Conv2D(128, kernel_size=3, strides=2,
62                     padding='same'))
63     model.add(LeakyReLU(alpha=0.2))
64     model.add(Dropout(0.25))

65
66     model.add(Conv2D(256, kernel_size=3, strides=2,
67                     padding='same'))
68     model.add(LeakyReLU(alpha=0.2))
69     model.add(Dropout(0.25))

70
71     # add the 1st convolution layer
72     model.add(Conv2D(512, kernel_size=3, strides=2,
73                     padding='same'))
74     model.add(LeakyReLU(alpha=0.2))
75     model.add(Dropout(0.25))

76
77     # add the 2nd convolution layer
78     model.add(Conv2D(1024, kernel_size=3, strides=2,
79                     padding='same'))
80     model.add(LeakyReLU(alpha=0.2))
81     model.add(Dropout(0.25))

82
83     model.add(Flatten())
84     model.add(Dense(1, activation='sigmoid'))

85
86     return model

87

88 (x_train, _), (_, _) = tf.keras.datasets.mnist.load_data()
89
90 x_train = (x_train - 127.5) / 127.5
91 x_train = np.expand_dims(x_train, axis=3)

92

93 def calculate_fid(real_images, fake_images):
94     act1 = real_images.reshape((real_images.shape[0], -1))
95
96     mu1, sigma1 = act1.mean(axis=0), np.cov(act1, rowvar=False)

```

```

92     act2 = fake_images.reshape((fake_images.shape[0], -1))
93     )
94     mu2, sigma2 = act2.mean(axis=0), np.cov(act2, rowvar=
95     False)
96
97
98     if np.iscomplexobj(covmean):
99         covmean = covmean.real
100
101    fid = ssdiff + np.trace(sigma1 + sigma2 - 2.0 *
102        covmean)
103
104
105    def train_gan(epochs=1000, batch_size=64, p_epoch=100):
106        generator = build_generator()
107        discriminator = build_discriminator()
108
109        discriminator.compile(loss='binary_crossentropy',
110                               optimizer=Adam(0.0002, 0.5), metrics=['accuracy'])
110        discriminator.trainable = False
111
112        gan_input = tf.keras.Input(shape=(100,))
113        gan_output = discriminator(generator(gan_input))
114        gan = tf.keras.Model(gan_input, gan_output)
115        gan.compile(loss='binary_crossentropy', optimizer=
116                               Adam(0.0002, 0.5))
117
117        half_batch = int(batch_size / 2)
118
119        d_losses = []
120        g_losses = []
121        d_acc = []
122        fid_scores = []
123
124        start_time = time.time() # Record the start time
125

```

```

126     for epoch in range(epochs):
127         # Select a random half batch of real images
128         idx = np.random.randint(0, x_train.shape[0],
129                               half_batch)
130         real_images = x_train[idx]
131
132         # Generate a half batch of new fake images
133         noise = np.random.normal(0, 1, (half_batch, 100))
134         fake_images = generator.predict(noise)
135
136         # Train the discriminator
137         real_labels = np.ones((half_batch, 1))
138         fake_labels = np.zeros((half_batch, 1))
139
140         d_loss_real = discriminator.train_on_batch(
141             real_images, real_labels)
142         d_loss_fake = discriminator.train_on_batch(
143             fake_images, fake_labels)
144         d_loss = 0.5 * np.add(d_loss_real, d_loss_fake)
145
146         # Train the generator
147         noise = np.random.normal(0, 1, (batch_size, 100))
148         valid_y = np.ones((batch_size, 1))
149
150         g_loss = gan.train_on_batch(noise, valid_y)
151
152         # Record the losses
153         d_losses.append(d_loss[0])
154         g_losses.append(g_loss)
155         d_acc.append(d_loss[1] * 100)
156
157         # Calculate FID every p_epoch epochs
158         if epoch % p_epoch == 0:
159             noise = np.random.normal(0, 1, (1000, 100))
160             fake_images = generator.predict(noise)

```

```

161             {100*d_loss[1]}%][Gloss:{g_loss}][
162             FID:{fid}]")
163
164         end_time = time.time() # Record the end time
165         total_time = end_time - start_time
166         print(f"Total training time: {total_time:.2f} seconds
167             ")
168
169         return generator, d_losses, g_losses, d_acc,
170             fid_scores

```

Listing A.11: Explore GAN with more convolutional layers 6

```

1 def build_generator():
2     model = Sequential()
3
4     model.add(Dense(7*7*128, input_dim=100))
5     model.add(LeakyReLU(alpha=0.2))
6     model.add(Reshape((7, 7, 128)))
7     model.add(BatchNormalization(momentum=0.8))
8
9     # add the 1st convolution layer
10    model.add(Conv2D(128, kernel_size=3, strides=1, padding='same'))
11    model.add(LeakyReLU(alpha=0.2))
12    model.add(BatchNormalization(momentum=0.8))
13
14    # add the 2nd convolution layer
15    model.add(Conv2D(128, kernel_size=3, strides=1, padding='same'))
16    model.add(LeakyReLU(alpha=0.2))
17    model.add(BatchNormalization(momentum=0.8))
18
19    # add the 3rd convolution layer
20    model.add(Conv2D(128, kernel_size=3, strides=1, padding='same'))
21    model.add(LeakyReLU(alpha=0.2))
22    model.add(BatchNormalization(momentum=0.8))
23
24    model.add(Conv2DTranspose(128, kernel_size=4, strides=2,
25        padding='same'))

```

```

25     model.add(LeakyReLU(alpha=0.2))
26     model.add(BatchNormalization(momentum=0.8))

27
28     model.add(Conv2DTranspose(64, kernel_size=4, strides=2,
29                             padding='same'))
30     model.add(LeakyReLU(alpha=0.2))
31     model.add(BatchNormalization(momentum=0.8))

32     model.add(Conv2D(1, kernel_size=7, activation='tanh',
33                      padding='same'))

34
35
36 def build_discriminator():
37     model = Sequential()

38
39     model.add(Conv2D(64, kernel_size=3, strides=2,
40                      input_shape=(28, 28, 1), padding='same'))
41     model.add(LeakyReLU(alpha=0.2))
42     model.add(Dropout(0.25))

43     model.add(Conv2D(128, kernel_size=3, strides=2, padding='
44                     same'))
45     model.add(LeakyReLU(alpha=0.2))
46     model.add(Dropout(0.25))

47     model.add(Conv2D(256, kernel_size=3, strides=2, padding='
48                     same'))
49     model.add(LeakyReLU(alpha=0.2))
50     model.add(Dropout(0.25))

51     # add the 1st convolution layer
52     model.add(Conv2D(512, kernel_size=3, strides=2, padding='
53                     same'))
54     model.add(LeakyReLU(alpha=0.2))
55     model.add(Dropout(0.25))

56     # add the 2nd convolution layer
57     model.add(Conv2D(1024, kernel_size=3, strides=2, padding=

```

```

    'same'))
58 model.add(LeakyReLU(alpha=0.2))
59 model.add(Dropout(0.25))

60
61 # add the 3rd convolution layer
62 model.add(Conv2D(2048, kernel_size=3, strides=2, padding=
    'same'))
63 model.add(LeakyReLU(alpha=0.2))
64 model.add(Dropout(0.25))

65
66 model.add(Flatten())
67 model.add(Dense(1, activation='sigmoid'))

68
69 return model

70
71
72
73
74 import tensorflow as tf
75 from tensorflow.keras.layers import Dense, Reshape, Flatten,
    Dropout, LeakyReLU, Conv2D, Conv2DTranspose,
    BatchNormalization
76 from tensorflow.keras.models import Sequential
77 from tensorflow.keras.optimizers import Adam
78 from scipy.linalg import sqrtm
79 import numpy as np
80 import time
81 import matplotlib.pyplot as plt

82
83 np.random.seed(1000)
84 tf.random.set_seed(1000)

85
86 (x_train, _), (_, _) = tf.keras.datasets.mnist.load_data()
87 x_train = (x_train - 127.5) / 127.5
88 x_train = np.expand_dims(x_train, axis=3)

89
90 def calculate_fid(real_images, fake_images):
91     act1 = real_images.reshape((real_images.shape[0], -1))
92     mu1, sigma1 = act1.mean(axis=0), np.cov(act1, rowvar=

```

```

        False)

93
94     act2 = fake_images.reshape((fake_images.shape[0], -1))
95     mu2, sigma2 = act2.mean(axis=0), np.cov(act2, rowvar=
96         False)
97
98     ssdiff = np.sum((mu1 - mu2)**2.0)
99     covmean = sqrtm(sigma1.dot(sigma2))
100
101    if np.iscomplexobj(covmean):
102        covmean = covmean.real
103
104    fid = ssdiff + np.trace(sigma1 + sigma2 - 2.0 * covmean)
105
106    return fid
107
108
109 def train_gan(epochs=1000, batch_size=64, p_epoch=100):
110     generator = build_generator()
111     discriminator = build_discriminator()
112
113     discriminator.compile(loss='binary_crossentropy',
114                           optimizer=Adam(0.0002, 0.5), metrics=['accuracy'])
115     discriminator.trainable = False
116
117     gan_input = tf.keras.Input(shape=(100,))
118     gan_output = discriminator(generator(gan_input))
119     gan = tf.keras.Model(gan_input, gan_output)
120     gan.compile(loss='binary_crossentropy', optimizer=Adam
121                  (0.0002, 0.5))
122
123     half_batch = int(batch_size / 2)
124
125     d_losses = []
126     g_losses = []
127     d_acc = []
128     fid_scores = []
129
130     start_time = time.time() # Record the start time

```

```

128     for epoch in range(epochs):
129         # Select a random half batch of real images
130         idx = np.random.randint(0, x_train.shape[0],
131                               half_batch)
132         real_images = x_train[idx]
133
134         # Generate a half batch of new fake images
135         noise = np.random.normal(0, 1, (half_batch, 100))
136         fake_images = generator.predict(noise)
137
138         # Train the discriminator
139         real_labels = np.ones((half_batch, 1))
140         fake_labels = np.zeros((half_batch, 1))
141
142         d_loss_real = discriminator.train_on_batch(
143             real_images, real_labels)
144         d_loss_fake = discriminator.train_on_batch(
145             fake_images, fake_labels)
146         d_loss = 0.5 * np.add(d_loss_real, d_loss_fake)
147
148
149         # Train the generator
150         noise = np.random.normal(0, 1, (batch_size, 100))
151         valid_y = np.ones((batch_size, 1))
152
153         g_loss = gan.train_on_batch(noise, valid_y)
154
155
156         # Record the losses
157         d_losses.append(d_loss[0])
158         g_losses.append(g_loss)
159         d_acc.append(d_loss[1] * 100)
160
161
162         # Calculate FID every p_epoch epochs
163         if epoch % p_epoch == 0:
164             noise = np.random.normal(0, 1, (1000, 100))
165             fake_images = generator.predict(noise)
166             fid = calculate_fid(x_train[:1000], fake_images)
167             fid_scores.append(fid)
168             print(f"epoch:[D_loss:{d_loss[0]},acc.:{100*d_loss[1]}%][G_loss:{g_loss}][FID:{fid}]")

```

```

        " )

163
164     end_time = time.time()    # Record the end time
165     total_time = end_time - start_time
166     print(f"Total\u00d7training\u00d7time:\u00d7{total_time:.2f}\u00d7seconds")
167     return generator, d_losses, g_losses, d_acc, fid_scores

```

Listing A.12: Apply new dataset

```

1
2     import os
3     import time
4     import numpy as np
5     import tensorflow as tf
6     import matplotlib.pyplot as plt
7     from tensorflow.keras.models import Sequential, Model
8     from tensorflow.keras.layers import Dense, Reshape,
9         BatchNormalization, LeakyReLU, Conv2D, Conv2DTranspose
10        , Flatten, Dropout, Input
11        from tensorflow.keras.optimizers import Adam
12        from tensorflow.keras.preprocessing.image import load_img
13        , img_to_array
14
15
16
17
18
19
20
21
22
23
24
25
26
27

```

```

1
2     import os
3     import time
4     import numpy as np
5     import tensorflow as tf
6     import matplotlib.pyplot as plt
7     from tensorflow.keras.models import Sequential, Model
8     from tensorflow.keras.layers import Dense, Reshape,
9         BatchNormalization, LeakyReLU, Conv2D, Conv2DTranspose
10        , Flatten, Dropout, Input
11        from tensorflow.keras.optimizers import Adam
12        from tensorflow.keras.preprocessing.image import load_img
13        , img_to_array
14
15
16
17
18
19
20
21
22
23
24
25
26
27

```

```

1
2     import os
3     import time
4     import numpy as np
5     import tensorflow as tf
6     import matplotlib.pyplot as plt
7     from tensorflow.keras.models import Sequential, Model
8     from tensorflow.keras.layers import Dense, Reshape,
9         BatchNormalization, LeakyReLU, Conv2D, Conv2DTranspose
10        , Flatten, Dropout, Input
11        from tensorflow.keras.optimizers import Adam
12        from tensorflow.keras.preprocessing.image import load_img
13        , img_to_array
14
15
16
17
18
19
20
21
22
23
24
25
26
27

```

```

1
2     import os
3     import time
4     import numpy as np
5     import tensorflow as tf
6     import matplotlib.pyplot as plt
7     from tensorflow.keras.models import Sequential, Model
8     from tensorflow.keras.layers import Dense, Reshape,
9         BatchNormalization, LeakyReLU, Conv2D, Conv2DTranspose
10        , Flatten, Dropout, Input
11        from tensorflow.keras.optimizers import Adam
12        from tensorflow.keras.preprocessing.image import load_img
13        , img_to_array
14
15
16
17
18
19
20
21
22
23
24
25
26
27

```

```

1
2     import os
3     import time
4     import numpy as np
5     import tensorflow as tf
6     import matplotlib.pyplot as plt
7     from tensorflow.keras.models import Sequential, Model
8     from tensorflow.keras.layers import Dense, Reshape,
9         BatchNormalization, LeakyReLU, Conv2D, Conv2DTranspose
10        , Flatten, Dropout, Input
11        from tensorflow.keras.optimizers import Adam
12        from tensorflow.keras.preprocessing.image import load_img
13        , img_to_array
14
15
16
17
18
19
20
21
22
23
24
25
26
27

```

```

1
2     import os
3     import time
4     import numpy as np
5     import tensorflow as tf
6     import matplotlib.pyplot as plt
7     from tensorflow.keras.models import Sequential, Model
8     from tensorflow.keras.layers import Dense, Reshape,
9         BatchNormalization, LeakyReLU, Conv2D, Conv2DTranspose
10        , Flatten, Dropout, Input
11        from tensorflow.keras.optimizers import Adam
12        from tensorflow.keras.preprocessing.image import load_img
13        , img_to_array
14
15
16
17
18
19
20
21
22
23
24
25
26
27

```

```

1
2     import os
3     import time
4     import numpy as np
5     import tensorflow as tf
6     import matplotlib.pyplot as plt
7     from tensorflow.keras.models import Sequential, Model
8     from tensorflow.keras.layers import Dense, Reshape,
9         BatchNormalization, LeakyReLU, Conv2D, Conv2DTranspose
10        , Flatten, Dropout, Input
11        from tensorflow.keras.optimizers import Adam
12        from tensorflow.keras.preprocessing.image import load_img
13        , img_to_array
14
15
16
17
18
19
20
21
22
23
24
25
26
27

```

```

28     cat_path = "/content/drive/MyDrive/gan/afhq/train/cat"
29     size = (128, 128)
30     x_train = load_images_as_rgb_matrices(cat_path, size)
31
32     # generator
33     # input 100
34     # output 128*128*3
35
36     def build_generator():
37         model = Sequential()
38
39         # Increase the dimension
40         model.add(Dense(16*16*256, input_dim=100))
41         model.add(LeakyReLU(alpha=0.2))
42         model.add(Reshape((16, 16, 256)))
43         model.add(BatchNormalization(momentum=0.8))
44
45         model.add(Conv2DTranspose(256, kernel_size=4, strides
46             =2, padding='same')) # 32x32
47         model.add(LeakyReLU(alpha=0.2))
48         model.add(BatchNormalization(momentum=0.8))
49
50         model.add(Conv2DTranspose(128, kernel_size=4, strides
51             =2, padding='same')) # 64x64
52         model.add(LeakyReLU(alpha=0.2))
53         model.add(BatchNormalization(momentum=0.8))
54
55         model.add(Conv2DTranspose(64, kernel_size=4, strides
56             =2, padding='same')) # 128x128
57         model.add(LeakyReLU(alpha=0.2))
58         model.add(BatchNormalization(momentum=0.8))
59
60         model.add(Conv2D(3, kernel_size=7, activation='tanh',
61             padding='same')) # 128x128x3
62
63     return model
64
65     # discriminator
66     # input 128*128*3

```

```

63 # output 1
64
65 def build_discriminator():
66     model = Sequential()
67
68     model.add(Conv2D(64, kernel_size=3, strides=2,
69                     input_shape=(128, 128, 3), padding='same'))
70     model.add(LeakyReLU(alpha=0.2))
71     model.add(Dropout(0.25))
72
73     model.add(Conv2D(128, kernel_size=3, strides=2,
74                     padding='same'))
75     model.add(LeakyReLU(alpha=0.2))
76     model.add(Dropout(0.25))
77
78     model.add(Conv2D(256, kernel_size=3, strides=2,
79                     padding='same'))
80     model.add(LeakyReLU(alpha=0.2))
81     model.add(Dropout(0.25))
82
83     model.add(Flatten())
84     model.add(Dense(1, activation='sigmoid'))
85
86
87     return model
88
89
90     # initial GAN
91     def build_gan(generator, discriminator):
92         discriminator.trainable = False
93         gan_input = Input(shape=(100,))
94         x = generator(gan_input)
95         gan_output = discriminator(x)
96         gan = Model(gan_input, gan_output)
97         gan.compile(loss='binary_crossentropy', optimizer=tf.
98                     keras.optimizers.legacy.Adam(0.0002, 0.5))

```

```

97     return gan
98
99
100    # define training step
101   def train(generator, discriminator, gan, x_train, epochs,
102             batch_size=128):
103       valid = np.ones((batch_size, 1))
104       fake = np.zeros((batch_size, 1))
105
106
107
108       for epoch in range(epochs):
109           start_time = time.time()
110
111
112       idx = np.random.randint(0, x_train.shape[0],
113                             batch_size)
114       imgs = x_train[idx]
115
116       noise = np.random.normal(0, 1, (batch_size, 100))
117       gen_imgs = generator.predict(noise)
118
119       d_loss_real = discriminator.train_on_batch(imgs,
120                                                   valid)
121       d_loss_fake = discriminator.train_on_batch(
122           gen_imgs, fake)
123       d_loss = 0.5 * np.add(d_loss_real, d_loss_fake)
124
125
126       noise = np.random.normal(0, 1, (batch_size, 100))
127       g_loss = gan.train_on_batch(noise, valid)
128
129
130       end_time = time.time()
131       epoch_time = end_time - start_time
132
133
134       print(f"Epoch: {epoch} [D loss: {d_loss[0]} | D accuracy: {100*d_loss[1]}] [G loss: {g_loss}] [Epoch time: {epoch_time:.2f} seconds]")
135
136
137
138   generator = build_generator()
139   discriminator = build_discriminator()

```

```
130     discriminator.compile(loss='binary_crossentropy',
131                             optimizer=tf.keras.optimizers.legacy.Adam(0.0002, 0.5)
132                             , metrics=['accuracy'])
133     gan = build_gan(generator, discriminator)
134
135     # gen images
136     def show_generated_images(generator, num_images=25, dim
137                               =(5, 5), figsize=(10, 10)):
138         noise = np.random.normal(0, 1, (num_images, 100))
139         gen_imgs = generator.predict(noise)
140         gen_imgs = 0.5 * gen_imgs + 0.5
141
142         plt.figure(figsize=figsize)
143         for i in range(num_images):
144             plt.subplot(dim[0], dim[1], i+1)
145             plt.imshow(gen_imgs[i])
146             plt.axis('off')
147         plt.tight_layout()
148         plt.show()
149
150     show_generated_images(generator)
```

Bibliography

- [1] I. Gulrajani, F. Ahmed, M. Arjovsky, V. Dumoulin, and A. Courville. Improved training of wasserstein gans. 2017.
- [2] Y. Jiang. Applications of generative adversarial networks in materials science. *Materials Genome Engineering Advances*, 2, 2024.
- [3] Y. Xin, E. Walia, and P. Babyn. Generative adversarial network in medical imaging: a review. *Medical Image Analysis*, 58:101552, 2019.
- [4] S. Kazeminia, C. Baur, A. Kuijper, B. Ginneken, N. Navab, S. Albarqouni, and A. Mukhopadhyay. Gans for medical image analysis. *Artificial Intelligence in Medicine*, 109:101938, 2020.
- [5] M. Frid-Adar, I. Diamant, E. Klang, M. M. Amitai, J. Goldberger, and H. Greenspan. Gan-based synthetic medical image augmentation for increased cnn performance in liver lesion classification. *Neurocomputing*, 321:321–331, 2018.
- [6] R. Abdal, Y. Qin, and P. Wonka. Image2stylegan: how to embed images into the stylegan latent space? *2019 IEEE/CVF International Conference on Computer Vision (ICCV)*, 2019.
- [7] K. Cho, T. Raiko, A. Ilin, and J. Karhunen. A two-stage pretraining algorithm for deep boltzmann machines. pages 106–113, 2013.
- [8] G. Montavon. Learning feature hierarchies with centered deep boltzmann machines. 2012.

- [9] S. Niazi, N. Aadit, M. Mohseni, S. Chowdhury, and Y. Qin. Training deep boltzmann networks with sparse ising machines. 2023.
- [10] R. Salakhutdinov and G. Hinton. An efficient learning procedure for deep boltzmann machines. *Neural Computation*, 24:1967–2006, 2012.
- [11] D. Kingma and M. Welling. An introduction to variational autoencoders. *Foundations and Trends® in Machine Learning*, 12:307–392, 2019.
- [12] B. Kiran, D. Thomas, and R. Parakkal. An overview of deep learning based methods for unsupervised and semi-supervised anomaly detection in videos. *Journal of Imaging*, 4:36, 2018.
- [13] C. Guo, J. Zhou, H. Chen, N. Ying, J. Zhang, and D. Zhou. Variational autoencoder with optimizing gaussian mixture model priors. *Ieee Access*, 8:43992–44005, 2020.
- [14] Y. Varolgunes, T. Bereau, and J. Rudzinski. Interpretable embeddings from molecular simulations using gaussian mixture variational autoencoders. *Machine Learning Science and Technology*, 1:015012, 2020.
- [15] S. Portillo. Dimensionality reduction of sdss spectra with variational autoencoders. 2020.
- [16] P. Munjal, A. Paul, and N. Krishnan. Implicit discriminator in variational autoencoder. 2019.
- [17] X. Bie, L. Girin, S. Leglaise, T. Hueber, and X. Alameda-Pineda. A benchmark of dynamical variational autoencoders applied to speech spectrogram modeling. 2021.
- [18] L. Han, Y. Huang, and T. Zhang. Candidates vs. noises estimation for large multi-class classification problem. 2017.
- [19] B. Liu, E. Rosenfeld, P. Ravikumar, and A. Risteski. Analyzing and improving the optimization landscape of noise-contrastive estimation. 2021.

- [20] B. Damavandi, S. Kumar, N. Shazeer, and A. Bruguier. Nn-grams: unifying neural network and n-gram language models for speech recognition. 2016.
- [21] M. Labeau and A. Allauzen. An experimental analysis of noise-contrastive estimation: the noise distribution matters. 2017.
- [22] Y. Song. How to train your energy-based models. 2021.
- [23] J. Parikh, T. Rumbell, X. Butova, T. Myachina, J. Acero, S. Khamzin, O. Solovyova, J. Kozloski, A. Khokhlova, and V. Gurev. Generative adversarial networks for construction of virtual populations of mechanistic models: simulations to study omecamtiv mecarbil action. *Journal of Pharmacokinetics and Pharmacodynamics*, 49:51–64, 2021.
- [24] Y. Saito, S. Takamichi, and H. Saruwatari. Statistical parametric speech synthesis incorporating generative adversarial networks. *Ieee/Acm Transactions on Audio Speech and Language Processing*, 26:84–96, 2018.
- [25] T. Miyato. Cgans with projection discriminator. 2018.
- [26] Ian J. Goodfellow, Jean Pouget-Abadie, Mehdi Mirza, Bing Xu, David Warde-Farley, Sherjil Ozair, Aaron Courville, and Yoshua Bengio. Generative adversarial networks. 2014.
- [27] T. Kynkäanniemi, T. Karras, M. Aittala, T. Aila, and J. Lehtinen. The role of imagenet classes in fréchet inception distance. 2022.
- [28] Y. Xu, T. Wu, J. Charlton, and K. Bennett. Gan training acceleration using fréchet descriptor-based coreset. *Applied Sciences*, 12:7599, 2022.
- [29] R. Xu, J. Wang, J. Liu, F. Ni, and B. Cao. Thermal infrared face image data enhancement method based on deep learning. 2023.
- [30] D. Berthelot, T. Schumm, and L. Metz.Began: boundary equilibrium generative adversarial networks. 2017.

- [31] H. Hyungrok, T. Jun, and D. Kim. Unbalanced gans: pre-training the generator of generative adversarial network using variational autoencoder. 2020.